
Table of Contents

Introduction	1.1
1. A Tour of Computer Systems	1.2

Part 1 Program Structure and Execution

2. Representing and Manipulating Information	2.1
2.55	2.1.1
2.56	2.1.2
2.57	2.1.3
2.58	2.1.4
2.59	2.1.5
2.60	2.1.6
2.61	2.1.7
2.62	2.1.8
2.63	2.1.9
2.64	2.1.10
2.65	2.1.11
2.66	2.1.12
2.67	2.1.13
2.68	2.1.14
2.69	2.1.15
2.70	2.1.16
2.71	2.1.17
2.72	2.1.18
2.73	2.1.19
2.74	2.1.20
2.75	2.1.21
2.76	2.1.22
2.77	2.1.23
2.78	2.1.24
2.79	2.1.25
2.80	2.1.26
2.81	2.1.27
2.82	2.1.28
2.83	2.1.29

2.84	2.1.30
2.85	2.1.31
2.86	2.1.32
2.87	2.1.33
2.88	2.1.34
2.89	2.1.35
2.90	2.1.36
2.91	2.1.37
2.92	2.1.38
2.93	2.1.39
2.94	2.1.40
2.95	2.1.41
2.96	2.1.42
2.97	2.1.43
3. Machine-Level Representation of Programs	2.2
3.58	2.2.1
3.59	2.2.2
3.60	2.2.3
3.61	2.2.4
3.62	2.2.5
3.63	2.2.6
3.64	2.2.7
3.65	2.2.8
3.66	2.2.9
3.67	2.2.10
3.68	2.2.11
3.69	2.2.12
3.70	2.2.13
3.71	2.2.14
3.72	2.2.15
3.73	2.2.16
3.74	2.2.17
3.75	2.2.18
4. Processor Architecture	2.3
4.45	2.3.1
4.46	2.3.2
4.47	2.3.3

4.48	2.3.4
4.49	2.3.5
4.50	2.3.6
4.51	2.3.7
4.52	2.3.8
4.53	2.3.9
4.54	2.3.10
4.55	2.3.11
4.56	2.3.12
4.57	2.3.13
4.58	2.3.14
4.59	2.3.15
5. Optimizing Program Performance	2.4
5.13	2.4.1
5.14	2.4.2
5.15	2.4.3
5.16	2.4.4
5.17	2.4.5
5.18	2.4.6
5.19	2.4.7
6. The Memory Hierarchy	2.5
6.22	2.5.1
6.23	2.5.2
6.24	2.5.3
6.25	2.5.4
6.26	2.5.5
6.27	2.5.6
6.28	2.5.7
6.29	2.5.8
6.30	2.5.9
6.31	2.5.10
6.32	2.5.11
6.33	2.5.12
6.34	2.5.13
6.35	2.5.14
6.36	2.5.15
6.37	2.5.16

6.38	2.5.17
6.39	2.5.18
6.40	2.5.19
6.41	2.5.20
6.42	2.5.21
6.43	2.5.22
6.44	2.5.23
6.45	2.5.24
6.46	2.5.25

Part 2 Running Programs on a System

7. Linking	3.1
7.6	3.1.1
7.7	3.1.2
7.8	3.1.3
7.9	3.1.4
7.10	3.1.5
7.11	3.1.6
7.12	3.1.7
7.13	3.1.8
8. Exceptional Control Flow	3.2
8.9	3.2.1
8.10	3.2.2
8.11	3.2.3
8.12	3.2.4
8.13	3.2.5
8.14	3.2.6
8.15	3.2.7
8.16	3.2.8
8.17	3.2.9
8.18	3.2.10
8.19	3.2.11
8.20	3.2.12
8.21	3.2.13
8.22	3.2.14
8.23	3.2.15
8.24	3.2.16

8.25	3.2.17
8.26	3.2.18
9. Virtual Memory	3.3
9.11	3.3.1
9.12	3.3.2
9.13	3.3.3
9.14	3.3.4
9.15	3.3.5
9.16	3.3.6
9.17	3.3.7
9.18	3.3.8
9.19	3.3.9
9.20	3.3.10

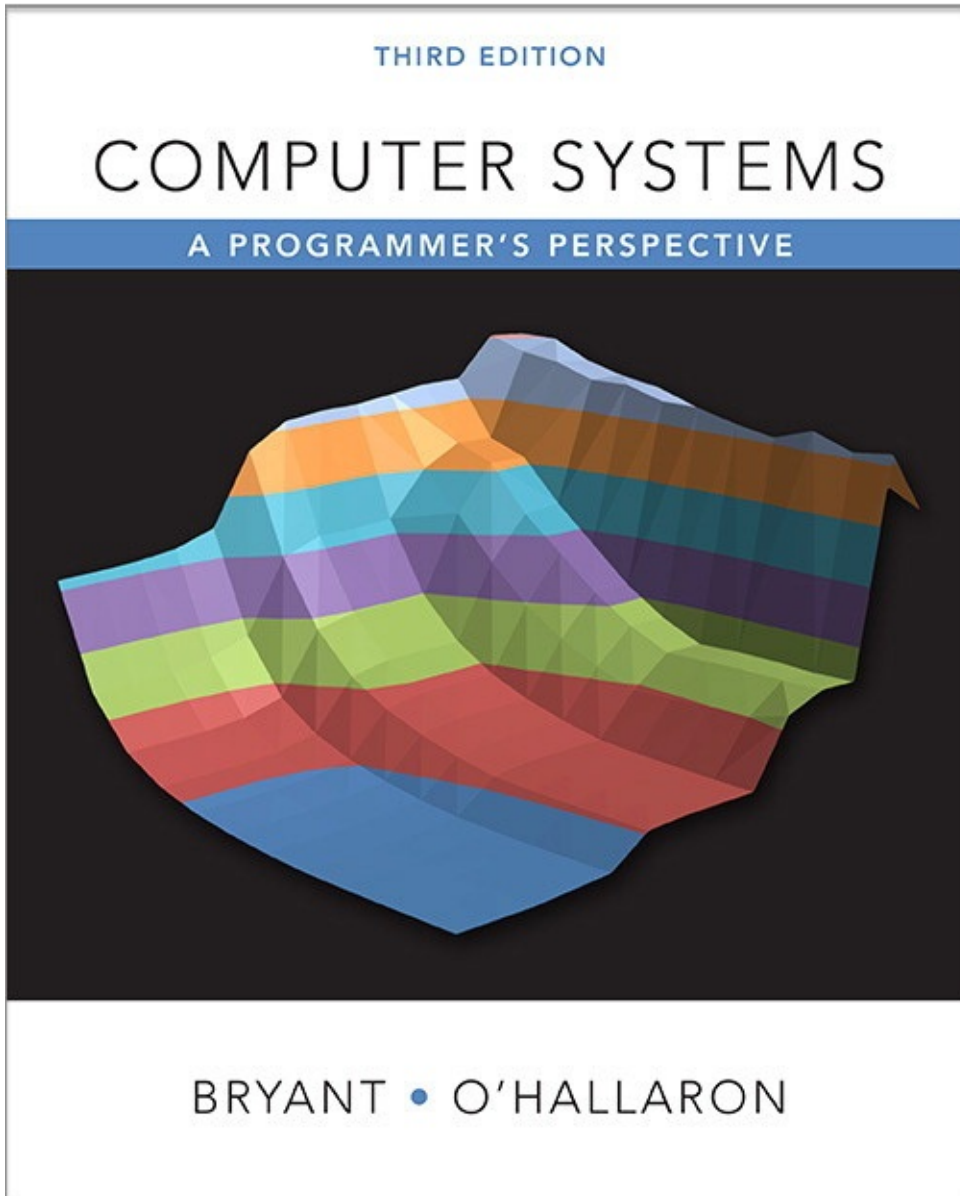
Part 3 Interaction and Communication between Programs

10. System-Level I/O	4.1
10.6	4.1.1
10.7	4.1.2
10.8	4.1.3
10.9	4.1.4
10.10	4.1.5
11. Network Programming	4.2
11.6	4.2.1
11.7	4.2.2
11.8	4.2.3
11.9	4.2.4
11.10	4.2.5
11.11	4.2.6
11.12	4.2.7
11.13	4.2.8
12. Concurrent Programming	4.3
12.16	4.3.1
12.17	4.3.2
12.18	4.3.3
12.19	4.3.4
12.20	4.3.5
12.21	4.3.6

12.22	4.3.7
12.23	4.3.8
12.24	4.3.9
12.25	4.3.10
12.26	4.3.11
12.27	4.3.12
12.28	4.3.13
12.29	4.3.14
12.30	4.3.15
12.31	4.3.16
12.32	4.3.17
12.33	4.3.18
12.34	4.3.19
12.35	4.3.20
12.36	4.3.21
12.37	4.3.22
12.38	4.3.23
12.39	4.3.24

CSAPP-3e-Solutions build passing

Computer Systems: A Programmer's Perspective Third Edition Solutions



at first

Almost all solutions has its own code piece in c/gas/yas and every code piece is tested!

Code files are classified by chapter. Please visit the index page of every chapter to see more info.

issues

Hurry makes work faulty and no improvement makes it disappointed.

Thanks every issue and pr, they really make this project better.

build

prerequisite

- x64 linux system
- docker

pull env image

```
sudo docker pull dreamanddead/csapp-3e-solutions
```

code

clone code

```
git clone https://github.com/DreamAndDead/CSAPP-3e-Solutions.git
cd CSAPP-3e-Solutions
```

compile

```
make
```

test

```
make test
```

clean

```
make clean
```

gitbook

must install gitbook plugins first before other gitbook actions

```
make plugin
```

serve book in <http://localhost:4000>

```
make serve
```

generate site in ./_book/

```
make html
```

generate E-books in ./


```
make pdf  
make mobi  
make epub
```

feedback

If you encounter some problem, you can [email me](#) or comment on Disqus in specific solution page

license

GPLv3

at last

I'll be :) if this little book helps you and make your life more convenient.

Long Live Open Source.

A Tour of Computer Systems

Computer science is no more about computers than astronomy is about telescopes.

by Edsger Dijkstra

no homework here.

Representing and Manipulating Information

Everything is physics and math.

by Katherine Johnson

2.1 - 2.54 visit book

2.55 - 2.97 visit here

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right

solution	code file	test way
2.55	show-bytes.c	output
2.56	show-bytes.c	output
2.57	show-bytes-more.c	output
2.58	is-little-endian.c	assert
2.59	generate-a-word.c	assert
2.60	replace-byte.c	assert
2.61	2.61.c	assert
2.62	int-shifts-are-arithmetic.c	assert
2.63	srl-sra.c	assert
2.64	any-odd-one.c	assert
2.65	odd-ones.c	assert
2.66	leftmost-one.c	assert
2.67	int-size-is-32.c	assert
2.68	lower-one-mask.c	assert
2.69	rotate-left.c	assert
2.70	fits-bits.c	assert
2.71	xbyte.c	assert
2.72	copy-int.c	assert
2.73	saturating-add.c	assert
2.74	tsub-ok.c	assert
2.75	unsigned-high-prod.c	assert
2.76	calloc.c	assert

2.77	2.77.c	assert
2.78	divide-power2.c	assert
2.79	mul3div4.c	assert
2.80	threeforths.c	assert
2.81	2.81.c	assert
2.82	2.81.c	assert
2.83	----	----
2.84	float-le.c	assert
2.85	----	----
2.86	----	----
2.87	----	----
2.88	----	----
2.89	2.89.c	assert
2.90	fpwr2.c	assert
2.91	----	----
2.92	floats/float-negate.c	assert
2.93	floats/float-absval.c	assert
2.94	floats/float-twice.c	assert
2.95	floats/float-half.c	assert
2.96	floats/float-f2i.c	assert
2.97	floats/float-i2f.c	assert

2.55

```
/*
 * show-bytes.c
 */

#include <stdio.h>

typedef unsigned char* byte_pointer;

void show_bytes(byte_pointer start, size_t len) {
    size_t i;
    for (i = 0; i < len; i++) {
        printf(" %.2x", start[i]);
    }
    printf("\n");
}

void show_int(int x) {
    show_bytes((byte_pointer) &x, sizeof(int));
}

void show_float(float x) {
    show_bytes((byte_pointer) &x, sizeof(float));
}

void show_pointer(void* x) {
    show_bytes((byte_pointer) &x, sizeof(void*));
}

void test_show_bytes(int val) {
    int ival = val;
    float fval = (float) ival;
    int* pval = &ival;

    show_int(ival);
    show_float(fval);
    show_pointer(pval);
}

int main(int argc, char* argv[]) {
    int test_num = 328;

    test_show_bytes(test_num);
    return 0;
}
```

uname -mr:

```
4.4.26-gentoo x86_64
```

compile:

```
gcc -m64 show-bytes.c -o show-bytes
```

run:

```
./show-bytes
```

output:

```
48 01 00 00
00 00 a4 43
a8 1e 71 ee fc 7f 00 00
```

2.56

change

```
int test_num = 1024;
```

run:

```
./show-bytes
```

output:

```
00 04 00 00
00 00 80 44
c8 fe 83 2f fc 7f 00 00
```

try more integers :)

2.57

```
/*
 * show-bytes.c
 */

#include <stdio.h>

typedef unsigned char* byte_pointer;

void show_bytes(byte_pointer start, size_t len) {
    size_t i;
    for (i = 0; i < len; i++) {
        printf(" %.2x", start[i]);
    }
    printf("\n");
}

void show_int(int x) {
    show_bytes((byte_pointer) &x, sizeof(int));
}

void show_float(float x) {
    show_bytes((byte_pointer) &x, sizeof(float));
}

void show_pointer(void* x) {
    show_bytes((byte_pointer) &x, sizeof(void*));
}

//=====
// 2.57 changes
//=====
void show_short(short x) {
    show_bytes((byte_pointer) &x, sizeof(short));
}

void show_long(long x) {
    show_bytes((byte_pointer) &x, sizeof(long));
}

void show_double(double x) {
    show_bytes((byte_pointer) &x, sizeof(double));
}

//=====
// 2.57 changes end
//=====

void test_show_bytes(int val) {
    int ival = val;
    float fval = (float) ival;
    int* pval = &ival;
```



```
show_int(ival);
show_float(fval);
show_pointer(pval);

//=====
// 2.57 changes
//=====
short sval = (short) ival;
long lval = (long) ival;
double dval = (double) ival;

show_short(sval);
show_long(lval);
show_double(dval);
//=====
// 2.57 changes end
//=====
}

int main(int argc, char* argv[]) {
    int test_num = 328;

    test_show_bytes(test_num);
    return 0;
}
```

uname -rm

```
4.4.0-21-generic x86_64
```

run

```
./show-bytes-more
```

output

```
48 01 00 00
00 00 a4 43
18 b7 2e 20 fd 7f 00 00
48 01
48 01 00 00 00 00 00 00
00 00 00 00 00 80 74 40
```

2.58

```
/*
 * is-little-endian.c
 */

#include <stdio.h>
#include <assert.h>

typedef unsigned char* byte_pointer;

int is_little_endian() {
    int test_num = 0xff;
    byte_pointer byte_start = (byte_pointer) &test_num;

    if (byte_start[0] == 0xff) {
        return 1;
    }
    return 0;
}

int main(int argc, char* argv[]) {
    assert(is_little_endian());
    return 0;
}
```

2.59

expression

```
(x & 0xFF) | (y & ~0xFF)
```

try it

```
/*
 * generate-a-word.c
 */

#include <stdio.h>
#include <assert.h>

int main(int argc, char* argv[]) {
    size_t mask = 0xff;
    size_t x = 0x89ABCDEF;
    size_t y = 0x76543210;

    size_t res = (x & mask) | (y & ~mask);
    assert(res == 0x765432EF);

    return 0;
}
```

2.60

```
/*
 * replace-byte.c
 */

#include <stdio.h>
#include <assert.h>

unsigned replace_byte(unsigned x, int i, unsigned char b) {
    if (i < 0) {
        printf("error: i is negetive\n");
        return x;
    }
    if (i > sizeof(unsigned)-1) {
        printf("error: too big i");
        return x;
    }

    // 1 byte has 8 bits, << 3 means * 8
    unsigned mask = ((unsigned) 0xFF) << (i << 3);
    unsigned pos_byte = ((unsigned) b) << (i << 3);

    return (x & ~mask) | pos_byte;
}

int main(int argc, char *argv[]) {
    unsigned rep_0 = replace_byte(0x12345678, 0, 0xAB);
    unsigned rep_3 = replace_byte(0x12345678, 3, 0xAB);

    assert(rep_0 == 0x123456AB);
    assert(rep_3 == 0xAB345678);
    return 0;
}
```

2.61

A

```
!~x
```

B

```
!x
```

C

```
!~(x | ~0xff)
```

D

```
!((x >> ((sizeof(int)-1) << 3)) & 0xff)
```

test it

```
/*
 * 2.61.c
 */

#include <stdio.h>
#include <assert.h>

int A(int x) {
    return !~x;
}

int B(int x) {
    return !x;
}

int C(int x) {
    return A(x | ~0xff);
}

int D(int x) {
    return B((x >> ((sizeof(int)-1) << 3)) & 0xff);
}

int main(int argc, char* argv[]) {
    int all_bit_one = ~0;
    int all_bit_zero = 0;

    assert(A(all_bit_one));
    assert(!B(all_bit_one));
    assert(C(all_bit_one));
    assert(!D(all_bit_one));

    assert(!A(all_bit_zero));
    assert(B(all_bit_zero));
    assert(!C(all_bit_zero));
    assert(D(all_bit_zero));

    // test magic number 0x1234ff
    assert(!A(0x1234ff));
    assert(!B(0x1234ff));
    assert(C(0x1234ff));
    assert(D(0x1234ff));

    // test magic number 0x1234
    assert(!A(0x1234));
    assert(!B(0x1234));
    assert(!C(0x1234));
    assert(D(0x1234));

    return 0;
}
```


2.62

```
/*
 * int-shifts-are-arithmetic.c
 */

#include <stdio.h>
#include <assert.h>

int int_shifts_are_arithmetic() {
    int num = -1;
    return !(num ^ (num >> 1));
}

int main(int argc, char* argv[]) {
    assert(int_shifts_are_arithmetic());
    return 0;
}
```



```
/*
 * srl-sra.c
 */

#include <stdio.h>
#include <assert.h>

unsigned srl(unsigned x, int k) {
    unsigned xsra = (int) x >> k;

    int w = sizeof(int) << 3;
    int mask = (int) -1 << (w - k);
    return xsra & ~mask;
}

int sra(int x, int k) {
    int xsrl = (unsigned) x >> k;

    int w = sizeof(int) << 3;
    int mask = (int) -1 << (w - k);
    //let mask remain unchanged when the first bit of x is 1, otherwise 0.
    int m = 1 << (w - 1);
    mask &= ! (x & m) - 1;
    return xsrl | mask;
}

int main(int argc, char* argv[]) {
    unsigned test_unsigned = 0x12345678;
    int test_int = 0x12345678;

    assert(srl(test_unsigned, 4) == test_unsigned >> 4);
    assert(sra(test_int, 4) == test_int >> 4);

    test_unsigned = 0x87654321;
    test_int = 0x87654321;

    assert (srl (test_unsigned, 4) == test_unsigned >> 4);
    assert (sra (test_int, 4) == test_int >> 4);

    return 0;
}
```

2.64

```
/*
 * any-odd-one.c
 */
#include <stdio.h>
#include <assert.h>

int any_odd_one(unsigned x) {
    return !(0xAAAAAAAA & x);
}

int main(int argc, char* argv[]) {
    assert(any_odd_one(0x2));
    assert(!any_odd_one(0x4));
    return 0;
}
```

2.65

thanks [this answer on stackoverflow](#)

```
/*
 * odd-ones.c
 */
#include <stdio.h>
#include <assert.h>

int odd_ones(unsigned x) {
    x ^= x >> 16;
    x ^= x >> 8;
    x ^= x >> 4;
    x ^= x >> 2;
    x ^= x >> 1;
    x &= 0x1;
    return x;
}

int main(int argc, char* argv[]) {
    assert(odd_ones(0x10101011));
    assert(!odd_ones(0x01010101));
    return 0;
}
```

2.66

```
/*
 * leftmost-one.c
 */
#include <stdio.h>
#include <assert.h>

/*
 * Generate mask indicating leftmost 1 in x. Assume w=32
 * For example, 0xFF00 -> 0x8000, and 0x6000 -> 0x4000.
 * If x = 0, then return 0
 */
int leftmost_one(unsigned x) {
    /*
     * first, generate a mask that all bits after leftmost one are one
     * e.g. 0xFF00 -> 0xFFFF, and 0x6000 -> 0x7FFF
     * If x = 0, get 0
     */
    x |= x >> 1;
    x |= x >> 2;
    x |= x >> 4;
    x |= x >> 8;
    x |= x >> 16;
    /*
     * then, do (mask >> 1) + (mask && 1), in which mask && 1 deals with case x = 0, reserve leftmost bit one
     * that's we want
     */
    return (x >> 1) + (x && 1);
}

int main(int argc, char* argv[]) {
    assert(leftmost_one(0xFF00) == 0x8000);
    assert(leftmost_one(0x6000) == 0x4000);
    assert(leftmost_one(0x0) == 0x0);
    assert(leftmost_one(0x80000000) == 0x80000000);
    return 0;
}
```

2.67

A.

In section 6.5.7 Bitwise shift operators of [c11 standard](#), it said

If the value of the right operand is negative or is greater than or equal to the width of the promoted left operand, the behavior is undefined.

B.

see function `int_size_is_32`

C.

see function `int_size_is_32_for_16bit`

```
/*
 * int-size-is-32.c
 */

#include <stdio.h>
#include <assert.h>

/* The following function does not run properly on some machine */
/*
int bad_int_size_is_32() {
    int set_msb = 1 << 31;
    int beyond_msb = 1 << 32;

    return set_msb && !beyond_msb;
}
*/

int int_size_is_32() {
    int set_msb = 1 << 31;
    int beyond_msb = set_msb << 1;

    return set_msb && !beyond_msb;
}

int int_size_is_32_for_16bit() {
    int set_msb = 1 << 15 << 15 << 1;
    int beyond_msb = set_msb << 1;

    return set_msb && !beyond_msb;
}

int main(int argc, char *argv[]) {
    assert(int_size_is_32());
    assert(int_size_is_32_for_16bit());
    return 0;
}
```


2.68

```
/*
 * lower-one-mask.c
 */
#include <stdio.h>
#include <assert.h>

/*
 * Mask with least significant n bits set to 1
 * Example: n = 6 -> 0x3F, n = 17 -> 0x1FFFF
 * Assume 1 <= n <= w
 */
int lower_one_mask(int n) {
    int w = sizeof(int) << 3;
    return (unsigned) -1 >> (w - n);
}

int main(int argc, char* argv[]) {
    assert(lower_one_mask(6) == 0x3F);
    assert(lower_one_mask(17) == 0x1FFFF);
    assert(lower_one_mask(32) == 0xFFFFFFFF);
    return 0;
}
```

2.69.md

```
/*
 * rotate-left.c
 */
#include <stdio.h>
#include <assert.h>

/*
 * Do rotate left shift. Assume 0 <= n < w
 * Example when x = 0x12345678 and w = 32:
 *   n = 4 -> 0x23456781, n = 20 -> 0x67812345
 */
unsigned rotate_left(unsigned x, int n) {
    int w = sizeof(unsigned) << 3;
    /* pay attention when n == 0 */
    return x << n | x >> (w - n - 1) >> 1;
}

int main(int argc, char* argv[]) {
    assert(rotate_left(0x12345678, 4) == 0x23456781);
    assert(rotate_left(0x12345678, 20) == 0x67812345);
    return 0;
}
```


2.70

```
/*
 * fits-bits.c
 */
#include <stdio.h>
#include <assert.h>

int fits_bits(int x, int n) {
    /*
     * 1 <= n <= w
     *
     * assume w = 8, n = 3
     * if x > 0
     * 0b00000010 is ok, 0b00001010 is not, and 0b00000110 is not yet (thanks itardc@16
3.com)
     * if x < 0
     * 0b11111100 is ok, 0b10111100 is not, and 0b11111000 is not yet
     *
     * the point is
     * x << (w-n) >> (w-n) must be equal to x itself.
     */
    int w = sizeof(int) << 3;
    int offset = w - n;
    return (x << offset >> offset) == x;
}

int main(int argc, char* argv[]) {
    assert(!fits_bits(0xFF, 8));
    assert(!fits_bits(-0xFF, 8));

    assert(fits_bits(0b0010, 3));
    assert(!fits_bits(0b1010, 3));
    assert(!fits_bits(0b0110, 3));

    assert(fits_bits(~0b11, 3));
    assert(!fits_bits(~0b01000011, 3));
    assert(!fits_bits(~0b111, 3));
    return 0;
}
```

2.71

A.

This function can't extract a negative byte number from word

B.

```
/*
 * xbyte.c
 */
#include <stdio.h>
#include <assert.h>

typedef unsigned packet_t;

int xbyte(packet_t word, int bytenum) {
    /*
     * pay attention when byte we want is negative
     *
     * Assume sizeof(unsigned) is 4
     * first shift left 8 * (4 - 1 - bytenum)
     * then arithmetic shift right 8 * (4 - 1) reserve significant bit
     */
    int size = sizeof(unsigned);
    int shift_left_val = (size - 1 - bytenum) << 3;
    int shift_right_val = (size - 1) << 3;
    return (int) word << shift_left_val >> shift_right_val;
}

int main(int argc, char* argv[]) {
    assert(xbyte(0xAABBCCDD, 1) == 0xFFFFFCC);
    assert(xbyte(0x00112233, 2) == 0x11);
    return 0;
}
```

2.72

A.

sizeof(val) return type size_t, it usually is a kind of unsigned type.

calculate maxbytes-sizeof(val) get size_t type value, and it always ≥ 0

B.

```
/*
 * copy-int.c
 */
#include <stdio.h>
#include <assert.h>
#include <string.h>
#include <stdlib.h>

void copy_int(int val, void* buf, int maxbytes) {
    /* compare two signed number, avoid someone set maxbytes a negetive value */
    if (maxbytes >= (int) sizeof(val)) {
        memcpy(buf, (void*)&val, sizeof(val));
    }
}

int main(int argc, char* argv[]) {
    int maxbytes = sizeof(int) * 10;
    void* buf = malloc(maxbytes);
    int val;

    val = 0x12345678;
    copy_int(val, buf, maxbytes);
    assert(*(int*)buf == val);

    val = 0xAABBCCDD;
    copy_int(val, buf, 0);
    assert(*(int*)buf != val);

    free(buf);
    return 0;
}
```

2.73

I can't figure out a elegant solution :(

thanks <https://zhangjunphy.github.io/csapp/chap2.html> :)

```
/*
 * saturating-add.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>

int saturating_add(int x, int y) {
    int sum = x + y;
    int sig_mask = INT_MIN;
    /*
     * if x > 0, y > 0 but sum < 0, it's a positive overflow
     * if x < 0, y < 0 but sum >= 0, it's a negetive overflow
     */
    int pos_over = !(x & sig_mask) && !(y & sig_mask) && (sum & sig_mask);
    int neg_over = (x & sig_mask) && (y & sig_mask) && !(sum & sig_mask);

    (pos_over && (sum = INT_MAX) || neg_over && (sum = INT_MIN));

    return sum;
}

int main(int argc, char* argv[]) {
    assert(INT_MAX == saturating_add(INT_MAX, 0x1234));
    assert(INT_MIN == saturating_add(INT_MIN, -0x1234));
    assert(0x11 + 0x22 == saturating_add(0x11, 0x22));
    return 0;
}
```

2.74

```
/*
 * tsub-ok.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>

/* Determine whether arguments can be subtracted without overflow */
int tsub_ok(int x, int y)
{
    int res = 1;

    (y == INT_MIN) && (res = 0);
    // if (y == INT_MIN) res = 0;

    int sub = x - y;
    int pos_over = x > 0 && y < 0 && sub < 0;
    int neg_over = x < 0 && y > 0 && sub > 0;

    res = res && !(pos_over || neg_over);

    return res;
}

int main(int argc, char* argv[]) {
    assert(!tsub_ok(0x00, INT_MIN));
    assert(tsub_ok(0x00, 0x00));
    return 0;
}
```

2.75.md

```
/*
 * unsigned-high-prod.c
 */
#include <stdio.h>
#include <assert.h>
#include <inttypes.h>

int signed_high_prod(int x, int y) {
    int64_t mul = (int64_t) x * y;
    return mul >> 32;
}

unsigned unsigned_high_prod(unsigned x, unsigned y) {
    /* TODO calculations */
    int sig_x = x >> 31;
    int sig_y = y >> 31;
    int signed_prod = signed_high_prod(x, y);
    return signed_prod + x * sig_y + y * sig_x;
}

/* a theoretically correct version to test unsigned_high_prod func */
unsigned another_unsigned_high_prod(unsigned x, unsigned y) {
    uint64_t mul = (uint64_t) x * y;
    return mul >> 32;
}

int main(int argc, char* argv[]) {
    unsigned x = 0x12345678;
    unsigned y = 0xFFFFFFFF;

    assert(another_unsigned_high_prod(x, y) == unsigned_high_prod(x, y));
    return 0;
}
```

2.76.md

```
/*
 * calloc.c
 */
#include <stdio.h>
#include <assert.h>
#include <stdlib.h>
#include <stdint.h>
#include <string.h>

/* rename to avoid conflict */
void* another_calloc(size_t nmemb, size_t size) {
    if (nmemb == 0 || size == 0) {
        return NULL;
    }
    size_t buf_size = nmemb * size;
    /* a good way to check overflow or not */
    if (nmemb == buf_size / size) {
        void* ptr = malloc(buf_size);
        if(ptr != NULL) {
            memset(ptr, 0, buf_size);
        }
        return ptr;
    }
    return NULL;
}

int main(int argc, char* argv[]) {
    void* p;
    p = another_calloc(0x1234, 1);
    assert(p != NULL);
    free(p);

    p = another_calloc(SIZE_MAX, 2);
    assert(p == NULL);
    free(p);
    return 0;
}
```

2.77

```
/*
 * 2.77.c
 */
#include <stdio.h>
#include <assert.h>

/* K = 17 */
int A(int x) {
    return (x << 4) + x;
}

/* K = -7 */
int B(int x) {
    return x - (x << 3);
}

/* K = 60 */
int C(int x) {
    return (x << 6) - (x << 2);
}

/* K = -112 */
int D(int x) {
    return (x << 4) - (x << 7);
}

int main(int argc, char* argv[]) {
    int x = 0x87654321;
    assert(A(x) == 17 * x);
    assert(B(x) == -7 * x);
    assert(C(x) == 60 * x);
    assert(D(x) == -112 * x);
    return 0;
}
```


2.78

```
/*
 * divide-power2.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>

/*
 * Divide by power of 2, -> x/2^k
 * Assume 0 <= k < w-1
 */
int divide_power2(int x, int k) {
    int is_neg = x & INT_MIN;
    (is_neg && (x = x + (1 << k) - 1));
    return x >> k;
}

int main(int argc, char* argv[]) {
    int x = 0x80000007;
    assert(divide_power2(x, 1) == x / 2);
    assert(divide_power2(x, 2) == x / 4);
    return 0;
}
```

2.79

```
/*
 * mul3div4.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>

/*
 * code from 2.78
 *
 * Divide by power of 2, -> x/2^k
 * Assume 0 <= k < w-1
 */
int divide_power2(int x, int k) {
    int is_neg = x & INT_MIN;
    (is_neg && (x = x + (1 << k) - 1));
    return x >> k;
}

int mul3div4(int x) {
    int mul3 = (x << 1) + x;
    return divide_power2(mul3, 2);
}

int main(int argc, char* argv[]) {
    int x = 0x87654321;
    assert(mul3div4(x) == x * 3 / 4);
    return 0;
}
```

```

/*
 * threeforths.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>

/*
 * calculate 3/4x, no overflow, round to zero
 *
 * no overflow means divide 4 first, then multiple 3, different from 2.79 here
 *
 * rounding to zero is a little complicated.
 * every int x, equals f(first 30 bit number) plus 1(last 2 bit number)
 *
 *   f = x & ~0x3
 *   l = x & 0x3
 *   x = f + l
 *   threeforths(x) = f/4*3 + l*3/4
 *
 * f doesn't care about round at all, we just care about rounding from l*3/4
 *
 *   lm3 = (l << 1) + 1
 *
 * when x > 0, rounding to zero is easy
 *
 *   lm3d4 = lm3 >> 2
 *
 * when x < 0, rounding to zero acts like divide_power2 in 2.78
 *
 *   bias = 0x3    // (1 << 2) - 1
 *   lm3d4 = (lm3 + bias) >> 2
 */
int threeforths(int x) {
    int is_neg = x & INT_MIN;
    int f = x & ~0x3;
    int l = x & 0x3;

    int fd4 = f >> 2;
    int fd4m3 = (fd4 << 1) + fd4;

    int lm3 = (l << 1) + 1;
    int bias = (1 << 2) - 1;
    (is_neg && (lm3 += bias));
    int lm3d4 = lm3 >> 2;

    return fd4m3 + lm3d4;
}

int main(int argc, char* argv[]) {
    assert(threeforths(8) == 6);
}

```

```
assert(threeforths(9) == 6);
assert(threeforths(10) == 7);
assert(threeforths(11) == 8);
assert(threeforths(12) == 9);

assert(threeforths(-8) == -6);
assert(threeforths(-9) == -6);
assert(threeforths(-10) == -7);
assert(threeforths(-11) == -8);
assert(threeforths(-12) == -9);
return 0;
}
```

2.81

A.

```
-1 << k
```

B.

```
~(-1 << k) << j
```

```
/*
 * 2.81.c
 */
#include <stdio.h>
#include <assert.h>

/* Assume 0 <= k < w */
int A(int k) {
    return -1 << k;
}

/* Assume 0 <= j, k < w */
int B(int k, int j) {
    return ~A(k) << j;
}

int main(int argc, char* argv[]) {
    assert(A(8) == 0xFFFFF00);
    assert(B(16, 8) == 0x00FFFF00);
    return 0;
}
```

2.82

A.

wrong, when x is INT_MIN

B.

right

C.

right

D.

right

E.

right

```
/*
 * 2.82.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>
#include "lib/random.h"

/* broken when x is INT_MIN */
int A(int x, int y) {
    return (x < y) == (-x > -y);
}

/*
 * right
 *
 * ((x + y) << 4) + y - x
 * =>
 * x << 4 - x + y << 4 + y
 * =>
 * x*16 - x + y*16 + y
 * whether overflow or not, =>
 * x*15 + y*17
 */
int B(int x, int y) {
    return ((x + y) << 4) + y - x == 17 * y + 15 * x;
}

/*
 * right
 *
 * ~x + ~y + 1
 * =>
```

```

*  $\sim x + 1 + \sim y + 1 - 1$ 
*  $\Rightarrow$ 
*  $-x + -y - 1$ 
*  $\Rightarrow$ 
*  $-(x + y) - 1$ 
*  $\Rightarrow$ 
*  $\sim(x + y) + 1 - 1$ 
*  $\Rightarrow$ 
*  $\sim(x + y)$ 
*/
int C(int x, int y) {
    return  $\sim x + \sim y + 1 == \sim(x + y)$ ;
}

/*
* right
*
*  $(ux - uy) == -(unsigned) (y - x)$ 
*  $\Rightarrow$ 
*  $-(ux - uy) == (unsigned) (y - x)$ 
*  $\Rightarrow$ 
*  $(ux - uy) == (unsigned) (x - y)$ 
*/
int D(int x, int y) {
    unsigned ux = (unsigned) x;
    unsigned uy = (unsigned) y;

    return  $(ux - uy) == -(unsigned) (y - x)$ ;
}

/*
* right
*
*  $x \gg 2 \ll 2$ 
*  $\Rightarrow$ 
*  $x \& \sim 0x3$ 
*  $\Rightarrow$ 
*  $x - \text{num}(00/01/10/11)$ 
*  $\Rightarrow$ 
*  $((x \gg 2) \ll 2) \leq x$ 
*/
int E(int x, int y) {
    return  $((x \gg 2) \ll 2) \leq x$ ;
}

int main(int argc, char* argv[]) {
    init_seed();
    int x = random_int();
    int y = random_int();

    assert(!A(INT_MIN, 0));
    assert(B(x, y));
    assert(C(x, y));
}

```

```
    assert(D(x, y));  
    assert(E(x, y));  
    return 0;  
}
```


2.83

A.

$$n = 0.yyyyy\dots$$

$$n \ll k = y.yyyyy\dots = Y + n$$

$$n \ll k - n = Y$$

$$n = Y/(2^k - 1)$$

B.

(a).

$$y = 101, Y = 5, k = 3$$

$$n = 5/7$$

(b).

$$y = 0110, Y = 6, k = 4$$

$$n = 2/5$$

(c).

$$y = 010011, Y = 19, k = 6$$

$$n = 19/63$$

2.84

thanks [czy1996](#)

```
/*
 * float-le.c
 */
#include <stdio.h>
#include <assert.h>

unsigned f2u(float x) {
    return *(unsigned*)&x;
}

int float_le(float x, float y) {
    unsigned ux = f2u(x);
    unsigned uy = f2u(y);

    unsigned sx = ux >> 31;
    unsigned sy = uy >> 31;

    // ref: https://github.com/DreamAndDead/CSAPP-3e-Solutions/issues/1
    return (ux << 1 == 0 && uy << 1 == 0) || /* both zeros */
        (sx && !sy) || /* x < 0, y >= 0 or x <= 0, y > 0 */
        (!sx && !sy && ux <= uy) || /* x > 0, y >= 0 or x >= 0, y > 0 */
        (sx && sy && ux >= uy); /* x < 0, y <= 0 or x <= 0, y < 0 */
}

int main(int argc, char* argv[]) {
    assert(float_le(-0, +0));
    assert(float_le(+0, -0));
    assert(float_le(0, 3));
    assert(float_le(-4, -0));
    assert(float_le(-4, 4));
    return 0;
}
```

2.85

$$\text{bias} = 2^{(k-1)} - 1$$

$$V = 2^E * M$$

A.

$$7.0 = 0b111.000\dots$$

$$M = 0b1.11, f = 0b0.11, E = 2, e = \text{bias} + E, V = 7.0$$

bits

```
0 10...01 110...
```

B.

Assume $\text{bias} \gg n$ biggest odd number, M must be $0b1.11111\dots$, $f = 0b0.1111111\dots(n \text{ bits } 1)$

$$E = n, V = 0b1111111\dots(n+1 \text{ bits } 1) = 2^{(n+1)} - 1$$

bits

```
0 bias+n 11111...
```

C.

least standard numberM must be $0b1.00\dots$, $f = 0b0.000\dots$, $E = 1 - \text{bias}$

$$V = 2^{(1-\text{bias})}$$

reciprocal

$$V = 2^{(\text{bias}-1)}$$

$$M = 0b1.0000\dots, f = 0b0.000\dots, E = \text{bias}-1, e = \text{bias} + E$$

bits

```
0 11...101 00000...
```

2.86

$\text{bias} = 2^{(15-1)} - 1$

description	binary	decimal
least positive unstandard	0 0000...(15) 0 000...(62)1	$2^{(1-\text{bias}-63)}$
least positive standard	0 000...(14)1 1 000....(63)	$2^{(1-\text{bias})}$
bigest standard	0 111...(14)0 1 111...(63)	$2^{\text{bias}} * (2-2^{-63})$

2.87

Desc	Hex	M	E	V	D
-0	0x8000	0	-14	-0	-0.0
>2 least	0x4001	1025/1024	1	1025/512	2.00195312
512	0x6000	1	9	512	512.0
biggest denormalized	0x03FF	1023/1024	-14	1023/(2^24)	6.09755516e-5
-∞	0xFC00	-	-	-∞	-∞
0x3BB0	0x3BB0	123/64	-1	123/128	0.9609375

2.88

A bit	A value	B bit	B value
1 01110 001	-9/16	1 0110 0010	-9/16
0 10110 101	$13 \cdot 2^4$	0 1110 1010	$13 \cdot 2^4$
1 00111 110	$-7/2^{10}$	1 0000 0111	$-7/2^{10}$
0 00000 101	$5/2^{11}$	0 0000 0001	$1/2^{10}$
1 11011 000	-2^{12}	1 1110 1111	$-31 \cdot 2^3$
0 11000 100	$3 \cdot 2^8$	0 1111 0000	+∞

```
/*
 * 2.89.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>
#include "lib/random.h"

/*
 * most important thing is that all double number come from ints
 */

/* right */
int A(int x, double dx) {
    return (float)x == (float)dx;
}

/* wrong when y is INT_MIN */
int B(int x, double dx, int y, double dy) {
    return dx-dy == (double)(x-y);
}

/* right */
int C(double dx, double dy, double dz) {
    return (dx+dy)+dz == dx+(dy+dz);
}

/*
 * wrong
 *
 * FIXME I don't know what conditions cause false
 */
int D(double dx, double dy, double dz) {
    return (dx*dy)*dz == dx*(dy*dz);
}

/* wrong when dx != 0 and dz == 0 */
int E(double dx, double dz) {
    return dx/dx == dz/dz;
}

int main(int argc, char* argv[]) {
    init_seed();

    int x = random_int();
    int y = random_int();
    int z = random_int();
    double dx = (double)x;
    double dy = (double)y;
    double dz = (double)z;
```

```
printf("%x %x %x\n", x, y, z);

assert(A(x, dx));
assert(!B(0, (double)(int)0, INT_MIN, (double)(int)INT_MIN));
assert(C(dx, dy, dz));
/* magic number, brute force attack */
assert(!D((double)(int)0x64e73387, (double)(int)0xd31cb264, (double)(int)0xd22f1fcd))
;
assert(!E(dx, (double)(int)0));
return 0;
}
```



```
/*
 * fpwr2.c
 */
#include <stdio.h>
#include <assert.h>
#include <math.h>

float u2f(unsigned x) {
    return *(float*) &x;
}

/* 2^x */
float fpwr2(int x) {
    /* Result exponent and fraction */
    unsigned exp, frac;
    unsigned u;

    if (x < 2-pow(2,7)-23) {
        /* too small. return 0.0 */
        exp = 0;
        frac = 0;
    } else if (x < 2-pow(2,7)) {
        /* Denormalized result */
        exp = 0;
        frac = 1 << (unsigned)(x - (2-pow(2,7)-23));
    } else if (x < pow(2,7)-1+1) {
        /* Normalized result */
        exp = pow(2,7)-1+x;
        frac = 0;
    } else {
        /* Too big, return +oo */
        exp = 0xFF;
        frac = 0;
    }

    /* pack exp and frac into 32 bits */
    u = exp << 23 | frac;
    /* Result as float */
    return u2f(u);
}

int main(int argc, char* argv[]) {
    assert(fpwr2(0) == powf(2,0));
    assert(fpwr2(100) == powf(2,100));
    assert(fpwr2(-100) == powf(2,-100));
    assert(fpwr2(10000) == powf(2,10000));
    assert(fpwr2(-10000) == powf(2,-10000));
    return 0;
}
```


2.91

A.

```
0x40490FDB
```

```
0 10000000 10010010000111111011011
```

float number

```
0b11.0010010000111111011011
```

B.

ref 2.83

```
0b11.001001(001)...
```

C.

9th

2.92

```
/*
 * float-negate.c
 */
#include <stdio.h>
#include <assert.h>
#include "float-negate.h"

float_bits float_negate(float_bits f) {
    unsigned sig = f >> 31;
    unsigned exp = f >> 23 & 0xFF;
    unsigned frac = f & 0x7FFFFFFF;

    int is_NAN = (exp == 0xFF) && (frac != 0);
    if (is_NAN) {
        return f;
    }

    return ~sig << 31 | exp << 23 | frac;
}
```

2.93

```
/*
 * float-absval.c
 */
#include <stdio.h>
#include <assert.h>
#include "float-absval.h"

float_bits float_absval(float_bits f) {
    unsigned sig = f >> 31;
    unsigned exp = f >> 23 & 0xFF;
    unsigned frac = f & 0x7FFFFFFF;

    int is_NAN = (exp == 0xFF) && (frac != 0);
    if (is_NAN) {
        return f;
    }

    return 0 << 31 | exp << 23 | frac;
}
```

2.94

```
/*
 * float-twice.c
 */
#include <stdio.h>
#include <assert.h>
#include "float-twice.h"

float_bits float_twice(float_bits f) {
    unsigned sig = f >> 31;
    unsigned exp = f >> 23 & 0xFF;
    unsigned frac = f & 0x7FFFFFFF;

    int is_NAN_or_oo = (exp == 0xFF);
    if (is_NAN_or_oo) {
        return f;
    }

    if (exp == 0) {
        /* Denormalized */
        frac <<= 1;
    } else if (exp == 0xFF - 1) {
        /* twice to oo */
        exp = 0xFF;
        frac = 0;
    } else {
        /* Normalized */
        exp += 1;
    }

    return sig << 31 | exp << 23 | frac;
}
```

```
/*
 * float-half.c
 */
#include <stdio.h>
#include <assert.h>
#include "float-half.h"

float_bits float_half(float_bits f) {
    unsigned sig = f >> 31;
    unsigned rest = f & 0x7FFFFFFF;
    unsigned exp = f >> 23 & 0xFF;
    unsigned frac = f & 0x7FFFFF;

    int is_NAN_or_oo = (exp == 0xFF);
    if (is_NAN_or_oo) {
        return f;
    }

    /*
     * round to even, we care about last 2 bits of frac
     *
     * 00 => 0 just >>1
     * 01 => 0 (round to even) just >>1
     * 10 => 1 just >>1
     * 11 => 1 + 1 (round to even) just >>1 and plus 1
     */
    int addition = (frac & 0x3) == 0x3;

    if (exp == 0) {
        /* Denormalized */
        frac >>= 1;
        frac += addition;
    } else if (exp == 1) {
        /* Normalized to denormalized */
        rest >>= 1;
        rest += addition;
        exp = rest >> 23 & 0xFF;
        frac = rest & 0x7FFFFF;
    } else {
        /* Normalized */
        exp -= 1;
    }

    return sig << 31 | exp << 23 | frac;
}
```

2.96

this is exactly how my machine does!!!

```

/*
 * float-f2i.c
 */
#include <stdio.h>
#include <assert.h>
#include "float-f2i.h"

/*
 * Compute (float) f
 * If conversion cause overflow or f is NaN, return 0x80000000
 */
int float_f2i(float_bits f) {
    unsigned sig = f >> 31;
    unsigned exp = f >> 23 & 0xFF;
    unsigned frac = f & 0x7FFFFFFF;
    unsigned bias = 0x7F;

    int num;
    unsigned E;
    unsigned M;

    /*
     * consider positive numbers
     *
     * 0 00000000 000000000000000000000000
     * ==>
     * 0 01111111 000000000000000000000000
     * 0 <= f < 1
     * get integer 0
     *
     * 0 01111111 000000000000000000000000
     * ==>
     * 0 (01111111+31) 000000000000000000000000
     * 1 <= f < 2^31
     * integer round to 0
     *
     * 0 (01111111+31) 000000000000000000000000
     * ==>
     * greater
     * 2^31 <= f < oo
     * return 0x80000000
     */
    if (exp >= 0 && exp < 0 + bias) {
        /* number less than 1 */
        num = 0;
    } else if (exp >= 31 + bias) {
        /* number overflow */
        /* or f < 0 and (int)f == INT_MIN */
        num = 0x80000000;
    }
}

```



```
    } else {  
        E = exp - bias;  
        M = frac | 0x800000;  
        if (E > 23) {  
            num = M << (E - 23);  
        } else {  
            /* whether sig is 1 or 0, round to zero */  
            num = M >> (23 - E);  
        }  
    }  
  
    return sig ? -num : num;  
}
```

```
/*
 * float-i2f.c
 */
#include <stdio.h>
#include <assert.h>
#include <limits.h>
#include "float-i2f.h"

/*
 * Assume i > 0
 * calculate i's bit length
 *
 * e.g.
 * 0x3 => 2
 * 0xFF => 8
 * 0x80 => 8
 */
int bits_length(int i) {
    if ((i & INT_MIN) != 0) {
        return 32;
    }

    unsigned u = (unsigned)i;
    int length = 0;
    while (u >= (1<<length)) {
        length++;
    }
    return length;
}

/*
 * generate mask
 * 00000...(32-1) 11111...(1)
 *
 * e.g.
 * 3  => 0x00000007
 * 16 => 0x0000FFFF
 */
unsigned bits_mask(int l) {
    return (unsigned) -1 >> (32-l);
}

/*
 * Compute (float) i
 */
float_bits float_i2f(int i) {
    unsigned sig, exp, frac, rest, exp_sig /* except sig */, round_part;
    unsigned bits, fbits;
    unsigned bias = 0x7F;

    if (i == 0) {
```

```

    sig = 0;
    exp = 0;
    frac = 0;
    return sig << 31 | exp << 23 | frac;
}
if (i == INT_MIN) {
    sig = 1;
    exp = bias + 31;
    frac = 0;
    return sig << 31 | exp << 23 | frac;
}

sig = 0;
/* 2's complatation */
if (i < 0) {
    sig = 1;
    i = -i;
}

bits = bits_length(i);
fbits = bits - 1;
exp = bias + fbits;

rest = i & bits_mask(fbits);
if (fbits <= 23) {
    frac = rest << (23 - fbits);
    exp_sig = exp << 23 | frac;
} else {
    int offset = fbits - 23;
    int round_mid = 1 << (offset - 1);

    round_part = rest & bits_mask(offset);
    frac = rest >> offset;
    exp_sig = exp << 23 | frac;

    /* round to even */
    if (round_part < round_mid) {
        /* nothing */
    } else if (round_part > round_mid) {
        exp_sig += 1;
    } else {
        /* round_part == round_mid */
        if ((frac & 0x1) == 1) {
            /* round to even */
            exp_sig += 1;
        }
    }
}

return sig << 31 | exp_sig;
}

```


Machine-Level Representation of Programs

To understand a program, you must become both the machine and the program.

by Alan Perlis

3.1 - 3.57 visit book

3.58 - 3.75 visit here

test

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right

solution	code file	test way
3.58	decode2/(decode.c, decode2.s, main.c)	assert
3.59	-----	-----
3.60	loop/loop.s, loop2.c, main.c	assert
3.61	cread/(cread-alt.c, cread-alt.s, cread.s)	assert
3.62	-----	-----
3.63	-----	-----
3.64	store-ele/(store-ele.s, store-ele2.c, main.c)	assert
3.65	-----	-----
3.66	sum-col/(sum-col.s, sum-col2.c, main.c)	assert
3.67	-----	-----
3.68	-----	-----
3.69	-----	-----
3.70	-----	-----
3.71	good-echo/good-echo.c	output
3.72	-----	-----
3.73	3.73.c	assert
3.74	3.74.c	assert
3.75	-----	-----

3.58

```
/*
 * decode.c
 */
long decode(long x, long y, long z) {
    long tmp = y - z;
    return (tmp * x) ^ (tmp << 63 >> 63);
}
```

3.59

assume

$$ux = x + x_{63}2^{64}$$

$$uy = y + y_{63}2^{64}$$

multiple

$$ux \cdot uy = (x + x_{63}2^{64}) \cdot (y + y_{63}2^{64})$$

$$ux \cdot uy = x \cdot y + (x_{63}y + y_{63}x)2^{64} + x_{63}y_{63}2^{128}$$

2¹²⁸ overflows, don't care about it

$$x \cdot y = ux \cdot uy - (x_{63}y + y_{63}x)2^{64}$$

```
# void store_prod(int128_t* dest, int64_t x, int64_t y)
# dest in %rdi, x in %rsi, y in %rdx
store_prod:
    movq %rdx, %rax    # %rax = y
    cqto               # (int128_t)y, %rdx = (-1)y_63
    movq %rsi, %rcx    # %rcx = x

    # x >> 63, if x == 1, %rcx = -1; if x_63 == 0, %rcx = 0
    # %rcx = (-1)x_63
    sarq $63, %rcx

    # pay attention, imulq behaves differently according to param number(1/2)
    imulq %rax, %rcx    # %rcx = y * -x_63
    imulq %rsi, %rdx    # %rdx = x * -y_63
    addq %rdx, %rcx     # %rcx = x * -y_63 + y * -x_63
    mulq %rsi           # %rdx:%rax <= ux * uy

    # lower 64 bits are same for x * y and ux * uy. ref (2.18) on book
    # %rdx = ux * uy(high 64 bits) - (x_{63}y + y_{63}x)2^{64}
    addq %rcx, %rdx

    movq %rax, (%rdi)   # set lower 64bits
    movq %rdx, 8(%rdi)  # set higher 64bits
    ret
```


3.60

A.

val	reg
x	%rdi
n	%esi
result	%rax
mask	%rdx

B.

```
result = 0

mask = 1
```

C.

```
mask != 0
```

D.

```
mask = mask << n
```

E.

```
/*
 * loop2.c
 */

long loop2(long x, int n) {
    long result = 0;
    long mask;
    for (mask = 1; mask != 0; mask <=< n) {
        result |= (x & mask);
    }
    return result;
}
```

3.61

function `cread_alt` should be

```
/*
 * cread-alt.c
 */
#include <stdio.h>
#include <assert.h>

long cread(long *xp) {
    return (xp ? *xp : 0);
}

long cread_alt(long *xp) {
    return (!xp ? 0 : *xp);
}

int main(int argc, char* argv[]) {
    long a = 0;
    assert(cread(&a) == cread_alt(&a));
    assert(cread(NULL) == cread_alt(NULL));
    return 0;
}
```

In book sample, compile function `cread` get

```
# long cread(long *xp)
# xp in %rdi
cread:
    movq (%rdi), %rax
    testq %rdi, %rdi
    movl $0, %edx
    cmovne %rdx, %rax
    ret
```

compile `cread_alt` may get

```
# long cread_alt(long* xp)
# xp in %rdi
cread_alt:
    movl $0, %eax
    testq %rdi, %rdi
    cmovne (%rdi), %rax
```

3.62

```
/*
 * 3.62.c
 */

/* Enumerated type creates set of constants numbered 0 and upward */
typedef enum { MODE_A, MODE_B, MODE_C, MODE_D, MODE_E } mode_t;

long switch3(long *p1, long *p2, mode_t action) {
    long result = 0;
    switch(action) {
        case MODE_A:
            result = *p2;
            *p2 = *p1;
            break;
        case MODE_B:
            *p1 = *p1 + *p2;
            result = *p1;
            break;
        case MODE_C:
            *p1 = 59;
            result = *p2;
            break;
        case MODE_D:
            *p1 = *p2;
            result = 27;
            break;
        case MODE_E:
            result = 27;
            break;
        default:
            result = 12;
            break;
    }
    return result;
}
```

3.63

```
/*
 * 3.63.c
 */

long switch_prob(long x, long n) {
    long result = x;
    switch(n) {
        /* Fill in code here */
        case 60:
        case 62:
            result = x * 8;
            break;
        case 63:
            result = x >> 3;
            break;
        case 64:
            x = x << 4 - x;
        case 65:
            x = x * x;
        default:
            result = x + 0x4B;
    }
    return result;
}
```

3.64

A.

3.1 in book

```
T D[R][C];

&D[i][j] = Xd + L(C*i + j)

T means type, D means data, R means row, C means column

L means sizeof(T), Xd means address of D
```

similarly, in 3d array

```
TYPE D[R][S][T]

&D[i][j][k] = Xd + L(S*T*i + T*j + k)
```

B.

```
.section .data
.global A
A:
    .fill 3640/8, 8, 121 # fill data 121

.section .text
.global store_ele
# long store_ele(long i, long j, long k, long *dest)
# i in %rdi, j in %rsi, k in %rdx, dest in %rcx
store_ele:
    leaq (%rsi, %rsi, 2), %rax    # t1 = j*3
    leaq (%rsi, %rax, 4), %rax    # t1 = j*13
    movq %rdi, %rsi              # t2 = i
    salq $6, %rsi                # t2 = i*64
    addq %rsi, %rdi               # t3 = i*65
    addq %rax, %rdi               # t3 = i*65 + j*13
    addq %rdi, %rdx               # t4 = i*65 + j*13 + k
    movq A(,%rdx,8), %rax         # t1 = *(A + 8*t4)
    movq %rax, (%rcx)             # *dest = t1
    movl $3640, %eax              # return 3640
    ret
```

base on comments,

```
S * T = 65
T = 13
8*R*S*T = 3640
```

so

R = 7
S = 5
T = 13

3.65

A.

```
&A[i][j] in %rdx
```

B.

```
&A[j][i] in %rax
```

C.

```
addq $8, %rdx # means A[i][j] -> A[i][j+1]
```

```
addq $120, %rax # means A[j][i] -> A[j+1][i], 120 == 8*M
```

M = 15

3.66

```

.section .text
.global sum_col
# long sum_col(long n, long A[NR(n)][NC(n)], long j)
# n in %rdi, A in %rsi, j in %rdx
sum_col:
    leaq 1(,%rdi,4), %r8      # t1 = n*4 + 1
    leaq (%rdi,%rdi,2), %rax   # t2 = n*3
    movq %rax, %rdi           # t3 = n*3
    testq %rax, %rax          # test n*3
    jle .L4                   # n*3 <= 0, jump .L4
    salq $3, %r8              # t1 = t1*8 = 8*(n*4 + 1)
    leaq (%rsi,%rdx,8), %rcx   # t4 = j*8 + A
    movl $0, %eax             # t2 = 0
    movl $0, %edx             # t5 = 0
.L3:
    addq (%rcx), %rax          # t2 = *(t4) = *(A + j*8)
    addq $1, %rdx              # t5 = t5+1
    addq %r8, %rcx             # t4 = t1+t4 = A + j*8 + 8*(n*4 + 1)
    cmpq %rdi, %rdx           # cmp t5 & t3
    jne .L3                   # if t5 != n*3, loop
    rep
    ret
.L4:
    movl $0, %eax             # return 0
    ret

```

base on comments in asm code

```

cmpq %rdi, %rdx      # cmp t5 & t3
jne .L3              # if t5 != n*3, loop

```

t5 is var i, so `NR(n) == n*3`

```

leaq 1(,%rdi,4), %r8      # t1 = n*4 + 1
.....
salq $3, %r8              # t1 = t1*8 = 8*(n*4 + 1)
.....
addq %r8, %rcx            # t4 = t1+t4 = A + j*8 + 8*(n*4 + 1)

```

in every loop, pointer move `8*(n*4 + 1)` bytes, so `NC(n) == n*4 + 1`

thanks [gonglinyuan](#)

3.67

```
/*
 * 3.67.c
 */

typedef struct {
    long a[2];
    long *p;
} strA;

typedef struct {
    long u[2];
    long q;
} strB;

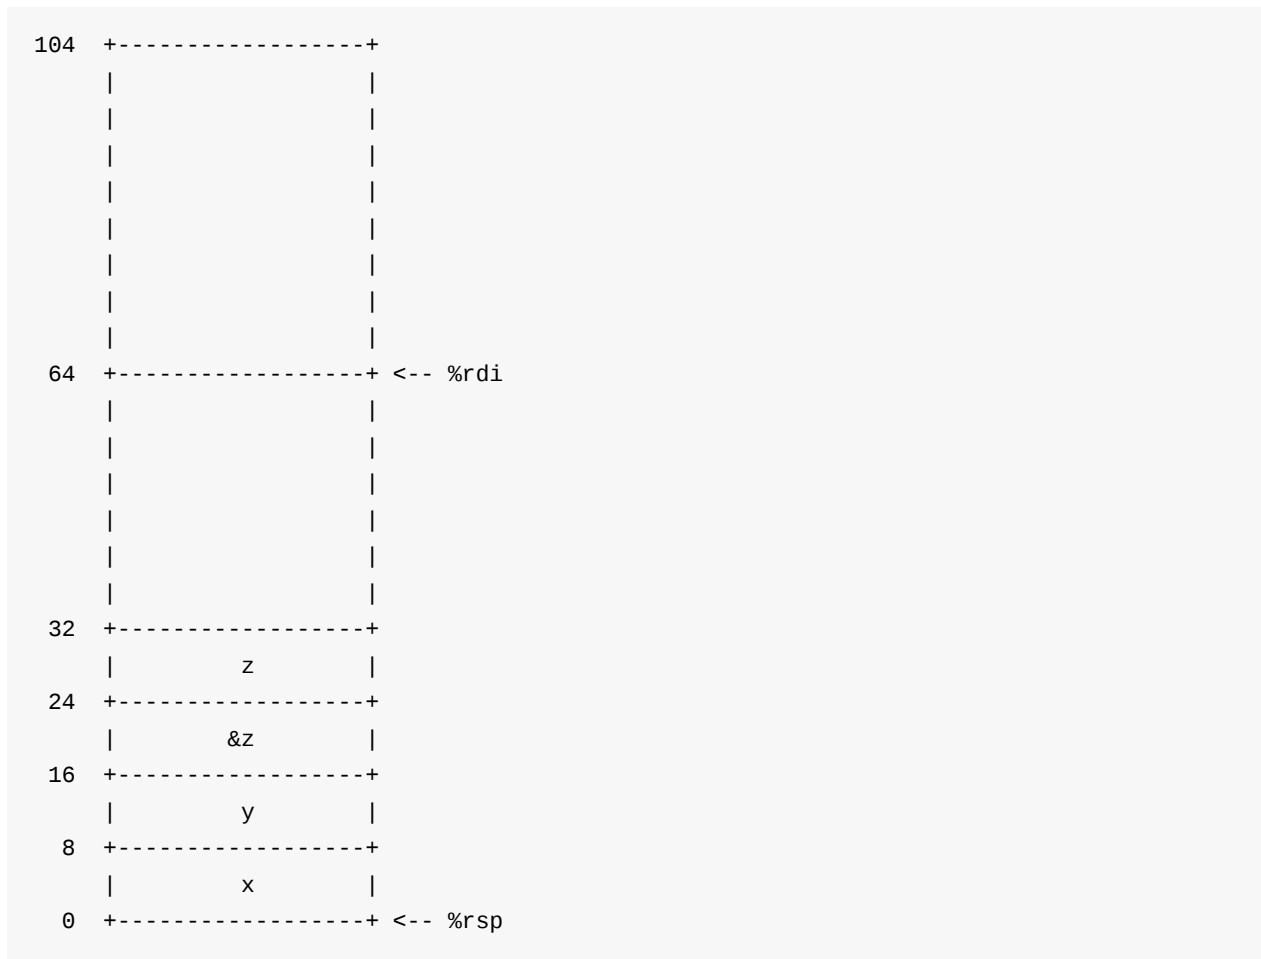
strB process(strA a) {
    strB r;
    r.u[0] = s.a[1];
    r.u[1] = s.a[0];
    r.q = *s.p;
    return r;
}

long eval(long x, long y, long z) {
    strA s;
    s.a[0] = x;
    s.a[1] = y;
    s.p = &z;
    strB r = process(s);
    return r.u[0] + r.u[1] + r.q;
}
```

```
# strB process(strA s)
# s in %rdi
process:
    movq %rdi, %rax
    movq 24(%rsp), %rdx
    movq (%rdx), %rdx
    movq 16(%rsp), %rcx
    movq %rcx, (%rdi)
    movq 8(%rsp), %rcx
    movq %rcx, 8(%rdi)
    movq %rdx, 16(%rdi)
    ret

# long eval(long x, long y, long z)
# x in %rdi, y in %rsi, z in %rdx
eval:
    subq $104, %rsp
    movq %rdx, 24(%rsp)
    leaq 24(%rsp), %rax
    movq %rdi, (%rsp)
    movq %rsi, 8(%rsp)
    movq %rax, 16(%rsp)
    leaq 64(%rsp), %rdi
    call process
    movq 72(%rsp), %rax
    addq 64(%rsp), %rax
    addq 80(%rsp), %rax
    addq $104, %rsp
    ret
```

A.



B.

eval pass a new address `%rsp+64` to process

C.

process access `s` by `%rsp+offset`, not by `%rdi`

D.

eval pass address `%rsp+64` to process, process store data from here as beginning, finally return this address

E.

```

104 +-----+
    |           |
    |           |
    |           |
    |           |
    |           |
    |           |
88  +-----+
    |         z         |
80  +-----+
    |         x         |
72  +-----+
    |         y         |
64  +-----+ <-- %rdi(eval pass in)
    |           | \
    |           |  -- %rax(process pass out)
    |           |
    |           |
    |           |
32  +-----+
    |         z         |
24  +-----+
    |        &z         |
16  +-----+
    |         y         |
8   +-----+
    |         x         |
0   +-----+ <-- %rsp in eval
    |           |
-8  +-----+ <-- %rsp in process

```

F.

caller find space and pass space address to callee, callee store data on this space area and return this address

3.68

```
# void setVal(str1* p, str2* q)
# p in %rdi, q in %rsi
setVal:
    # 8(%rsi) fetch q->t, int t is aligned by 4, so 4 < B <=8
    movslq 8(%rsi), %rax

    # 32(%rsi) fetch q->u, long u is aligned by 8
    # offset q->s is offset q->t + 4, so 24 < 12 + A*2 <= 32
    addq 32(%rsi), %rax

    # 184(%rdi) fetch p->v, long y is aligned by 8, so 176 < A*B*4 <= 184
    movq %rax, 184(%rdi)
    ret
```

```
4 < B <= 8
5 < A <= 10
44 < A*B <= 46
```

only

A = 9

B = 5

3.69

```

/*
 * 3.69.c
 */

typedef struct {
    int first;
    a_struct a[CNT];
    int last;
} b_struct;

void test(long i, b_struct *bp) {
    int n = bp->first + bp->last;
    a_struct *ap = &bp->a[i];
    ap->x[ap->idx] = n;
}

```

```

# void test(long i, b_struct *bp)
# i in %rdi, bp in %rsi
test:
    mov 0x120(%rsi), %ecx      # bp+0x120 fetch bp->last
    add (%rsi), %ecx          # bp->first + bp->last
    lea (%rdi,%rdi,4), %rax    # i*5
    lea (%rsi,%rax,8), %rax    # bp+i*40

    # ap = &bp->a[i] = bp+i*40+8, +8 means skip int first
    # so a_struct is aligned by 8, size is 40
    # check last instruction, %rdx here saves value ap->idx!!!
    # so in a_struct, idx is first element
    mov 0x8(%rax), %rdx

    movslq %ecx, %rcx          # n = bp->first + bp->last, convert to long

    # save n to address 8*(ap->idx) + bp+i*40+0x8 + 0x8 (0x10)
    # bp+i*40+0x8 means ap
    # ap + 0x8 means &(ap->x)
    # ap + 0x8 + 8*(ap->idx) means &(ap->x[ap->idx])
    # second element of a_struct is x, x is an array of long
    # 8*(ap->idx) means idx is also long type
    mov %rcx, 0x10(%rax,%rdx,8)

    # size of a_struct is 40 and aligned by 8
    # so array x has 4 long elements
    # finally, a_struct is
    #     typedef struct {
    #         long idx,
    #         long x[4]
    #     } a_struct
    retq

```

A.

```
7*40 + 8 = 288 = 0x120 , so
```

CNT = 7

thanks <https://github.com/zagortenay333>

B.

```
typedef struct {  
    long idx,  
    long x[4]  
} a_struct
```

3.70

A.

val	offset
e1.p	0
e1.y	8
e2.x	0
e2.next	8

B.

16

C.

```
# void proc(union ele *up)
# up in %rdi
proc:
    # %rax = *(up+8), don't know it's next or y
    movq 8(%rdi), %rax

    # %rdx = *( *(up+8) ), %rax stands for a pointer
    # so *( *(up+8) ) means *(up->e2.next)
    movq (%rax), %rdx

    # %rdx = *( *(up->e2.next) )
    # %rdx is treated as a pointer
    # so %rdx stores *( *(up->e2.next).e1.p )
    movq (%rdx), %rdx

    # %rax stores *(up+8)
    # %rax is treated as a pointer
    # so %rax = *( up->e2.next ), stands for another union ele's address
    #
    # in subq, %rdx is a long number
    # *( *(up->e2.next)+8 ) must be a long number
    # so 8(%rax) means *(up->e2.next).e1.y
    subq 8(%rax), %rdx

    # %rdi never changes in previous instructions
    # instruction below is the final assignment
    # so (%rdi) means up->e2.x
    movq %rdx, (%rdi)
    ret
```

base on comments


```
/*
 * 3.70.c
 */

union ele {
    struct {
        long *p;
        long y;
    } e1;
    struct {
        long x;
        union ele *next;
    } e2;
};

void proc(union ele *up) {
    /* up->  = *(      ) -      ; */
    up->e2.x = *( *(up->e2.next).e1.p ) - *(up->e2.next).e1.y
}
```

3.71

```
/*
 * good-echo.c
 */
#include <stdio.h>
#include <assert.h>
#define BUF_SIZE 12

void good_echo(void) {
    char buf[BUF_SIZE];
    while(1) {
        /* function fgets is interesting */
        char* p = fgets(buf, BUF_SIZE, stdin);
        if (p == NULL) {
            break;
        }
        printf("%s", p);
    }
    return;
}

int main(int argc, char* argv[]) {
    good_echo();
    return 0;
}
```

3.72

```

/*
 * 3.72.c
 */

#include <alloca.h>

long aframe(long n, long idx, long *q) {
    long i;
    long **p = alloca(n * sizeof(long*));
    p[0] = &i;
    for (i = 1; i < n; i++) {
        p[i] = q;
    }
    return *p[idx];
}

```

```

# long aframe(long n, long idx, long *q)
# n in %rdi, idx in %rsi, q in %rdx
aframe:
    pushq %rbp
    movq %rsp, %rbp
    subq $16, %rsp
    leaq 30(,%rdi,8), %rax
    andq $-16, %rax
    subq %rax, %rsp
    leaq 15(%rsp), %r8
    andq $-16, %r8

```

A.

$$s_2 = s_1 - [(n * 8 + 30) \& 0xFFFFFFFF0]$$

when n is odd

$$s_2 = s_1 - (n * 8 + 24)$$

when n is even

$$s_2 = s_1 - (n * 8 + 16)$$

B.

$$p = (s_2 + 15) \& 0xFFFFFFFF0$$

the least multiple of 16 that greater than s2

C.

which	e1	n	s1
least	1	even	$n \% 16 == 1$
greatest	24	odd	$n \% 16 == 0$

least:

e1 can't be 0, if $e1 == 0$, p should equal to s2

when n is even, $e1 + e2 == 16$, if e2 is 15, e1 will be least that can reach, so $s1 == n$ that $n \% 16 == 1$

greatest:

when n is odd, $e1 + e2 == 24$. if $p == s2$, $e2 == 0$, so e1 will be the greatest value 24 that can reach. $s1 == n$ that $n \% 16 == 0$

D.

p is aligned by 16

s2 is the least multiple of 16 that preserve $8 * n$ size space

3.73

```
/*
 * 3.73.c
 */
#include <stdio.h>
#include <assert.h>

typedef enum {NEG, ZERO, POS, OTHER} range_t;

range_t find_range(float x) {
    __asm__(
        "vxorps %xmm1, %xmm1, %xmm1\n\t"
        "vucomiss %xmm1, %xmm0\n\t"
        "jp .P\n\t"
        "ja .A\n\t"
        "jb .B\n\t"
        "je .E\n\t"
        ".A:\n\t"
        "movl $2, %eax\n\t"
        "jmp .Done\n\t"
        ".B:\n\t"
        "movl $0, %eax\n\t"
        "jmp .Done\n\t"
        ".E:\n\t"
        "movl $1, %eax\n\t"
        "jmp .Done\n\t"
        ".P:\n\t"
        "movl $3, %eax\n\t"
        ".Done:\n\t"
    );
}

int main(int argc, char* argv[]) {
    range_t n = NEG, z = ZERO, p = POS, o = OTHER;
    assert(o == find_range(0.0/0.0));
    assert(n == find_range(-2.3));
    assert(z == find_range(0.0));
    assert(p == find_range(3.33));
    return 0;
}
```

3.74

```
/*
 * 3.74.c
 */
#include <stdio.h>
#include <assert.h>

typedef enum {NEG, ZERO, POS, OTHER} range_t;

range_t find_range(float x) {
    __asm__(
        "vxorps %xmm1, %xmm1, %xmm1\n\t"
        "movq $1, %rax\n\t"
        "movq $2, %r8\n\t"
        "movq $0, %r9\n\t"
        "movq $3, %r10\n\t"
        "vucomiss %xmm1, %xmm0\n\t"
        "cmovaq %r8, %rax\n\t"
        "cmovbq %r9, %rax\n\t"
        "cmovpq %r10, %rax\n\t"
        );
}

int main(int argc, char* argv[]) {
    range_t n = NEG, z = ZERO, p = POS, o = OTHER;
    assert(o == find_range(0.0/0.0));
    assert(n == find_range(-2.3));
    assert(z == find_range(0.0));
    assert(p == find_range(3.33));
    return 0;
}
```

3.75

A.

param n	real	img
1	%xmm0	%xmm1
2	%xmm2	%xmm3
3	%xmm4	%xmm5
n

B.

return %xmm0 as real part, %xmm1 as img part

Processor Architecture

The speed at which modern CPUs perform computations still blows my mind daily.

by Markus Persson

4.1 - 4.44 visit book

4.45 - 4.59 visit here

yas simulation

This chapter focus on processor architecture and design a little simple processor and yas -- a assemble language -- designed running on it.

you can access the processor simulation code and simulation manual from csapp official site.

[simulation code](#)

[simulator manual](#)

I have saved the simulation code in directory chapter4/code/sim.

Highly recommend you read the manual and README in code to know how it works and how to test yas code.

test

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right
- yas means using simulator test script to test simulator itself

solution	code file	test way
4.45	--	--
4.46	--	--
4.47	bubble-sort/bubble-sort-pointer.*	output, assert
4.48	bubble-sort/bubble-sort-pointer-3-cmove.*	output
4.49	bubble-sort/bubble-sort-pointer-1-cmove.*	output
4.50	switch/*	output
4.51	--	--
4.52	sim/seq/seq-full.hcl	yas
4.53	sim/pipe/pip-stall.hcl	yas
4.54	sim/pipe/pip-full.hcl	yas
4.55	sim/pipe/pip-nt.hcl	yas
4.56	sim/pipe/pip-btfnt.hcl	yas
4.57	sim/pipe/pip-lf.hcl	yas
4.58	sim/pipe/pip-1w.hcl	yas
4.59	--	--

4.45

```
subq $8, %rsp  
movq REG, (%rsp)
```

A.

No

if REG is %rsp, we push %rsp-8 instead of %rsp into stack

B.

```
movq REG, -8(%rsp)  
subq $8, %rsp
```

4.46

```
movq (%rsp), REG  
addq $8, %rsp
```

A.

No

if REG is %rsp, `movq (%rsp), REG` pop the right value into %rsp, but `addq $8, %rsp` modify it

B.

```
addq $8, %rsp  
movq 8(%rsp), REG
```

4.47

A.

Change into pointer version

```
/*
 * bubble-sort-pointer.c
 */

void bubble_p(long* data, long count) {
    long *i, *last;
    for (last = data+count-1; last > data; last--) {
        for (i = data; i < last; i++) {
            if (*(i+1) < *i) {
                /* swap adjacent elements */
                long t = *(i+1);
                *(i+1) = *i;
                *i = t;
            }
        }
    }
}
```

B.

Compile bubble-sort-pointer.c

```
gcc -S -Og bubble-sort-pointer.c
```

get gas file

```

.file      "bubble-sort-pointer.c"
.text
.globl     bubble_p
.type      bubble_p, @function
bubble_p:
.LFB0:
.cfi_startproc
    leaq    -8(%rdi,%rsi,8), %rsi
    jmp     .L2
.L4:
    movq    8(%rax), %rdx
    movq    (%rax), %rcx
    cmpq    %rcx, %rdx
    jge     .L3
    movq    %rcx, 8(%rax)
    movq    %rdx, (%rax)
.L3:
    addq    $8, %rax
    jmp     .L5
.L6:
    movq    %rdi, %rax
.L5:
    cmpq    %rsi, %rax
    jb     .L4
    subq    $8, %rsi
.L2:
    cmpq    %rdi, %rsi
    ja     .L6
    rep ret
.cfi_endproc
.LFE0:
.size      bubble_p, .-bubble_p
.ident     "GCC: (Ubuntu 5.4.0-6ubuntu1~16.04.5) 5.4.0 20160609"
.section   .note.GNU-stack,"",@progbits

```

Change it into ys version

```

/* bubble-sort-pointer.ys */
.pos 0
    irmovq stack, %rsp
    call main
    halt

# Array of 4 elements
.align 8
data:
    .quad 0x0000000000000004
    .quad 0x0000000000000003
    .quad 0x0000000000000002
data_end:
    .quad 0x0000000000000001

```

```

main:
    irmovq data,%rdi
    irmovq data_end,%rsi
    call ysBubbleP
    ret

# long ys_bubble_p(long *data, long *end)
# data in %rdi, end in %rsi
ysBubbleP:
    jmp L2
L4:
    mrmovq 8(%rax), %r9
    mrmovq (%rax), %r10
    rrmovq %r9, %r8
    subq %r10, %r8
    jge L3
    rmmovq %r10, 8(%rax)
    rmmovq %r9, (%rax)
L3:
    irmovq $8, %r8
    addq %r8, %rax
    jmp L5
L6:
    rrmovq %rdi, %rax
L5:
    rrmovq %rsi, %r8
    subq %rax, %r8
    jg L4
    irmovq $8, %r8
    subq %r8, %rsi
L2:
    rrmovq %rsi, %r8
    subq %rdi, %r8
    jg L6
    ret

.pos 0x200
stack:

```

main difference is second param of function is `long* end` instead of `long count` for convinient, `long* end` marks the position of last element.

using `yas` assemble it and `yis` run it, see output

```

../sim/misc/yas bubble-sort-pointer.y
../sim/misc/yis bubble-sort-pointer.yo
Stopped in 117 steps at PC = 0x13.  Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%rax:  0x0000000000000000  0x0000000000000020
%rsp:  0x0000000000000000  0x0000000000000200
%rsi:  0x0000000000000000  0x0000000000000018
%rdi:  0x0000000000000000  0x0000000000000018
%r9:   0x0000000000000000  0x0000000000000001
%r10:  0x0000000000000000  0x0000000000000002

Changes to memory:
0x0018:  0x0000000000000004  0x0000000000000001
0x0020:  0x0000000000000003  0x0000000000000002
0x0028:  0x0000000000000002  0x0000000000000003
0x0030:  0x0000000000000001  0x0000000000000004
0x01f0:  0x0000000000000000  0x0000000000000055
0x01f8:  0x0000000000000000  0x0000000000000013

```

initial array order is 4,3,2,1, now 1,2,3,4

4.48

a little easy changes.

```

/* bubble-sort-pointer.yys */
.pos 0
    irmovq stack, %rsp
    call main
    halt

# Array of 4 elements
.align 8
data:
    .quad 0x0000000000000004
    .quad 0x0000000000000003
    .quad 0x0000000000000002
data_end:
    .quad 0x0000000000000001

main:
    irmovq data,%rdi
    irmovq data_end,%rsi
    call ysBubbleP
    ret

# long ys_bubble_p(long *data, long *end)
# data in %rdi, end in %rsi
ysBubbleP:
    jmp L2
L4:
    mrmovq 8(%rax), %r9
    mrmovq (%rax), %r10
    rrmovq %r9, %r8
    subq %r10, %r8
#####
# begin differences
#####
    cmovl %r9, %r11
    cmovl %r10, %r9
    cmovl %r11, %r10
    rmmovq %r9, 8(%rax)
    rmmovq %r10, (%rax)
#####
# end
#####
    irmovq $8, %r8
    addq %r8, %rax
    jmp L5
L6:
    rrmovq %rdi, %rax
L5:
    rrmovq %rsi, %r8
    subq %rax, %r8

```



```

    jg L4
    irmovq $8, %r8
    subq %r8, %rsi
L2:
    rrmovq %rsi, %r8
    subq %rdi, %r8
    jg L6
    ret

.pos 0x200
stack:

```

run test, output

```

../sim/misc/yas bubble-sort-pointer-3-cmove.y
../sim/misc/yis bubble-sort-pointer-3-cmove.yo
Stopped in 129 steps at PC = 0x13.  Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%rax:    0x0000000000000000    0x0000000000000020
%rsp:    0x0000000000000000    0x0000000000000200
%rsi:    0x0000000000000000    0x0000000000000018
%rdi:    0x0000000000000000    0x0000000000000018
%r9:     0x0000000000000000    0x0000000000000002
%r10:    0x0000000000000000    0x0000000000000001
%r11:    0x0000000000000000    0x0000000000000001

Changes to memory:
0x0018:   0x0000000000000004    0x0000000000000001
0x0020:   0x0000000000000003    0x0000000000000002
0x0028:   0x0000000000000002    0x0000000000000003
0x0030:   0x0000000000000001    0x0000000000000004

```

4.49

it's a brilliant idea, 3 xor is a normal swap

if $x = 9$, $y = 10$

```
x = x ^ y  x:9^10; y:10
y = x ^ y  x:9^10; y:9
x = x ^ y  x:10; y:9
```

if change y after step 1

```
tmp = x      x:9; y:10; tmp:9  # store x
x = x ^ y    x:9^10; y:10
y = tmp      x:9^10; y:9      # set y to origin x
y = x ^ y    x:9^10; y:10
x = x ^ y    x:9; y:10
```

no swap happens. that's core in code below

```
/* bubble-sort-pointer.yys */
.pos 0
    irmovq stack, %rsp
    call main
    halt

# Array of 4 elements
.align 8
data:
    .quad 0x0000000000000004
    .quad 0x0000000000000003
    .quad 0x0000000000000002
data_end:
    .quad 0x0000000000000001

main:
    irmovq data,%rdi
    irmovq data_end,%rsi
    call ysBubbleP
    ret

# long ys_bubble_p(long *data, long *end)
# data in %rdi, end in %rsi
ysBubbleP:
    jmp L2
L4:
    mrmovq 8(%rax), %r9
    mrmovq (%rax), %r10
#####
# begin differences
#####
    rrmovq %r9, %r8
```

```

    rrmovq %r10, %r11
    xorq %r9, %r10
    subq %r11, %r8
    cmovge %r11, %r9
    xorq %r10, %r9
    xorq %r9, %r10
    rmmovq %r9, 8(%rax)
    rmmovq %r10, (%rax)
#####
# end
#####
    irmovq $8, %r8
    addq %r8, %rax
    jmp L5
L6:
    rrmovq %rdi, %rax
L5:
    rrmovq %rsi, %r8
    subq %rax, %r8
    jg L4
    irmovq $8, %r8
    subq %r8, %rsi
L2:
    rrmovq %rsi, %r8
    subq %rdi, %r8
    jg L6
    ret

.pos 0x200
stack:

```

test and output, watch memory changes:

```

../sim/misc/yas bubble-sort-pointer-1-cmove.ys
../sim/misc/yis bubble-sort-pointer-1-cmove.yo
Stopped in 141 steps at PC = 0x13.  Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%rax:    0x0000000000000000    0x0000000000000020
%rsp:    0x0000000000000000    0x0000000000000200
%rsi:    0x0000000000000000    0x0000000000000018
%rdi:    0x0000000000000000    0x0000000000000018
%r9:     0x0000000000000000    0x0000000000000002
%r10:    0x0000000000000000    0x0000000000000001
%r11:    0x0000000000000000    0x0000000000000002

Changes to memory:
0x0018:  0x0000000000000004    0x0000000000000001
0x0020:  0x0000000000000003    0x0000000000000002
0x0028:  0x0000000000000002    0x0000000000000003
0x0030:  0x0000000000000001    0x0000000000000004

```


4.50

```
/* switch.ys */
.pos 0
    irmovq stack, %rsp
    call main
    halt

# Array of 4 elements
.align 8
array:
    .quad 0x0000000000000000
    .quad 0x0000000000000000
    .quad 0x0000000000000000
    .quad 0x0000000000000000

main:
    # test number 1, -1, 3, 5
    irmovq array, %r10

    irmovq $1,%rdi
    call switchv
    rmmovq %rax, (%r10)

    irmovq $-1,%rdi
    call switchv
    rmmovq %rax, 8(%r10)

    irmovq $3,%rdi
    call switchv
    rmmovq %rax, 16(%r10)

    irmovq $5,%rdi
    call switchv
    rmmovq %rax, 24(%r10)
    ret

table:
    .quad LD # default branch
    .quad L0 # idx == 0
    .quad L1 # idx == 1
    .quad L2 # idx == 2
    .quad L3 # idx == 3
    .quad L4 # idx == 4
    .quad L5 # idx == 5

# long switchv(long idx)
# idx in %rdi
switchv:
    # constant number
    irmovq $8, %r8
    irmovq $0, %r10
    irmovq $1, %r11
```

```
    irmovq $0, %rax
    irmovq table, %rcx # table address
    rrmovq %rdi, %rdx
    subq %r8, %rdx
    jg def             # idx > 5
    subq %r10, %rdi
    jl def             # idx < 0
mul: # calculate 8 * %rdi
    subq %r10, %rdi
    je addr
    addq %r8, %rcx
    subq %r11, %rdi
    jmp mul
addr: # jump using table address
    addq %r8, %rcx
    mrmovq (%rcx), %rdi
    pushq %rdi
    ret
def: # default branch
    irmovq table, %rcx
    mrmovq (%rcx), %rdi
    pushq %rdi
    ret
L0:
    irmovq $0xaaa, %rax
    ret
L1:
    jmp LD
L2:
    jmp L5
L3:
    irmovq $0xccc, %rax
    ret
L4:
    jmp LD
L5:
    irmovq $0xbbb, %rax
    ret
LD:
    irmovq $0xddd, %rax
    ret

.pos 0x200
stack:
```

test function switchv in main, using idx 1,-1,3,5, store result into array

test and output, watch changes to memory:

```
../sim/misc/yas switch.ys
../sim/misc/yis switch.yo
Stopped in 133 steps at PC = 0x13.  Status 'HLT', CC Z=0 S=0 O=0
Changes to registers:
%rax:  0x0000000000000000  0x0000000000000bbb
%rcx:  0x0000000000000000  0x0000000000000e7
%rdx:  0x0000000000000000  0xfffffffffffffffd
%rsp:  0x0000000000000000  0x0000000000000200
%rdi:  0x0000000000000000  0x00000000000001a8
%r8:   0x0000000000000000  0x0000000000000008
%r11:  0x0000000000000000  0x0000000000000001

Changes to memory:
0x0000:  0x000000000200f430  0x0000000000000ddd
0x0008:  0x0000000003880000  0x0000000000000ddd
0x0010:  0x0000000000000000  0x0000000000000ccc
0x0018:  0x0000000000000000  0x0000000000000bbb
```

4.51

phase	iaddq V,rB
F	icode;ifun = M1[PC]; rA:rB = M1[PC+1]; valC = M8[PC+2]; valP = PC + 10;
D	valB = R[rB]
E	valE = valB + valC; set CC
M	
W	R[rB] = valE
PC	PC = valP

4.52

easy changes if you make 4.51 right.

check whole file `./chapter4/code/sim/seq/seq-full.hcl` with `iaddq` functionality

only watch changes with origin `seq-full.hcl` file

```
(cd ./chapter4/code/sim/seq; diff -u origin-seq-full.hcl seq-full.hcl)
```

```
--- origin-seq-full.hcl    2017-11-09 02:57:43.671935966 +0000
+++ seq-full.hcl          2017-11-09 02:57:43.671935966 +0000
@@ -107,16 +107,16 @@

    bool instr_valid = icode in
        { INOP, IHALT, IRRMOVQ, IIRMOVQ, IRMMOVQ, IMRMVQ,
-          IOPQ, IJXX, ICALL, IRET, IPUSHQ, IPOPQ };
+          IOPQ, IJXX, ICALL, IRET, IPUSHQ, IPOPQ, IIADDQ };

    # Does fetched instruction require a regid byte?
    bool need_regids =
        icode in { IRRMOVQ, IOPQ, IPUSHQ, IPOPQ,
-          IIRMOVQ, IRMMOVQ, IMRMVQ };
+          IIRMOVQ, IRMMOVQ, IMRMVQ, IIADDQ };

    # Does fetched instruction require a constant word?
    bool need_valC =
-    icode in { IIRMOVQ, IRMMOVQ, IMRMVQ, IJXX, ICALL };
+    icode in { IIRMOVQ, IRMMOVQ, IMRMVQ, IJXX, ICALL, IIADDQ };

    ##### Decode Stage #####

@@ -129,7 +129,7 @@

    ## What register should be used as the B source?
    word srcB = [
-    icode in { IOPQ, IRMMOVQ, IMRMVQ } : rB;
+    icode in { IOPQ, IRMMOVQ, IMRMVQ, IIADDQ } : rB;
    icode in { IPUSHQ, IPOPQ, ICALL, IRET } : RRSP;
    1 : RNONE; # Don't need register
    ];
@@ -137,7 +137,7 @@

    ## What register should be used as the E destination?
    word dstE = [
        icode in { IRRMOVQ } && Cnd : rB;
-    icode in { IIRMOVQ, IOPQ } : rB;
+    icode in { IIRMOVQ, IOPQ, IIADDQ } : rB;
    icode in { IPUSHQ, IPOPQ, ICALL, IRET } : RRSP;
    1 : RNONE; # Don't write any register
    ];
@@ -153,7 +153,7 @@

    ## Select input A to ALU
```

```

word aluA = [
    icode in { IRRMOVQ, IOPQ } : valA;
-   icode in { IIRMOVQ, IRMMOVQ, IMRMVQ } : valC;
+   icode in { IIRMOVQ, IRMMOVQ, IMRMVQ, IIADDQ } : valC;
    icode in { ICALL, IPUSHQ } : -8;
    icode in { IRET, IPOPQ } : 8;
    # Other instructions don't need ALU
@@ -162,7 +162,7 @@
    ## Select input B to ALU
    word aluB = [
        icode in { IRMMOVQ, IMRMVQ, IOPQ, ICALL,
-            IPUSHQ, IRET, IPOPQ } : valB;
+            IPUSHQ, IRET, IPOPQ, IIADDQ } : valB;
        icode in { IRRMOVQ, IIRMOVQ } : 0;
        # Other instructions don't need ALU
    ];
@@ -174,7 +174,7 @@
    ];

    ## Should the condition codes be updated?
-bool set_cc = icode in { IOPQ };
+bool set_cc = icode in { IOPQ, IIADDQ };

##### Memory Stage #####

```

4.53

1. data hazard

something handled by data-forward must be handled by stall if no data-forward anymore. so when

```
d_srcA in { e_dstE, M_dstM, M_dstE, W_dstM, W_dstE } ||  
d_srcB in { e_dstE, M_dstM, M_dstE, W_dstM, W_dstE }
```

data hazard happens, we have to insert bubble in phase E and stall phase F&D

load/use hazard

load/use hazard only happens when we use data-forward. no load/use hazard if no data-forward.

pay attention

- previous load/use toggle condition `d_srcA == E_dstM || d_srcB == E_dstM` now toggles data hazard.
- in the beginning, `d_srcA == RNONE == e_dstE == E_dstM == ...etc`, that's not data hazard

so we get:

```
situation: data_hazard  
bool s_data_hazard =  
(  
  (  
    d_srcA != RNONE  &&  
    d_srcA in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }  
  ) ||  
  (  
    d_srcB != RNONE  &&  
    d_srcB in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }  
  )  
)
```

2. ret situation

keep same

```
situation: ret  
bool s_ret = IRET in { D_icode, E_icode, M_icode };
```

3. jxx error

keep same

```
situation: jxx error  
bool s_jxx_error = (E_icode == IJXX && !e_Cnd);
```

4. hazard composition

X means nothing to do; stall means stall; bubble means inserting bubble

num	data	ret	jxx	F	D	E	M	W
1	0	0	0	X	X	X	X	X
2	0	0	1	X	bubble	bubble	X	X
3	0	1	0	stall	bubble	X	X	X
4	1	0	0	stall	stall	bubble	X	X

situation 1: nothing happens, every thing is fine

situation 2: just jxx error, keep same with book

situation 3: just ret, keep same with book

situation 4: just data hazard, stall phase F and D, insert bubble in phase E, M and W keep same

what if two/three of them happen same time?

ret:

M				M				M		ret	
E				E		ret		E		nop	
D		ret		D		nop		D		nop	
		1				2				3	

one of them

jxx error:

M			
E		jxx	
D			

when jxx error happens, E_icode must be jxx

data hazard:

```

                <--+
+-----+ |
M |         |<--+
+-----+ |
E |         |<--+
+-----+ |
D |    xxx   +--+
+-----+

```

when data hazard happens, D_icode is not sure, xxx means any instruction

composition 1: ret and jxx

keep same on book

composition 2: ret and data hazard

they two happens same time, so xxx must be ret

```

                <--+
+-----+ |
M |         |<--+
+-----+ |
E |         |<--+
+-----+ |
D |    ret   +--+
+-----+

```

when they happen same time, data hazard is prior to ret because if ret doesn't stall to avoid data hazard, we get wrong answer with ISA

composition 3: jxx error and data hazard

```

                <--+
+-----+ |
M |         |<--+
+-----+ |
E |    jxx   |<--+
+-----+ |
D |    xxx   +--+
+-----+

```

when they two happens same time, jxx error is prior to data hazard because next 2 instructions is canceled when jxx error, xxx is canceled anymore.

composition 4: ret & jxx error & data hazard

```

          <--+
+-----+ |
M |      | <--+
+-----+ |
E |   jxx  | <--+
+-----+ |
D |   ret   +---+
+-----+

```

same like composition 3.

finally:

num	data	ret	jxx	F	D	E	M	W
1	0	0	0	X	X	X	X	X
2	0	0	1	X	bubble	bubble	X	X
3	0	1	0	stall	bubble	X	X	X
4	1	0	0	stall	stall	bubble	X	X
5	0	1	1	X	bubble	bubble	X	X
6	1	0	1	X	bubble	bubble	X	X
7	1	1	0	stall	stall	bubble	X	X
8	1	1	1	X	bubble	bubble	X	X

F:

- stall: (data || ret) && !jxx
- bubble: 0

D:

- stall: data && !jxx
- bubble: jxx || (!data && ret)

E:

- stall: 0
- bubble: jxx || data

M:

keep same

W:

keep same

finally:

check file `./chapter4/code/sim/pipe/pipe-stall.hcl`

watch changes with origin pipe-stall.hcl file

```
(cd ../chapter4/code/sim/pipe; diff -u origin-pipe-stall.hcl pipe-stall.hcl)
```

```
--- origin-pipe-stall.hcl    2017-11-09 02:57:43.671935966 +0000
+++ pipe-stall.hcl          2017-11-09 02:57:43.671935966 +0000
@@ -303,40 +303,95 @@
];

##### Pipeline Register Control #####
+# situation: ret
+# bool s_ret = IRET in { D_icode, E_icode, M_icode };
+#
+# situation: jxx error
+# bool s_jxx_error = (E_icode == IJXX && !e_Cnd);
+#
+# situation: data_hazard
+# bool s_data_hazard =
+#   (
+#     (
+#       d_srcA != RNONE &&
+#       d_srcA in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+#     ) ||
+#     (
+#       d_srcB != RNONE &&
+#       d_srcB in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+#     )
+#   )

# Should I stall or inject a bubble into Pipeline Register F?
# At most one of these can be true.
-bool F_stall =
-  # Modify the following to stall the update of pipeline register F
-  0 ||
-  # Stalling at fetch while ret passes through pipeline
-  IRET in { D_icode, E_icode, M_icode };
+# bool F_stall = (s_ret || s_data_hazard) && !s_jxx_error;
+bool F_stall = (
+  (IRET in { D_icode, E_icode, M_icode }) ||
+  (
+    (
+      d_srcA != RNONE &&
+      d_srcA in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+    ) ||
+    (
+      d_srcB != RNONE &&
+      d_srcB in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+    )
+  )
+  ) &&
+  !(E_icode == IJXX && !e_Cnd);

bool F_bubble = 0;
```

```

# Should I stall or inject a bubble into Pipeline Register D?
# At most one of these can be true.
-bool D_stall =
-    # Modify the following to stall the instruction in decode
-    0;
+# bool D_stall = s_data_hazard && !s_jxx_error;
+# bool D_bubble = s_jxx_error || (!s_data_hazard && s_ret)
+bool D_stall = (
+    (
+        d_srcA != RNONE &&
+        d_srcA in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+    ) ||
+    (
+        d_srcB != RNONE &&
+        d_srcB in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+    )
+ ) &&
+ !(E_icode == IJXX && !e_Cnd);

bool D_bubble =
-    # Mispredicted branch
-    (E_icode == IJXX && !e_Cnd) ||
-    # Stalling at fetch while ret passes through pipeline
-    !(E_icode in { IMMOVQ, IPOPQ } && E_dstM in { d_srcA, d_srcB }) &&
-    # but not condition for a generate/use hazard
-    !0 &&
-    IRET in { D_icode, E_icode, M_icode };
+ (E_icode == IJXX && !e_Cnd) ||
+ (
+    !(
+        (
+            d_srcA != RNONE &&
+            d_srcA in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+        ) ||
+        (
+            d_srcB != RNONE &&
+            d_srcB in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }
+        )
+    ) &&
+    (IRET in { D_icode, E_icode, M_icode })
+ );

# Should I stall or inject a bubble into Pipeline Register E?
# At most one of these can be true.
+# bool E_stall = 0;
+# bool E_bubble = s_jxx_error || s_data_hazard
bool E_stall = 0;
bool E_bubble =
-    # Mispredicted branch
-    (E_icode == IJXX && !e_Cnd) ||
-    # Modify the following to inject bubble into the execute stage
-    0;

```



```
+ (E_icode == IJXX && !e_Cnd) ||  
+ (  
+   (  
+     d_srcA != RNONE &&  
+     d_srcA in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }  
+   ) ||  
+   (  
+     d_srcB != RNONE &&  
+     d_srcB in { e_dstE, E_dstM, M_dstM, M_dstE, W_dstM, W_dstE }  
+   )  
+ );  
+
```

Should I stall or inject a bubble into Pipeline Register M?

At most one of these can be true.

4.54

almost same like 4.51 seq-full.hcl

check file `./chapter4/code/sim/pipe/pipe-full.hcl` with iaddq functionality

watch changes with origin pipe-full.hcl file

```
(cd ./chapter4/code/sim/pipe; diff -u origin-pipe-full.hcl pipe-full.hcl)
```

```

--- origin-pipe-full.hcl      2017-11-09 02:57:43.671935966 +0000
+++ pipe-full.hcl            2017-11-09 02:57:43.671935966 +0000
@@ -158,7 +158,7 @@
# Is instruction valid?
bool instr_valid = f_icode in
    { INOP, IHALT, IRRMOVQ, IIRMOVQ, IRMMOVQ, IMRMVQ,
-    IOPQ, IJXX, ICALL, IRET, IPUSHQ, IPOPQ };
+    IOPQ, IJXX, ICALL, IRET, IPUSHQ, IPOPQ, IIADDQ };

# Determine status code for fetched instruction
word f_stat = [
@@ -171,11 +171,11 @@
# Does fetched instruction require a regid byte?
bool need_regids =
    f_icode in { IRRMOVQ, IOPQ, IPUSHQ, IPOPQ,
-    IIRMOVQ, IRMMOVQ, IMRMVQ };
+    IIRMOVQ, IRMMOVQ, IMRMVQ, IIADDQ };

# Does fetched instruction require a constant word?
bool need_valC =
-    f_icode in { IIRMOVQ, IRMMOVQ, IMRMVQ, IJXX, ICALL };
+    f_icode in { IIRMOVQ, IRMMOVQ, IMRMVQ, IJXX, ICALL, IIADDQ };

# Predict next value of PC
word f_predPC = [
@@ -195,14 +195,14 @@

## What register should be used as the B source?
word d_srcB = [
-    D_icode in { IOPQ, IRMMOVQ, IMRMVQ } : D_rB;
+    D_icode in { IOPQ, IRMMOVQ, IMRMVQ, IIADDQ } : D_rB;
    D_icode in { IPUSHQ, IPOPQ, ICALL, IRET } : RRSP;
    1 : RNONE; # Don't need register
];

## What register should be used as the E destination?
word d_dstE = [
-    D_icode in { IRRMOVQ, IIRMOVQ, IOPQ } : D_rB;
+    D_icode in { IRRMOVQ, IIRMOVQ, IOPQ, IIADDQ } : D_rB;
    D_icode in { IPUSHQ, IPOPQ, ICALL, IRET } : RRSP;
    1 : RNONE; # Don't write any register
];

```

```

@@ -239,7 +239,7 @@
## Select input A to ALU
word aluA = [
    E_icode in { IRRMOVQ, IOPQ } : E_valA;
-   E_icode in { IIRMOVQ, IRMMOVQ, IMRMVQ } : E_valC;
+   E_icode in { IIRMOVQ, IRMMOVQ, IMRMVQ, IIADDQ } : E_valC;
    E_icode in { ICALL, IPUSHQ } : -8;
    E_icode in { IRET, IPOPQ } : 8;
    # Other instructions don't need ALU
@@ -248,7 +248,7 @@
## Select input B to ALU
word aluB = [
    E_icode in { IRMMOVQ, IMRMVQ, IOPQ, ICALL,
-   IPUSHQ, IRET, IPOPQ } : E_valB;
+   IPUSHQ, IRET, IPOPQ, IIADDQ } : E_valB;
    E_icode in { IRRMOVQ, IIRMOVQ } : 0;
    # Other instructions don't need ALU
];
@@ -260,7 +260,7 @@
];

## Should the condition codes be updated?
-bool set_cc = E_icode == IOPQ &&
+bool set_cc = (E_icode == IOPQ || E_icode == IIADDQ) &&
    # State changes only during normal operation
    !m_stat in { SADR, SINS, SHLT } && !W_stat in { SADR, SINS, SHLT };

```

4.55

Problem mismatch with skeleton code, we follow the code here: **change J_YES to UNCOND**

one point is

```
M_icode == IJXX && M_ifun != UNCOND
```

means all jxx except jmp instruction

another important point is Mispredicted branch condition

origin

```
(E_icode == IJXX && !e_Cnd)
```

now

```
(E_icode == IJXX && E_ifun != UNCOND && e_Cnd)
```

if e_Cnd means misprediction

check file `./chapter4/code/sim/pipe/pipe-nt.hcl`

watch changes with origin pipe-nt.hcl

```
(cd ./chapter4/code/sim/pipe; diff -u origin-pipe-nt.hcl pipe-nt.hcl)
```

```

--- origin-pipe-nt.hcl    2017-11-09 02:57:43.671935966 +0000
+++ pipe-nt.hcl          2017-11-09 02:57:43.671935966 +0000
@@ -139,7 +139,7 @@
## What address should instruction be fetched at
word f_pc = [
    # Mispredicted branch.  Fetch at incremented PC
-   M_icode == IJXX && !M_Cnd : M_valA;
+   M_icode == IJXX && M_ifun != UNCOND && M_Cnd : M_valA;
    # Completion of RET instruction
    W_icode == IRET : W_valM;
    # Default: Use predicted value of PC
@@ -183,6 +183,7 @@
    # Predict next value of PC
word f_predPC = [
    # BNT: This is where you'll change the branch prediction rule
+   f_icode == IJXX && f_ifun != UNCOND : f_valP;
    f_icode in { IJXX, ICALL } : f_valC;
    1 : f_valP;
];
@@ -273,7 +274,11 @@
    !m_stat in { SADR, SINS, SHLT } && !W_stat in { SADR, SINS, SHLT };

## Generate valA in execute stage
-word e_valA = E_valA;    # Pass valA through stage
+## pass branch address back by M_valA
+word e_valA = [
+   E_icode == IJXX && E_ifun != UNCOND : E_valC;
+   1 : E_valA;    # Pass valA through stage
+];

## Set dstE to RNONE in event of not-taken conditional move
word e_dstE = [
@@ -343,7 +348,7 @@

bool D_bubble =
    # Mispredicted branch
-   (E_icode == IJXX && !e_Cnd) ||
+   (E_icode == IJXX && E_ifun != UNCOND && e_Cnd) ||
    # Stalling at fetch while ret passes through pipeline
    # but not condition for a load/use hazard
    !(E_icode in { IMRMVQ, IPOPQ } && E_dstM in { d_srcA, d_srcB }) &&
@@ -354,7 +359,7 @@
bool E_stall = 0;
bool E_bubble =
    # Mispredicted branch
-   (E_icode == IJXX && !e_Cnd) ||
+   (E_icode == IJXX && E_ifun != UNCOND && e_Cnd) ||
    # Conditions for a load/use hazard
    E_icode in { IMRMVQ, IPOPQ } &&
    E_dstM in { d_srcA, d_srcB};

```


4.56

Problem mismatch with skeleton code, we follow the code here: **change J_YES to UNCOND**

almost same like 4.55. main differences are

- pay attention to whether valC is greater or lower than valP
- pass both valC and valP back by M registers because if misprediction happens we need judge jumping back to valC or valP

check file `./chapter4/code/sim/pipe/pipe-btfnt.hcl`

only watch changes with origin pipe-btfnt.hcl

```
(cd ./chapter4/code/sim/pipe; diff -u origin-pipe-btfnt.hcl pipe-btfnt.hcl)
```

```
--- origin-pipe-btfnt.hcl    2017-11-09 02:57:43.667935995 +0000
+++ pipe-btfnt.hcl        2017-11-09 02:57:43.671935966 +0000
@@ -139,7 +139,10 @@
  ## What address should instruction be fetched at
  word f_pc = [
    # Mispredicted branch.  Fetch at incremented PC
-   M_icode == IJXX && !M_Cnd : M_valA;
+   # backward taken error
+   M_icode == IJXX && M_ifun != UNCOND && M_valE < M_valA && !M_Cnd : M_valA;
+   # forward not-taken error
+   M_icode == IJXX && M_ifun != UNCOND && M_valE >= M_valA && M_Cnd : M_valE;
    # Completion of RET instruction
    W_icode == IRET : W_valM;
    # Default: Use predicted value of PC
@@ -183,6 +186,8 @@
  # Predict next value of PC
  word f_predPC = [
    # BBTfnt: This is where you'll change the branch prediction rule
+   f_icode == IJXX && f_ifun != UNCOND && f_valC < f_valP : f_valC;
+   f_icode == IJXX && f_ifun != UNCOND && f_valC >= f_valP : f_valP;
    f_icode in { IJXX, ICALL } : f_valC;
    1 : f_valP;
  ];
@@ -244,12 +249,15 @@
  # way to get valC into pipeline register M, so that
  # you can correct for a mispredicted branch.

+## pass valC by M_valE, pass valP by M_valA
+
  ## Select input A to ALU
  word aluA = [
    E_icode in { IRRMOVQ, IOPQ } : E_valA;
    E_icode in { IIRMOVQ, IRMMOVQ, IMRMVQ } : E_valC;
    E_icode in { ICALL, IPUSHQ } : -8;
    E_icode in { IRET, IPOPQ } : 8;
+   E_icode in { IJXX } : E_valC;
    # Other instructions don't need ALU
```

```

];

@@ -258,6 +266,7 @@
    E_icode in { IRMMOVQ, IMRMVQ, IOPQ, ICALL,
                IPUSHQ, IRET, IPOPQ } : E_valB;
    E_icode in { IRRMOVQ, IIRMOVQ } : 0;
+   E_icode in { IJXX } : 0;
    # Other instructions don't need ALU
];

@@ -343,7 +352,11 @@

bool D_bubble =
    # Mispredicted branch
-   (E_icode == IJXX && !e_Cnd) ||
+   # backward taken error or forward not-taken error
+   (
+   (E_icode == IJXX && E_ifun != UNCOND && E_valC < E_valA && !e_Cnd) ||
+   (E_icode == IJXX && E_ifun != UNCOND && E_valC >= E_valA && e_Cnd)
+   ) ||
    # BBTFNT: This condition will change
    # Stalling at fetch while ret passes through pipeline
    # but not condition for a load/use hazard
@@ -355,7 +368,11 @@
bool E_stall = 0;
bool E_bubble =
    # Mispredicted branch
-   (E_icode == IJXX && !e_Cnd) ||
+   # backward taken error or forward not-taken error
+   (
+   (E_icode == IJXX && E_ifun != UNCOND && E_valC < E_valA && !e_Cnd) ||
+   (E_icode == IJXX && E_ifun != UNCOND && E_valC >= E_valA && e_Cnd)
+   ) ||
    # BBTFNT: This condition will change
    # Conditions for a load/use hazard
    E_icode in { IMRMVQ, IPOPQ } &&

```


4.57

A.

consider load/use hazard

```
E_icode in { IMRMOVQ, IPOPQ } && E_dstM in { d_srcA, d_srcB };
```

situation	1	2	3	4
E_dstM == d_srcA	1	1	0	0
E_dstM == d_srcB	1	0	1	0

situation 4:

normal, no hazard happens

situation 1,2,3:

load/use hazard happens

if E_icode in { IMRMOVQ, IPOPQ }, then E_dstM must not be RNONE

consider situation 1 and 3, E_dstM == d_srcB, then d_srcB is not RNONE and must be used in phase E, so load-forward can't work. load-forward only work in situation 2!

consider all the instructions that d_srcA is not RNONE

instructions	d_srcA	valA used in phase E?	load-forward works?
rrmovq	rA	Y	N
rmmovq	rA	N	Y
opq	rA	Y	N
pushq	rA	N	Y
ret	rsp	N	N
popq	rsp	N	N

ret and popq can't work because d_srcA == d_srcB == %rsp, that's not in situation 2!!

finally, load/use only work in condition:

```
E_icode in { IMRMOVQ, IPOPQ } &&
(
  E_dstM == d_srcB ||
  (
    E_dstM == d_srcA && !(D_icode in { IRMMOVQ, IPUSHQ })
  )
);
```

B.

check file `./chapter4/code/sim/pipe/pipe-lf.hcl`

only watch changes with origin pipe-lf.hcl

```
(cd ./chapter4/code/sim/pipe; diff -u origin-pipe-lf.hcl pipe-lf.hcl)
```

```
--- origin-pipe-lf.hcl      2017-11-09 02:57:43.671935966 +0000
+++ pipe-lf.hcl            2017-11-09 02:57:43.671935966 +0000
@@ -271,6 +271,7 @@
##   from memory stage when appropriate
## Here it is set to the default used in the normal pipeline
word e_valA = [
+   E_icode in { IRMMOVQ, IPUSHQ } && E_srcA == M_dstM : m_valM;
+   1 : E_valA; # Use valA from stage pipe register
];

@@ -329,7 +330,13 @@
bool F_stall =
    # Conditions for a load/use hazard
    ## Set this to the new load/use condition
-   0 ||
+   E_icode in { IMRMOVQ, IPOPQ } &&
+   (
+   E_dstM == d_srcB ||
+   (
+   E_dstM == d_srcA && !(D_icode in { IRMMOVQ, IPUSHQ })
+   )
+   ) ||
    # Stalling at fetch while ret passes through pipeline
    IRET in { D_icode, E_icode, M_icode };

@@ -338,15 +345,29 @@
bool D_stall =
    # Conditions for a load/use hazard
    ## Set this to the new load/use condition
-   0;
+   E_icode in { IMRMOVQ, IPOPQ } &&
+   (
+   E_dstM == d_srcB ||
+   (
+   E_dstM == d_srcA && !(D_icode in { IRMMOVQ, IPUSHQ })
+   )
+   );

bool D_bubble =
    # Mispredicted branch
    (E_icode == IJXX && !e_Cnd) ||
    # Stalling at fetch while ret passes through pipeline
    # but not condition for a load/use hazard
-   !(E_icode in { IMRMOVQ, IPOPQ } && E_dstM in { d_srcA, d_srcB }) &&
-   IRET in { D_icode, E_icode, M_icode };
+   !(
+   E_icode in { IMRMOVQ, IPOPQ } &&
+   (
```

```

+      E_dstM == d_srcB ||
+      (
+          E_dstM == d_srcA && !(D_icode in { IRMMOVQ, IPUSHQ })
+      )
+  )
+ ) &&
+ IRET in { D_icode, E_icode, M_icode };

```

Should I stall or inject a bubble into Pipeline Register E?

At most one of these can be true.

```
@@ -356,7 +377,13 @@
```

```

      (E_icode == IJXX && !e_Cnd) ||
      # Conditions for a load/use hazard
      ## Set this to the new load/use condition
-      0;
+      E_icode in { IMRMOVQ, IPOPQ } &&
+      (
+          E_dstM == d_srcB ||
+          (
+              E_dstM == d_srcA && !(D_icode in { IRMMOVQ, IPUSHQ })
+          )
+      );

```

Should I stall or inject a bubble into Pipeline Register M?

At most one of these can be true.

4.58

it's not a hard problem, focus on key points:

1. fetch popq instruction twice.
2. first time fetch popq, works like iaddq; second time fetch popq2, works like mrmovq
3. load/use condition should be popq2 not popq.

phase	popq rA	popq2 rA
works like	iadd \$8, %rsp	mrmovq -8(%rsp), rA
F	valP = PC	valP = PC + 2
D	valB=R[rsp]	valB=R[rsp]
E	valE=valB+8	valE=valB-8
M		valM=M8[valE]
W	R[rsp]=valE	R[rA]=valM

check file `./chapter4/code/sim/pipe/pipe-1w.hcl`

only watch changes with origin pipe-1w.hcl

```
(cd ./chapter4/code/sim/pipe; diff -u origin-pipe-1w.hcl pipe-1w.hcl)
```

```
--- origin-pipe-1w.hcl    2017-11-09 02:57:43.667935995 +0000
+++ pipe-1w.hcl          2017-11-09 02:57:43.671935966 +0000
@@ -157,6 +157,7 @@
## so that it will be IPOP2 when fetched for second time.
word f_icode = [
    imem_error : INOP;
+   D_icode == IPOPQ : IPOP2;
    1: imem_icode;
];

@@ -169,7 +170,7 @@
# Is instruction valid?
bool instr_valid = f_icode in
    { INOP, IHALT, IRRMOVQ, IIRMOVQ, IRMMOVQ, IMRMOVQ,
-     IOPQ, IJXX, ICALL, IRET, IPUSHQ, IPOPQ };
+     IOPQ, IJXX, ICALL, IRET, IPUSHQ, IPOPQ, IPOP2 };

# Determine status code for fetched instruction
word f_stat = [
@@ -182,7 +183,7 @@
# Does fetched instruction require a regid byte?
bool need_regids =
    f_icode in { IRRMOVQ, IOPQ, IPUSHQ, IPOPQ,
-               IIRMOVQ, IRMMOVQ, IMRMOVQ };
+               IIRMOVQ, IRMMOVQ, IMRMOVQ, IPOP2 };

# Does fetched instruction require a constant word?
```

```

bool need_valC =
@@ -192,6 +193,7 @@
word f_predPC = [
    f_icode in { IJXX, ICALL } : f_valC;
    ## 1W: Want to refetch popq one time
+   f_icode == IPOPQ : f_pc;
    1 : f_valP;
];

@@ -204,14 +206,14 @@
## What register should be used as the A source?
word d_srcA = [
    D_icode in { IRRMOVQ, IRMMOVQ, IOPQ, IPUSHQ } : D_rA;
-   D_icode in { IPOPQ, IRET } : RRSP;
+   D_icode in { IRET } : RRSP;
    1 : RNONE; # Don't need register
];

## What register should be used as the B source?
word d_srcB = [
    D_icode in { IOPQ, IRMMOVQ, IMRMVQ } : D_rB;
-   D_icode in { IPUSHQ, IPOPQ, ICALL, IRET } : RRSP;
+   D_icode in { IPUSHQ, IPOPQ, IPOP2, ICALL, IRET } : RRSP;
    1 : RNONE; # Don't need register
];

@@ -224,7 +226,7 @@

## What register should be used as the M destination?
word d_dstM = [
-   D_icode in { IMRMVQ, IPOPQ } : D_rA;
+   D_icode in { IMRMVQ, IPOP2 } : D_rA;
    1 : RNONE; # Don't write any register
];

@@ -255,7 +257,7 @@
word aluA = [
    E_icode in { IRRMOVQ, IOPQ } : E_valA;
    E_icode in { IIRMOVQ, IRMMOVQ, IMRMVQ } : E_valC;
-   E_icode in { ICALL, IPUSHQ } : -8;
+   E_icode in { ICALL, IPUSHQ, IPOP2 } : -8;
    E_icode in { IRET, IPOPQ } : 8;
    # Other instructions don't need ALU
];

@@ -263,7 +265,7 @@
## Select input B to ALU
word aluB = [
    E_icode in { IRMMOVQ, IMRMVQ, IOPQ, ICALL,
-   IPUSHQ, IRET, IPOPQ } : E_valB;
+   IPUSHQ, IRET, IPOPQ, IPOP2 } : E_valB;
    E_icode in { IRRMOVQ, IIRMOVQ } : 0;
    # Other instructions don't need ALU
];

```

```

@@ -292,13 +294,13 @@

## Select memory address
word mem_addr = [
-   M_icode in { IRMMOVQ, IPUSHQ, ICALL, IMRMVQ } : M_valE;
-   M_icode in { IPOPQ, IRET } : M_valA;
+   M_icode in { IRMMOVQ, IPUSHQ, ICALL, IMRMVQ, IPOP2 } : M_valE;
+   M_icode in { IRET } : M_valA;
    # Other instructions don't need address
];

## Set read control signal
-bool mem_read = M_icode in { IMRMVQ, IPOPQ, IRET };
+bool mem_read = M_icode in { IMRMVQ, IRET, IPOP2 };

## Set write control signal
bool mem_write = M_icode in { IRMMOVQ, IPUSHQ, ICALL };
@@ -350,7 +352,7 @@
bool F_bubble = 0;
bool F_stall =
    # Conditions for a load/use hazard
-   E_icode in { IMRMVQ, IPOPQ } &&
+   E_icode in { IMRMVQ, IPOP2 } &&
    E_dstM in { d_srcA, d_srcB } ||
    # Stalling at fetch while ret passes through pipeline
    IRET in { D_icode, E_icode, M_icode };
@@ -359,7 +361,7 @@
# At most one of these can be true.
bool D_stall =
    # Conditions for a load/use hazard
-   E_icode in { IMRMVQ, IPOPQ } &&
+   E_icode in { IMRMVQ, IPOP2 } &&
    E_dstM in { d_srcA, d_srcB };

bool D_bubble =
@@ -367,7 +369,7 @@
    (E_icode == IJXX && !e_Cnd) ||
    # Stalling at fetch while ret passes through pipeline
    # but not condition for a load/use hazard
-   !(E_icode in { IMRMVQ, IPOPQ } && E_dstM in { d_srcA, d_srcB }) &&
+   !(E_icode in { IMRMVQ, IPOP2 } && E_dstM in { d_srcA, d_srcB }) &&
    # 1W: This condition will change
    IRET in { D_icode, E_icode, M_icode };

@@ -378,7 +380,7 @@
# Mispredicted branch
(E_icode == IJXX && !e_Cnd) ||
# Conditions for a load/use hazard
-   E_icode in { IMRMVQ, IPOPQ } &&
+   E_icode in { IMRMVQ, IPOP2 } &&
    E_dstM in { d_srcA, d_srcB };

# Should I stall or inject a bubble into Pipeline Register M?

```


4.59

4.47 is the better one

loop part in 4.47

```
L4:
    mrmovq 8(%rax), %r9
    mrmovq (%rax), %r10
    rrmovq %r9, %r8
    subq %r10, %r8
    jge L3
    rmmovq %r10, 8(%rax)
    rmmovq %r9, (%rax)
```

50% jge is right, run 5 instructions; 50% jge is wrong, run 7 instructions and 2 nop bubble. so Cycles Per Loop is $50\% 5 + (7 + 2) 50\% = 7$

loop part in 4.48

```
L4:
    mrmovq 8(%rax), %r9
    mrmovq (%rax), %r10
    rrmovq %r9, %r8
    subq %r10, %r8
    cmovl %r9, %r11
    cmovl %r10, %r9
    cmovl %r11, %r10
    rmmovq %r9, 8(%rax)
    rmmovq %r10, (%rax)
```

Cycles Per Loop is 9

loop part in 4.49

```
L4:
    mrmovq 8(%rax), %r9
    mrmovq (%rax), %r10
    rrmovq %r9, %r8
    rrmovq %r10, %r11
    xorq %r9, %r10
    subq %r11, %r8
    cmovge %r11, %r9
    xorq %r10, %r9
    xorq %r9, %r10
    rmmovq %r9, 8(%rax)
    rmmovq %r10, (%rax)
```

Cycles Per Loop is 11

Optimizing Program Performance

We are all tasked to balance and optimize ourselves.

by Mae Jemison

5.1 - 5.12 visit book

5.13 - 5.19 visit here

test

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right

solution	code file	test way
5.13	5.13.c	assert
5.14	5.14.c	assert
5.15	5.15.c	assert
5.16	5.16.c	assert
5.17	5.17.c	assert
5.18	5.18.c	assert
5.19	5.19.c	assert

prof

prerequisite

- google [gperftools](#)

to 5.18 & 5.19, you can measure its performance.

```
(cd chapter5/code; make 5.18.prof)
(cdd chapter5/code; make 5.18.prof)
```

5.13

A.

```

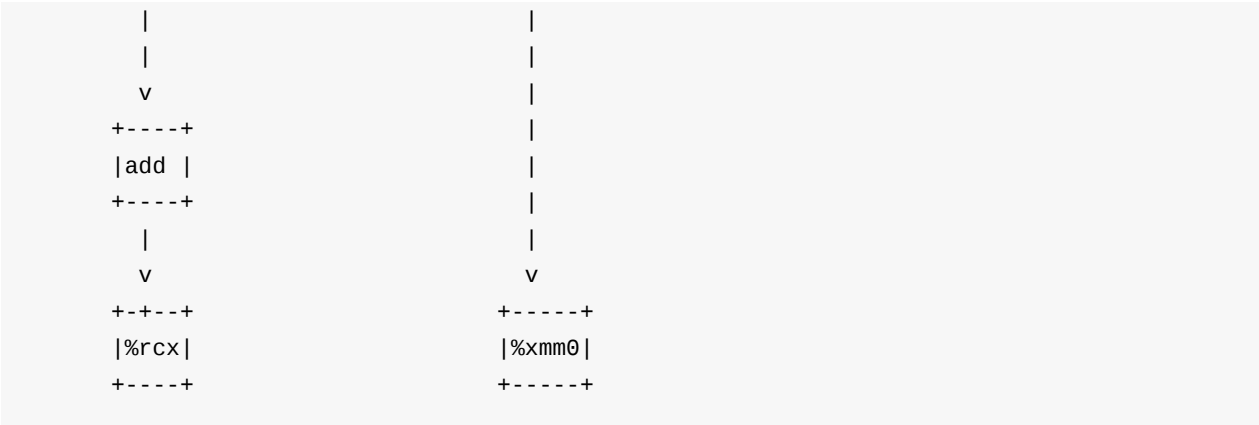
+-----+-----+-----+-----+-----+-----+
| %rbp | %rcx | %rax | %rbx | %xmm1 | %xmm0 |
+-----+-----+-----+-----+-----+-----+
| +----|----|----|-----|--->|      |
|      +----|----|-----|--->|load|      | vmovad 0(%rbp,%rcx,8),%xmm1
|      |    |    |      +----|----|      |
|      |    |    |    |      +----+
|      +----|----|-----|--->|      |
|      |    |    |    |      |load|----+
|      |    +----|-----|--->|      |      | vmulsd (%rax,%rcx,8),%xmm1,%xmm0
|      |    |    |    |      +----+      |
|      |    |    |    |      |<---+
|      |    |    |      +----|--->|mul|      |
|      |    |    |      +----|----|      |
|      |    |    |    |      +----+
|      |    |    |      +----|--->|      |
|      |    |    |    |      +--->|add|      | vaddsd %xmm1,%xmm0,%xmm0
|      |    |    |    |      +----|      |
|      |    |    |    |      +----+
|      +----|----|-----|--->|      |
|      |    |    |    |      |add|      | addq $1, %rcx
|      +----|----|-----|----|      |
|      |    |    |    |      +----+
|      +----|----|-----|--->|      |
|      |    |    |    |      |cmp|----+ cmpq %rbx, %rcx
|      |    |      +----|-----|--->|      | |
|      |    |    |    |      +----+      |
|      |    |    |    |      |    |      |
|      |    |    |    |      |jne|<---+ jne .L15
|      |    |    |    |      |    |
|      |    |    |    |      +----+
v      v      v      v      v      v
+-----+-----+-----+-----+-----+
| %rbp | %rcx | %rax | %rbx | %xmm1 | %xmm0 |
+-----+-----+-----+-----+-----+

```

```

+-----+          +-----+
| %rcx |          | %xmm0 |
+-----+          +-----+
|          |          |
|      +----+          | <----- key path
+----->|load|-----+
|      +----+          |
|          v          v
|      +----+ +----+ +----+
+----->|load|--->|mul|--->|add|
|      +----+ +----+ +----+

```



B.

5-12

float add cell, CPE is 3.0

C.

5-12

long add cell, 1.0

D.

only float add on key path

```
/*
 * 5.13.c
 */
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include "../lib/vec.h"

/* inner product. accumulate in temporary */
void inner4(vec_ptr u, vec_ptr v, data_t *dest) {
    long i;
    long length = vec_length(u);
    data_t *udata = get_vec_start(u);
    data_t *vdata = get_vec_start(v);
    data_t sum = (data_t) 0;

    for (i = 0; i < length; i++) {
        sum = sum + udata[i] * vdata[i];
    }
    *dest = sum;
}

int main(int argc, char* argv[]) {
    vec_ptr u = new_vec(4);
    vec_ptr v = new_vec(4);

    data_t *arr = (data_t*) malloc(sizeof(data_t) * 4);
    arr[0] = 0;
    arr[1] = 1;
    arr[2] = 2;
    arr[3] = 3;

    set_vec_start(u, arr);
    set_vec_start(v, arr);

    data_t res;
    inner4(u, v, &res);

    assert(res == 1+4+9);
    return 0;
}
```

5.14

```

+-----+               +-----+
| sum |               | i |
+-----+               +-----+
|               |
v               |
+-----+   +-----+ +-----+   |
|add |<-----|mul |<-|load|<-----+
+-----+   +-----+ +-----+   |
|               |
v               |
+-----+   +-----+ +-----+   |
|add |<-----|mul |<-|load|<-----+
+-----+   +-----+ +-----+   |
|               |
v               |
+-----+   +-----+ +-----+   |
|add |<-----|mul |<-|load|<-----+
+-----+   +-----+ +-----+   |
|               |
v               |
+-----+   +-----+ +-----+   |
|add |<-----|mul |<-|load|<-----+
+-----+   +-----+ +-----+   |
|               |
v               |
+-----+   +-----+ +-----+   |
|add |<-----|mul |<-|load|<-----+
+-----+   +-----+ +-----+   |
|               |
|               |
|               |
| <----- key path               |
|               |
|               |
v               v
+-----+       +-----+
| sum |       | i |
+-----+       +-----+

```

A.

every element has 6 long/float add

element count is $n/6$

so $n/6 * 6 = n$

CPE bound == 1.0

B.

same like A

```

/*
 * 5.14.c
 */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <assert.h>
#include "../lib/vec.h"

#define LEN 24

/* inner product. accumulate in temporary */
void inner4(vec_ptr u, vec_ptr v, data_t *dest) {
    long i;
    long length = vec_length(u);
    data_t *udata = get_vec_start(u);
    data_t *vdata = get_vec_start(v);
    data_t sum = (data_t) 0;

    for (i = 0; i < length-6; i+=6) {
        sum = sum + udata[i] * vdata[i] +
            udata[i+1] * vdata[i+1] +
            udata[i+2] * vdata[i+2] +
            udata[i+3] * vdata[i+3] +
            udata[i+4] * vdata[i+4] +
            udata[i+5] * vdata[i+5];
    }
    for(; i < length; i++) {
        sum = sum + udata[i] * vdata[i];
    }
    *dest = sum;
}

int main(int argc, char* argv[]) {
    vec_ptr u = new_vec(LEN);
    vec_ptr v = new_vec(LEN);

    data_t *arr = (data_t*) malloc(sizeof(data_t) * LEN);
    memset(arr, 0, sizeof(data_t) * LEN);
    arr[0] = 0;
    arr[1] = 1;
    arr[2] = 2;
    arr[3] = 3;

    set_vec_start(u, arr);

```

```
set_vec_start(v, arr);  
  
data_t res;  
inner4(u, v, &res);  
  
assert(res == 1+4+9);  
return 0;  
}
```


5.15

maybe

1. float mul capacity limit(less than 6)
2. register renaming limit

```

/*
 * 5.15.c
 */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <assert.h>
#include "../lib/vec.h"

#define LEN 24

/* inner product. accumulate in temporary */
void inner4(vec_ptr u, vec_ptr v, data_t *dest) {
    long i;
    long length = vec_length(u);
    data_t *udata = get_vec_start(u);
    data_t *vdata = get_vec_start(v);
    data_t sum = (data_t) 0;
    data_t sum1 = (data_t) 0;
    data_t sum2 = (data_t) 0;
    data_t sum3 = (data_t) 0;
    data_t sum4 = (data_t) 0;
    data_t sum5 = (data_t) 0;

    for (i = 0; i < length-6; i+=6) {
        sum = sum + udata[i] * vdata[i];
        sum1 = sum1 + udata[i+1] * vdata[i+1];
        sum2 = sum2 + udata[i+2] * vdata[i+2];
        sum3 = sum3 + udata[i+3] * vdata[i+3];
        sum4 = sum4 + udata[i+4] * vdata[i+4];
        sum5 = sum5 + udata[i+5] * vdata[i+5];
    }
    for(; i < length; i++) {
        sum = sum + udata[i] * vdata[i];
    }
    *dest = sum + sum1 + sum2 + sum3 + sum4 + sum5;
}

int main(int argc, char* argv[]) {
    vec_ptr u = new_vec(LEN);
    vec_ptr v = new_vec(LEN);

    data_t *arr = (data_t*) malloc(sizeof(data_t) * LEN);
    memset(arr, 0, sizeof(data_t) * LEN);
    arr[0] = 0;
    arr[11] = 1;
}

```

```
arr[2] = 2;
arr[23] = 3;

set_vec_start(u, arr);
set_vec_start(v, arr);

data_t res;
inner4(u, v, &res);

assert(res == 1+4+9);
return 0;
}
```

5.16

```
/*
 * 5.16.c
 */
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <assert.h>
#include "../lib/vec.h"

#define LEN 24

/* inner product. accumulate in temporary */
void inner4(vec_ptr u, vec_ptr v, data_t *dest) {
    long i;
    long length = vec_length(u);
    data_t *udata = get_vec_start(u);
    data_t *vdata = get_vec_start(v);
    data_t sum = (data_t) 0;

    for (i = 0; i < length-6; i+=6) {
        sum = sum +
            (
                udata[i] * vdata[i] +
                udata[i+1] * vdata[i+1] +
                udata[i+2] * vdata[i+2] +
                udata[i+3] * vdata[i+3] +
                udata[i+4] * vdata[i+4] +
                udata[i+5] * vdata[i+5]
            );
    }
    for(; i < length; i++) {
        sum = sum + udata[i] * vdata[i];
    }
    *dest = sum;
}

int main(int argc, char* argv[]) {
    vec_ptr u = new_vec(LEN);
    vec_ptr v = new_vec(LEN);

    data_t *arr = (data_t*) malloc(sizeof(data_t) * LEN);
    memset(arr, 0, sizeof(data_t) * LEN);
    arr[0] = 0;
    arr[11] = 1;
    arr[2] = 2;
    arr[23] = 3;

    set_vec_start(u, arr);
    set_vec_start(v, arr);

    data_t res;
```

```
    inner4(u, v, &res);  
  
    assert(res == 1+4+9);  
    return 0;  
}
```

```
/*
 * 5.17.c
 */
#include <stdio.h>
#include <stdint.h>
#include <stdlib.h>
#include <string.h>
#include <assert.h>

void* basic_memset(void *s, int c, size_t n) {
    size_t cnt = 0;
    unsigned char *schar = s;
    while (cnt < n) {
        *schar++ = (unsigned char) c;
        cnt++;
    }
    return s;
}

/*
 * K = sizeof(unsigned long)
 * cs store K chars for memset
 */
void* effective_memset(void *s, unsigned long cs, size_t n) {
    /* align to K */
    size_t K = sizeof(unsigned long);
    size_t cnt = 0;
    unsigned char *schar = s;
    while (cnt < n) {
        if ((size_t)schar % K == 0) {
            break;
        }
        *schar++ = (unsigned char)cs;
        cnt++;
    }

    /* set K chars one time */
    unsigned long *slong = (unsigned long *)schar;
    size_t rest = n - cnt;
    size_t loop = rest / K;
    size_t tail = rest % K;

    for (size_t i = 0; i < loop; i++) {
        *slong++ = cs;
    }

    /* pad the tail part */
    schar = (unsigned char *)slong;
    for (size_t i = 0; i < tail; i++) {
        *schar++ = (unsigned char)cs;
    }
}
```

```
    return s;
}

int main(int argc, char* argv[]) {
    size_t space = sizeof(char) * 65537;
    // careful! malloc SIZE_MAX size memory will make sys slow
    // or crash down
    // size_t space = SIZE_MAX;

    void* basic_space = malloc(space);
    void* effective_space = malloc(space);

    int basic_fill = 0xFF;
    unsigned long effective_fill = ~0;

    basic_memset(basic_space, basic_fill, space);
    effective_memset(effective_space, effective_fill, space);

    assert(memcmp(basic_space, effective_space, space) == 0);

    free(basic_space);
    free(effective_space);
    return 0;
}
```

5.18

try to write 6*3a version for function holy.

this solution is not fully tested on my machine, sorry for that :(

```
uname -p
```

```
Intel(R) Core(TM) i5-3210M CPU @ 2.50GHz
```

if your machine core is Intel Core i7 Haswell like book, you can try install google [gperftools](#) for profile

```
(cd chapter5/code; make 5.18.prof)
```

loop more times is a good idea :)

```
/*
 * 5.18.c
 */
#include <stdio.h>
#include <assert.h>

/* calculate a0 + a1*x + a2*x^2 + ... + an*x^n */
double poly(double a[], double x, long degree) {
    long i;
    double result = a[0];
    double xpwr = x;
    for (i = 1; i <= degree; i++) {
        result += a[i] * xpwr;
        xpwr = x * xpwr;
    }
    return result;
}

/* version 6*3a */
double poly_6_3a(double a[], double x, long degree) {
    long i = 1;
    double result = a[0];
    double result1 = 0;
    double result2 = 0;

    double xpwr = x;
    double xpwr1 = x * x * x;
    double xpwr2 = x * x * x * x * x * x;

    double xpwr_step = x * x * x * x * x * x * x;
    for (; i <= degree - 6; i+=6) {
        result = result + (a[i]*xpwr + a[i+1]*xpwr*x);
        result1 = result1 + (a[i+2]*xpwr1 + a[i+3]*xpwr1*x);
        result2 = result2 + (a[i+4]*xpwr2 + a[i+5]*xpwr2*x);

        xpwr *= xpwr_step;
    }
}
```

```

    xpw1 *= xpw_step;
    xpw2 *= xpw_step;
}

for (; i <= degree; i++) {
    result = result + a[i]*xpw;
    xpw *= x;
}

return result + result1 + result2;
}

/* apply horner's method */
double polyh(double a[], double x, long degree) {
    long i;
    double result = a[degree];
    for (i = degree-1; i >= 0; i--) {
        result = a[i] + x*result;
    }
    return result;
}

#define LOOP 1000
#define LEN 1000

int main(int argc, char* argv[]) {
    double a[10 + 1] = { 0, 1, 1, 1, 1, 1,
                        1, 1, 1, 1, 1};

    double x = 2;
    long degree = 10;

    assert(poly(a, x, degree) == 2+4+8+16+32+(2+4+8+16+32)*32);
    assert(poly_6_3a(a, x, degree) == 2+4+8+16+32+(2+4+8+16+32)*32);
    assert(poly(a, x, degree) == polyh(a, x, degree));

    x = 1;
    degree = LEN;
    double b[LEN + 1];

    for (int c = 0; c < LOOP; c++) {
        poly(b, x, degree);
    }
    for (int c = 0; c < LOOP; c++) {
        poly_6_3a(b, x, degree);
    }
    for (int c = 0; c < LOOP; c++) {
        polyh(b, x, degree);
    }
    return 0;
}

```


5.19

this solution's performance is not fully measured.

if you're interested, you can try install google's [gperftools](#) for profile

```
(cd chapter5/code; make 5.19.prof)
```

loop more times is a good idea :)

```
/*
 * 5.19.c
 */
#include <stdio.h>
#include <assert.h>

void psum1a(float a[], float p[], long n) {
    long i;
    float last_val, val;
    last_val = p[0] = a[0];
    for (i = 1; i < n; i++) {
        val = last_val + a[i];
        p[i] = val;
        last_val = val;
    }
}

/* version 4*1a */
void psum_4_1a(float a[], float p[], long n) {
    long i;
    float val, last_val;
    float tmp, tmp1, tmp2, tmp3;
    last_val = p[0] = a[0];

    for (i = 1; i < n - 4; i++) {
        tmp = last_val + a[i];
        tmp1 = tmp + a[i+1];
        tmp2 = tmp1 + a[i+2];
        tmp3 = tmp2 + a[i+3];

        p[i] = tmp;
        p[i+1] = tmp1;
        p[i+2] = tmp2;
        p[i+3] = tmp3;

        /* key point */
        last_val = last_val + (a[i] + a[i+1] + a[i+2] + a[i+3]);
    }

    for (; i < n; i++) {
        last_val += a[i];
        p[i] = last_val;
    }
}
```

```
    }  
}  
  
#define LOOP 1000  
#define LEN 1000  
  
int main(int argc, char* argv[]) {  
    float a[5] = { 1, 2, 3, 4, 5 };  
    float p[5];  
    psum1a(a, p, 5);  
    assert(p[4] == 15);  
  
    float q[5];  
    psum_4_1a(a, q, 5);  
    assert(q[4] == 15);  
  
    /* for prof */  
    for (int i = 0; i < LOOP; i++) {  
        float s[LEN];  
        float d[LEN];  
        psum1a(s, d, LEN);  
        psum_4_1a(s, d, LEN);  
    }  
    return 0;  
}
```

The Memory Hierarchy

The more storage you have, the more stuff you accumulate.

by Alexis Stewart

6.1 - 6.21 visit book

6.22 - 6.46 visit here

test

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right

solution	code file	test way
6.22	----	----
6.23	----	----
6.24	----	----
6.25	----	----
6.26	----	----
6.27	----	----
6.28	----	----
6.29	----	----
6.30	----	----
6.31	----	----
6.32	----	----
6.33	----	----
6.34	----	----
6.35	----	----
6.36	----	----
6.37	----	----
6.38	----	----
6.39	----	----
6.40	----	----
6.41	----	----
6.42	----	----
6.43	----	----
6.44	----	----
6.45	transpose.c	no test, mainly measure performance
6.46	convert.c	no test, mainly measure performance

6.22

assume

Bits Per Track

$$bpt = x * r * K$$

Track Count

$$tc = (1-x) * r * M$$

M, K are constant

so

Bit Count

$$bc = K * M * r^2 * (1-x) * x$$

when $x = 1/2$, bc is maximum

6.23

$$T_{\text{avg_seek}} = 4\text{ms}$$

$$T_{\text{avg_rotation}} = 1/2 \cdot 1/15000 \cdot 60\text{s/min} \cdot 1000\text{ms/s} = 2\text{ms}$$

$$T_{\text{avg_transfer}} = 1/15000 \cdot 1/800 \cdot 60\text{s/min} \cdot 1000\text{ms/s} = 0.005\text{ms}$$

so

$$T_{\text{access}} = 6.005\text{ms}$$

6.24

almost same like problem 6.4 on book

A.

best case: blocks are mapped sequential and on same cylinder. just seek data once.

```
T_avg_seek = 4ms  
T_avg_rotation = 2ms
```

file size 2MB, block size 512B, block count $2\text{MB}/512\text{B} = 4000$

Block Per Track = 1000, so we need rotate 4 loop to read all data

```
T_transfer = T_rotation = T_max_rotation * 4 = 16ms
```

so

```
T_access = 22ms
```

B.

worse case: blocks are random.

```
T_access = 4000 * (T_avg_seek + T_avg_rotation) = 24s
```


6.25

m	c	B	E	S	t	s	b
32	1024	4	4	64	24	6	2
32	1024	4	256	1	30	0	2
32	1024	8	1	128	22	7	3
32	1024	8	128	1	29	0	3
32	1024	32	1	32	22	5	5
32	1024	32	4	8	24	3	5

6.26

m	c	B	E	S	t	s	b
32	2048	8	1	256	21	8	3
32	2048	4	4	128	23	7	2
32	1024	2	8	64	25	6	1
32	1024	32	2	16	23	4	5

6.27

A.

 $t = 0x45 = 0b01000101$, $s = 0b001$, $b = xx$ (xx is 00/01/10/11)

address may be

```
01000101 001 xx
```

```
format
```

```
0 1000 1010 01xx
```

address range: 0x08A4 - 0x08A7

 $t = 0x38$

address range: 0x0704 - 0x0707

B.

0x1238 - 0x123B

6.28

same like 6.27

A.

None

B.

0x18F0 - 0x18F3

0x00B0 - 0x00B3

C.

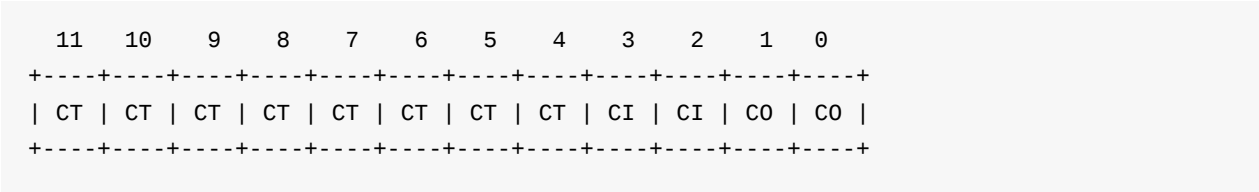
0x0E34 - 0x0E37

D.

0x1BDC - 0x1BDF

6.29

A.



B.

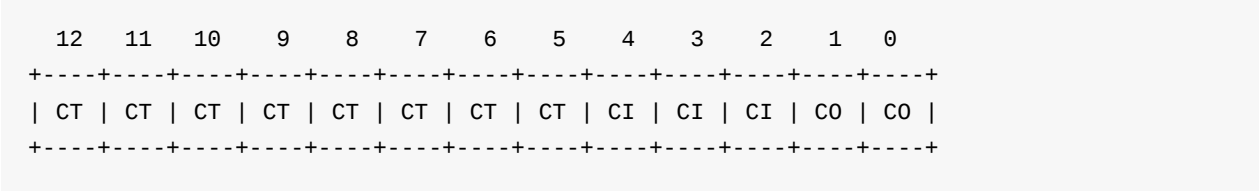
read/write	addr	hit?	value(or unknown)
read	0x834	No	-
write	0x836	Yes	unknown
read	0xFFD	Yes	0xC0

6.30

A.

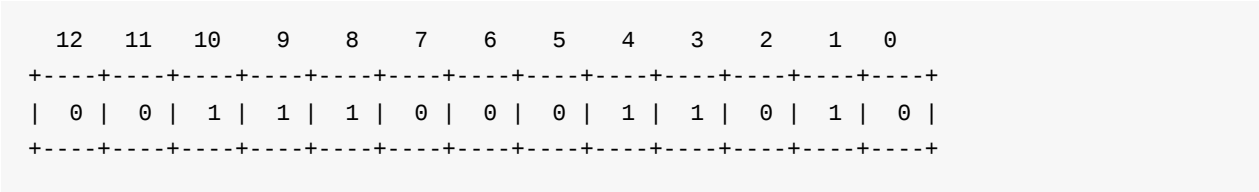
$C = E \ B \ S = 128B$

B.



6.31

A.



B.

param	value
CO	0x02
CI	0x06
CT	0x38
hit?	Yes
return	0xEB

6.32

A.

12	11	10	9	8	7	6	5	4	3	2	1	0
+	+	+	+	+	+	+	+	+	+	+	+	+
	1		0		1		1		0		1	
+	+	+	+	+	+	+	+	+	+	+	+	+

B.

param	value
CO	0x00
CI	0x02
CT	0xB7
hit?	No
return	--

6.33

same like 6.27

0x1788 - 0x178B

0x16C8 - 0x16CB

6.34

src:

	c0	c1	c2	c3
r0	m	m	h	m
r1	m	h	m	h
r2	m	m	h	m
r3	m	h	m	h

dst:

	c0	c1	c2	c3
r0	m	m	m	m
r1	m	m	m	m
r2	m	m	m	m
r3	m	m	m	m

6.35

src:

	c0	c1	c2	c3
r0	m	h	h	h
r1	m	h	h	h
r2	m	h	h	h
r3	m	h	h	h

dst:

	c0	c1	c2	c3
r0	m	h	h	h
r1	m	h	h	h
r2	m	h	h	h
r3	m	h	h	h

6.36

```

int x[2][128];
int i;
int sum = 0;

for (i = 0; i < 128; i++) {
    sum += x[0][i] * x[1][i];
}

```

A.

C = 512, E = 1, B = 16, S = 32

total read count: $2 * 128$ x[0][i] address: $i * 4$ x[1][i] address: $(128+i)4 = 512 + i4$

so x[0][i] and x[1][i] are cached into same block.

miss rate 100%

B.

C = 1024, E = 1, B = 16, S = 64

 $\text{sizeof}(x) = 2 * 128 * 4 = C$

whole array can be cached.

miss rate is depended on block size B.

B = 16, $\text{sizeof}(\text{int}) = 4$, so

miss rate is 25%

C.

C = 512, E = 2, B = 16, S = 16

total read count: $2 * 128$ x[0][i] address: $i * 4$ x[1][i] address: $(128+i)4 = 512 + i4$

so x[0][i] and x[1][i] are cached into different block in same set.

in first half, all elements can be cached. miss rate is 25%.

```

for (i = 0; i < 64; i++)
    sum += x[0][i] * x[1][i];

```

in second half

```
for (i = 64; i < 128; i++)  
    sum += x[0][i] * x[1][i];
```

$x[0][i]$ is not in cache. according to LRU strategy, cache $x[0][i]$ into the same block with $x[0][i-64]$, cache $x[1][i]$ into the same block with $x[1][i-64]$. miss rate is 25%.

finally, miss rate is still 25%.

D.

No

if B is still 16, $\text{sizeof}(\text{int}) = 4$, block can only cache 4 int one time.

read int first time toggle memory cache, miss; next 3 time read hit.

25% miss rate is lowest.

E.

Yes

assume $B = 32$, block cache 8 int at one time.

only 1 miss in 8 time int read.

miss rate can be 12.5%.

6.37

```
/*
 * 6.37.c
 */

typedef int array_t[N][N];

int sumA(array_t a) {
    int i, j;
    int sum = 0;
    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];

    return sum;
}

int sumB(array_t a) {
    int i, j;
    int sum = 0;
    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            sum += a[j][i];

    return sum;
}

int sumC(array_t a) {
    int i, j;
    int sum = 0;
    for (i = 0; i < N; i+=2)
        for (j = 0; j < N; j+=2)
            sum += (a[j][i] + a[j][i+1] + a[j+1][i] + a[j+1][i+1])
}
}
```

C = 4096, B = 16, E = 1, S = 256

N = 64

sizeof(array_t) == 64 64 == 4096 == 4C

memory-cache graph

memory address start from 0 and end to 4096×4 .

cell size is 16B. number(0-255) in cell means cache block number that the cell will be cached.

```

0      +-----+
      |    0    |
16     +-----+
      |    1    |
32     +-----+
      |    2    |
48     +-----+
      |    .    |
      |    .    |
      |    .    |
      |    .    |
      |    .    |
4096-16 +-----+
      |   255   |
4096   +-----+
      |    0    |
4096+16 +-----+
      |    1    |
4096+32 +-----+
      |    .    |
      |    .    |
      |    .    |
      |    .    |
      |    .    |
4096*4-16 +-----+
      |   255   |
4096*4  +-----+

```

A. sumA

```
sum += a[i][j];
```

read memory address order:

0, 4, 8, 12,, $4096 \times 4 - 4$

read cache order:

0, 0, 0, 0, 1, 1, 1, 1,, 255, 255, 255, 255, 0, 0, 0, 0,, 255, 255, 255, 255

first cache read miss, next 3 time read hit.

miss rate: 25%

B. sumB

```
sum += a[j][i];
```

read memory address order:

0, 644, 6442, 64463, 4, 644+4, 6442+4, 4096*4-4

read cache order:

0, 16, 32, 48, ... 240,(4 times) 1, 17, 33, ... 241,(4 times) 15, 31, 47, ... 255(4 times)

let's see first read loop:

read cache order loop 4 times

0, 16, 32, 48, ... 240,(4 times)

first loop all miss, next 3 loop all hit

so miss rate is 25%.

C. sumC

```
for (i = 0; i < N; i+=2)
    for (j = 0; j < N; j+=2)
        sum += (a[j][i] + a[j][i+1] + a[j+1][i] + a[j+1][i+1]);
```

easy to see that read $a[j][i+1]$ and $a[j+1][i+1]$ always hit

same like

```
for (i = 0; i < N; i+=2)
    for (j = 0; j < N; j+=2)
        sum += (a[j][i] + a[j+1][i]);
```

same like

```
for (i = 0; i < N; i+=2)
    for (j = 0; j < N; j++)
        sum += a[j][i];
```

total read count = 64×64

because of $i+=2$,

read cache order only loop 2 times

0, 16, 32, 48, ... 240,(2 times)

so miss rate is 50%

total read miss count = $64/2 \times 64 \times 50\% = 64 \times 64/4$

so miss rate is still 25%.

N = 60

A. sumA

```
sum += a[i][j];
```

read memory by step 1

miss rate 25%

B. sumB

```
for (i = 0; i < N; i++)  
    for (j = 0; j < N; j++)  
        sum += a[j][i];
```

it's interesting.

let's see first inner loop $a[0][0] \rightarrow a[59][0]$

read memory address order:

0, 604, 6042, $604 \cdot 59$

read cache order:

0, 15, 30,, 225, (17 numbers) 255, 14, 29,, 224, (17 numbers) 254, 13, 28,, 223, (17 numbers) 253, 12, 27, 42, 57, 72, 87, 102, 117 (9 numbers)

all read miss and store into different blocks

next 3 loops: $a[0][1] \rightarrow a[59][1]$, $a[0][2] \rightarrow a[59][2]$, $a[0][3] \rightarrow a[59][3]$

all hit

although N is smaller and not power of 2, miss rate is still 25%

C. sumC

same as miss rate when $N = 64$

25%

6.38

A.

 $4 * 16 * 16$

B.

`sizeof(point_color) == 16, B = 32`

```
square[i][j].c = 0
```

miss, cache 2 point_color, then

```
square[i][j].m = 0  
square[i][j].y = 0  
square[i][j].k = 0  
square[i][j+1].c = 0  
square[i][j+1].m = 0  
square[i][j+1].y = 0  
square[i][j+1].k = 0
```

all hit

so miss count is $4 * 16 * 16 * 1/8$

C.

 $1/8$

6.39

A.

 $4 * 16 * 16$

B.

`sizeof(point_color) == 16, B = 32`

```
square[j][i].c = 0
```

miss, cache 2 point_color, then

```
square[j][i].m = 0  
square[j][i].y = 0  
square[j][i].k = 0
```

all hit.

next loop

```
square[j+1][i] - square[j][i] == 16*16 == 256
```

square[j+1][i] miss, cache block not conflict with square[j][i]

```
square[j+8][i] - square[j][i] == 16*16*8 == 2048
```

square[j+8][i] miss, cache block overwrite square[j][i] block. so when we reach square[j][i+1], still miss.

so miss count is $4 * 16 * 16 * 1/4$

C.

 $1/4$

6.40

A.

4 16 16

B.

first loop, same like 6.38, but

write count is 16×16 , miss rate is $1/2$.

second loop, same like 6.39, but

write count is 16×163 , miss rate is $1/6$.

miss count is

$$16 \times 16 \times 1/2 + 16 \times 163 \times 1/6$$

C.

 $1/4$

6.41

every loop

```
buffer[i][j].r = 0;
```

always miss, then cache one pixel, so

```
buffer[i][j].g = 0;  
buffer[i][j].b = 0;  
buffer[i][j].a = 0;
```

all hit

miss rate is 1/4

6.42

same like

```
for (i = 0; i < 640; i++)  
  for (j = 0; j < 480; j++)  
    buffer[i][j].r = 0;  
    buffer[i][j].g = 0;  
    buffer[i][j].b = 0;  
    buffer[i][j].a = 0;
```

C = 64KB, B = 4B, sizeof(pixel) = 4

```
buffer[i][j].r = 0;
```

miss, cache one pixel, so

```
buffer[i][j].g = 0;  
buffer[i][j].b = 0;  
buffer[i][j].a = 0;
```

all hit

miss rate is 1/4

6.43

same like

```
for (i = 0; i < 640; i++)  
    for (j = 0; j < 480; j++)  
        (int*)&buffer[i][j] = 0;
```

every loop,

```
(int*)&buffer[i][j] = 0;
```

always miss

miss rate 100%

6.44

download mountain program from [here](#)

I've download it into chapter6/code dir

run

```
(cd chapter6/code/mountain; ./mountain)
```

see result

Clock frequency is approx. 2500.0 MHz

Memory mountain (MB/sec)

	s1	s2	s3	s4	s5	s6	s7	s8
	s9	s10	s11	s12	s13	s14	s15	
128m	12824	7552	5119	3776	2981	2452	2080	
1799	1663	1563	1483	1410	1372	1334	1304	
64m	12880	7575	5121	3790	2976	2456	2079	
1797	1663	1563	1489	1415	1366	1333	1302	
32m	12849	7635	5137	3785	2996	2456	2075	
1795	1671	1570	1486	1419	1372	1321	1299	
16m	12906	7656	5174	3826	3000	2504	2090	
1808	1700	1579	1506	1442	1397	1373	1357	
8m	13045	7832	5321	4072	3121	2577	2194	
1885	1787	1706	1663	1633	1622	1621	1703	
4m	13303	8352	5602	4210	3326	2765	2352	
2033	1931	1859	1824	1812	1817	1906	2076	
2m	16265	11003	9150	7368	6024	5014	4310	
3816	3795	3761	3848	3877	3909	3946	4044	
1024k	16708	11597	10180	8150	6605	5544	4772	
4169	4109	4150	4136	4127	4113	4127	4119	
512k	16674	11613	10179	8160	6595	5543	4787	
4182	4245	4242	4300	4360	4416	4508	4668	
256k	16929	12872	11826	10661	8538	7360	6383	
5724	6516	6951	7181	7297	7041	7634	7768	
128k	16992	15381	14141	13145	12497	11444	10028	
9734	9642	9638	9671	9696	9606	9438	9117	
64k	17109	15238	14163	13563	12263	11302	9736	
9412	9175	9162	9239	9200	9267	9229	12245	
32k	18060	17655	17503	16954	16785	16489	16718	
14382	16135	16512	16471	15799	14318	16250	15337	
16k	17809	17534	17326	16516	16783	16406	16250	
13333	14043	15280	13890	12546	13347	12394	17062	

6.45

assume matrix size $N = 4$, cache block size $B = 8\text{Byte}$, $\text{sizeof}(\text{int}) = 4$

```
matrix size 4*4
+---+---+---+
|0 |1 |2 |3 |
+---+---+---+
|4 |5 |6 |7 |
+---+---+---+
|8 |9 |10|11|
+---+---+---+
|12|13|14|15|
+---+---+---+
```

function transpose

```
void transpose(int* dst, int* src, int N) {
    int i, j;

    for (i = 0; i <= N-1; i++)
        for (j = 0; j <= N-1; j++)
            dst[j*N+i] = src[i*N+j];
}
```

every cache block can store 2 int numbers

traverse src by step 1, order:

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15

order 0: cache miss, load src[0],src[1] order 1: cache hit because of order 0 order 2: cache miss, load src[2],src[3]
order 15: cache hit

if B is greater, hit rate will be greater too.

traverse dst by step 4, order:

0 4 8 12 1 5 9 13 2 6 10 14 3 7 11 15

element 0: miss, load dst[0],dst[1] element 4: miss, load dst[4],dst[5] element 8: miss, load dst[8],dst[9] element 12:
miss, load dst[12],dst[13] element 1: interesting, order 0 has loaded dst[1] into cache, hit. but in many cases, element
4/8/12 may use the same cache block with element 0. so miss is also possible. element 15: hit / miss

let's assume worst case: all element miss cache

if we split matrix by 2*2

split matrix by size 2*2 block

```
+---+---+---+
|   |   |   |
+ 0 + 1 +
|   |   |   |
+---+---+---+
|   |   |   |
+ 2 + 3 +
|   |   |   |
+---+---+---+
```

transpose block 0 itself

transpose block 1 with 2

transpose block 3 itself

code

```
for (i = 0; i <= N-2; i+=2)
    for (j = 0; j <= N-2; j+=2) {
        dst[j*N+i] = src[i*N+j];
        dst[j*N+i+1] = src[(i+1)*N+j];
        dst[(j+1)*N+i] = src[i*N+j+1];
        dst[(j+1)*N+i+1] = src[(i+1)*N+j+1];
    }
```

when i = 0, j = 0, transpose block 0 itself

```
+---+---+      +---+---+
|0 |1 |        |0 |4 |
+---+---+  =>  +---+---+
|4 |5 |        |1 |5 |
+---+---+      +---+---+
```

```
dst[0] = src[0];
dst[1] = src[4];
dst[4] = src[1];
dst[5] = src[5];
```

if element 0 is miss, element 1 must hit; if element 4 is miss, element 5 must hit;

50% is the highest hit rate in such low cache block size.

if B is greater and cache size C is larger, we can split matrix into 44, 88 or more larger. theoretically we will archive the highest hit rate.

finally code:

```
/*
 * transpose.c
```

```

*/
#include <assert.h>
#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>

// a large prime number
#define MATRIX_N 9973
#define MEM_SIZE (sizeof(int) * MATRIX_N * MATRIX_N)
#define LOOP 1000
#define BLOCK 16

void randomize(void *mem, size_t size) {
    int rnd = open("/dev/urandom", O_RDONLY);
    read(rnd, mem, size);
    close(rnd);
}

void transpose(int *dst, int *src, int N) {
    int i, j;

    for (i = 0; i <= N - 1; i++)
        for (j = 0; j <= N - 1; j++)
            dst[j * N + i] = src[i * N + j];
}

void effective_transpose(int *dst, int *src, int N) {
    int i, j, a, b;

    for (i = 0; i <= N - BLOCK; i += BLOCK)
        for (j = 0; j <= N - BLOCK; j += BLOCK)
            for (a = i; a < i + BLOCK; a++)
                for (b = j; b < j + BLOCK; b++)
                    dst[b * N + a] = src[a * N + b];

    int offset = i;

    for (i = offset; i <= N - 1; i++)
        for (j = 0; j < offset; j += BLOCK)
            for (b = j; b < j + BLOCK; b++)
                dst[b * N + i] = src[i * N + b];

    for (i = 0; i <= N - 1; i++)
        for (j = offset; j <= N - 1; j++)
            dst[j * N + i] = src[i * N + j];
}

void test(void) {
    int *d = (int *)malloc(MEM_SIZE);
    int *s = (int *)malloc(MEM_SIZE);
    randomize((void *)s, MEM_SIZE);

```

```

transpose(d, s, MATRIX_N);

for (int i = 0; i < MATRIX_N; i++)
    for (int j = 0; j < MATRIX_N; j++)
        assert(s[i * MATRIX_N + j] == d[j * MATRIX_N + i]);

memset(d, 0, MEM_SIZE);
effective_transpose(d, s, MATRIX_N);

for (int i = 0; i < MATRIX_N; i++)
    for (int j = 0; j < MATRIX_N; j++)
        assert(s[i * MATRIX_N + j] == d[j * MATRIX_N + i]);

free((void *)d);
free((void *)s);
}

void prof(void) {
    int *d = (int *)malloc(MEM_SIZE);
    int *s = (int *)malloc(MEM_SIZE);

    for (int c = 0; c < LOOP; c++) transpose(d, s, MATRIX_N);

    free((void *)d);
    free((void *)s);
}

void prof_effect(void) {
    int *d = (int *)malloc(MEM_SIZE);
    int *s = (int *)malloc(MEM_SIZE);

    for (int c = 0; c < LOOP; c++) effective_transpose(d, s, MATRIX_N);

    free((void *)d);
    free((void *)s);
}

int main(int argc, char *argv[]) {
    test();

    /* prof(); */
    /* prof_effect(); */

    return 0;
}

```

in code, matrix size 1024*1024, loop 1000 times to measure program run time.

change BLOCK from 2 to 16, record time statistics

origin function run time: 16.46s

BLOCK	time(s)
2	9.99
3	7.16
4	5.6
5	5.66
6	5.34
7	5.39
8	5.38
9	5.48
10	6.21
11	7.9
12	10.17
13	11.14
14	11.88
15	12.11
16	11.85

tip: look up cpu cache info

cat /sys/devices/system/cpu/cpu0/cache/*

6.46

same like 6.45, pay attention to brilliant comment :)

```

/*
 * convert.c
 */
#include <assert.h>
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <string.h>
#include <unistd.h>
#include <time.h>

// a large prime number
#define MATRIX_N 9973
#define MEM_SIZE (sizeof(int) * MATRIX_N * MATRIX_N)
#define LOOP 1000
#define BLOCK 16

void randomize(int *arr, int N) {
    srand(time(0));

    int i, j;
    for (i = 0; i <= N - 1; i++)
        for (j = 0; j <= N - 1; j++)
            arr[i * N + j] = rand() % 2;
}

void convert(int *src, int N) {
    int i, j;

    for (i = 0; i <= N - 1; i++)
        for (j = 0; j <= N - 1; j++)
            src[j * N + i] = src[i * N + j] || src[j * N + i];
}

void effective_convert(int *src, int N) {
    int i, j, a, b, tmp;

    for (i = 0; i <= N - BLOCK; i += BLOCK)
        /* not j = 0 here */
        for (j = i; j <= N - BLOCK; j += BLOCK)
            for (a = i; a < i + BLOCK; a++)
                for (b = j; b < j + BLOCK; b++) {
                    /* brilliant! store two value in one loop */
                    tmp = src[b * N + a] || src[a * N + b];
                    src[b * N + a] = tmp;
                    src[a * N + b] = tmp;
                }
}

```

```

int offset = i;

for (i = offset; i <= N - 1; i++)
    for (j = 0; j < offset; j += BLOCK)
        for (b = j; b < j + BLOCK; b++) {
            tmp = src[b * N + i] || src[i * N + b];
            src[b * N + i] = tmp;
            src[i * N + b] = tmp;
        }

for (i = offset; i <= N - 1; i++)
    for (j = i; j <= N - 1; j++) {
        tmp = src[j * N + i] || src[i * N + j];
        src[j * N + i] = tmp;
        src[i * N + j] = tmp;
    }
}

void test(void) {
    int *s = (int *)malloc(MEM_SIZE);
    int *e = (int *)malloc(MEM_SIZE);

    randomize(s, MATRIX_N);
    memcpy(e, s, MEM_SIZE);

    convert(s, MATRIX_N);
    effective_convert(e, MATRIX_N);

    for (int i = 0; i < MATRIX_N; i++)
        for (int j = 0; j < MATRIX_N; j++)
            assert(s[i * MATRIX_N + j] == e[i * MATRIX_N + j]);

    free((void *)s);
    free((void *)e);
}

void prof(void) {
    int *s = (int *)malloc(MEM_SIZE);

    for (int c = 0; c < LOOP; c++)
        convert(s, MATRIX_N);

    free((void *)s);
}

void prof_effect(void) {
    int *e = (int *)malloc(MEM_SIZE);

    for (int c = 0; c < LOOP; c++)
        effective_convert(e, MATRIX_N);

    free((void *)e);
}

```

```
}

int main(int argc, char *argv[]) {
    test();

    /* prof(); */
    /* prof_effect(); */

    return 0;
}
```

measure time statistics:

origin function run time: 30.28s

BLOCK	time(s)
1	14.26
2	12.01
3	7.43
4	6.20
5	6.08
6	5.86
7	5.70
8	5.67
9	6.30
10	6.39
11	6.21
12	6.18
13	5.9
14	6.3
15	5.88
16	5.92

Linking

Our ultimate goal is extensible programming. By this, we mean the construction of hierarchies of modules, each module adding new functionality to the system.

by Niklaus Wirth

7.1 - 7.5 visit book

7.6 - 7.13 visit here

test

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right

solution	code file	test way
7.6	7.6/swap.c	----
7.7	7.7/*	output
7.8	----	----
7.9	7.9/*	output
7.10	----	----
7.11	----	----
7.12	----	----
7.13	7.13/*	----

7.6

```

/*
 * swap.c
 */

extern int buf[];

int* bufp0 = &buf[0];
static int* bufp1;

static void incr() {
    static int count=0;
    count++;
}

void swap() {
    int temp;

    incr();
    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}

```

	in .symtab?	type	module	section
buf	Yes	external	m	.data
bufp0	Yes	global	swap	.data
bufp1	Yes	local	swap	.bss
swap	Yes	global	swap	.text
temp	No	-----	-----	-----
incr	Yes	local	swap	.text
count	Yes	local	swap	.bss

```
(cd chapter7/code/7.6; make && make sym)
```

output:

```
gcc -c swap.c
objdump -t swap.o
```

```
swap.o      elf64-x86-64
```

SYMBOL TABLE:

0000000000000000 l	0	.bss	0000000000000008 bufp1
0000000000000000 l	F	.text	0000000000000015 incr
0000000000000008 l	0	.bss	0000000000000004 count.1747
0000000000000000 g	0	.data	0000000000000008 bufp0
0000000000000000		*UND*	0000000000000000 buf
0000000000000015 g	F	.text	0000000000000049 swap

7.7

```
/*  
 * bar5.c  
 */  
double x;  
  
void f() {  
    /*x = -0.0;*/  
}
```

delete line

```
x = -0.0;
```

7.8

A.

main.1

main.2

B.

unknown

unknown

C.

error

error

7.9

function main print 0x55 on my machine.

modify char main to unsigned int main here

```
/*
 * bar6.c
 */
#include <stdio.h>

unsigned int main;

void p2() {
    printf("0x%x\n", main);
}
```

```
/*
 * foo6.c
 */
void p2(void);

void offset(void) {
    return;
}

int main(int argc, char* argv[]) {
    p2();
    return 0;
}
```

```
(cd chapter7/code/7.9; make && ./main)
```

output:

```
gcc foo6.c bar6.c -o main
/usr/lib/gcc/x86_64-pc-linux-gnu/4.9.4/../../../../x86_64-pc-linux-gnu/bin/ld: Warning:
alignment 1 of symbol `main' in /tmp/ccTBhRjm.o is smaller than 4 in /tmp/ccc3SjbF.o

0xe5894855
```

using objdump inspect

```
objdump -d main
```

find function main

```

000000000040055d <main>:
  40055d:  55                      push   %rbp
  40055e:  48 89 e5                mov    %rsp,%rbp
  400561:  48 83 ec 10             sub    $0x10,%rsp
  400565:  89 7d fc                mov    %edi,-0x4(%rbp)
  400568:  48 89 75 f0             mov    %rsi,-0x10(%rbp)
  40056c:  e8 07 00 00 00         callq  400578 <p2>
  400571:  b8 00 00 00 00         mov    $0x0,%eax
  400576:  c9                      leaveq
  400577:  c3                      retq

```

0xe5894855 is first 2 instructions content of function main.

works same like

```

/*
 * another-bar6.c
 */
#include <stdio.h>

int main(int argc, char* argv[]);

void p2() {
    printf("0x%x\n", * (unsigned int *)main);
}

```

7.10

A.

```
gcc p.o libx.a
```

B.

```
gcc p.o libx.a liby.a libx.a
```

C.

```
gcc p.o libx.a liby.a libx.a libz.a
```


space for section .bss

7.12

A.

 $\text{ADDR}(s) = \text{ADDR}(.text) = 0x4004e0$ $\text{ADDR}(r.symbol) = \text{ADDR}(swap) = 0x4004f8$ $\text{refaddr} = \text{ADDR}(s) + r.offset = 0x4004ea$ $*\text{refptr} = (\text{unsigned}) (\text{ADDR}(r.symbol) + r.addend - \text{refaddr}) = 0xa$

B.

 $\text{ADDR}(s) = \text{ADDR}(.text) = 0x4004d0$ $\text{ADDR}(r.symbol) = \text{ADDR}(swap) = 0x400500$ $\text{refaddr} = \text{ADDR}(s) + r.offset = 0x4004da$ $*\text{refptr} = (\text{unsigned}) (\text{ADDR}(r.symbol) + r.addend - \text{refaddr}) = 0x22$

7.13

A.

libm.a path

```
whereis libm.a
```

output:

```
libm: /usr/lib64/libm.a /usr/lib64/libm.so
```

libm.a files

```
ar t /usr/lib64/libm.a
```

output:

```
s_lib_version.o
s_matherr.o
s_signgam.o
fclrexcpt.o
fgetexcptflg.o
fraiseexcpt.o
fsetexcptflg.o
ftestexcept.o
fegetround.o
fesetround.o
fegetenv.o
feholdexcpt.o
fesetenv.o
feupdateenv.o
t_exp.o
fedisblxcpt.o
feenablcpt.o
fegetexcept.o
powl_helper.o
e_acos.o
e_acosh.o
e_asin.o
e_atan2.o
e_atanh.o
....
```

similar way for libc.a

B.

compile code with -Og and -Og -g

```
/*  
 * little.c  
 */  
int main(int argc, char* argv[]) {  
    return 0;  
}
```

```
(cd chapter7/code/7.13; make && make dump-exe-code)  
# objdump -d og-little  
# objdump -d dog-little
```

they are same

C.

```
ldd og-little
```

output:

```
linux-vdso.so.1 (0x00007ffef51d3000)  
libc.so.6 => /lib64/libc.so.6 (0x00007f27c6b8b000)  
/lib64/ld-linux-x86-64.so.2 (0x00007f27c6f24000)
```

Exceptional Control Flow

Nature provides exceptions to every rule.

by Margaret Fuller

`csapp.h` & `csapp.c` is downloaded from [csapp site](#)

test

8.1 - 8.8 visit book

8.9 - 8.26 visit [here](#)

test

code directory: `./code`

test way:

- output means to watch code output to judge if it works right

solution	code file	test way
8.9	----	----
8.10	----	----
8.11	8.11.c	output
8.12	8.12.c	output
8.13	8.13.c	output
8.14	8.14.c	output
8.15	8.15.c	output
8.16	8.16.c	output
8.17	----	----
8.18	8.18.c	output
8.19	8.19.c	output
8.20	8.20.c	output
8.21	8.21.c	output
8.22	mssystem.c, exit-code.c, wait-sig.c	output
8.23	8.23.c	output
8.24	8.24.c	output
8.25	8.25.c	output
8.26	shell/*	output

8.9

process pair	cocurrent?
AB	No
AC	Yes
AD	Yes
BC	Yes
BD	Yes
CD	Yes

8.10

A. call once, return twice

fork

B. call once, never return

longjmp, execve

C. call once, return 1 or more times

setjmp

8.11

4 lines

```

      c
    +-----+
    |       | "hello"
    |       |
  c |       | p
+---+-----+-----+
|   fork   | "hello"
|           |
|           | c
|           | +-----+
|           | | "hello"
|           | |
|   p       | | p
+---+-----+-----+
main fork   fork   "hello"
```

8.12

8 lines

```

      c
      +-----+-----+
      | "hello"  "hello"
      |
      c | p
+-----+-----+-----+
|      fork  "hello"  "hello"
|
|      c
|      +-----+-----+
|      | "hello"  "hello"
|      |
|      p | p
+-----+-----+-----+
main fork  fork  "hello"  "hello"
```

8.13

```
x=4  
x=3  
x=2
```

pay attention, parent process and child process don't share global x, they have own private x.

8.14

3 lines

```
          c
        +-----+
        |      "hello"
        |
      c  |  p
    +-----+-----+
    |      fork  "hello"
    |
    |  p
  +-----+-----+-----+
main fork  return "hello"
```

8.15

5 lines

```
          c
        +-----+-----+
        |      "hello"  "hello"
        |
      c  |  p
    +-----+-----+-----+
    |      fork  "hello"  "hello"
    |
    |
    |  p
+-----+-----+-----+
main fork  return "hello"
```

```
counter = 2
```

child process has its own counter.

8.17

```
hello 0 1 Bye 2 Bye
hello 1 0 Bye 2 Bye
hello 1 Bye 0 2 Bye
```

8.18

```

              c
            +-----+-----+
            |           "0"      exit "2"
            |
            |
          c  |  p
        +-----+-----+-----+
        |      fork      "1"      exit "2"
        |    (atexit)
        |              c
        |      +-----+-----+
        |      |           "0"      exit
        |      |
        |      |  p      |  p
        +-----+-----+-----+
    main  fork      fork      "1"      exit

```

2 must be behind 0/1

B & D is impossible.

8.19

 2^n

8.20

```
/*
 * 8.20.c
 */
#include <stdio.h>
#include "csapp.h"

int main(int argc, char* argv[], char* env[]) {
    if (execve("/bin/ls", argv, env) == -1) {
        fprintf(stderr, "execve error: %s\n", strerror(errno));
        exit(1);
    }
}
```

8.21

abc

or

bac

8.22

```
/*
 * mysystem.c
 */
#include <stdio.h>
#include "csapp.h"

int mysystem(char* command) {
    pid_t pid;
    int status;

    if ((pid = Fork()) == 0) {
        /* child process */
        char* argv[4] = { "", "-c", command, NULL };
        execve("/bin/sh", argv, environ);
    }

    /* print child pid so we can kill it */
    printf("child pid: %d\n", pid);

    if (Waitpid(pid, &status, 0) > 0) {
        /* exit normally */
        if (WIFEXITED(status))
            return WEXITSTATUS(status);

        /* exit by signal */
        if (WIFSIGNALED(status))
            return WTERMSIG(status);
    }
}

int main(int argc, char* argv[]) {
    int code;

    code = mysystem("./exit-code");
    printf("normally exit, code: %d\n", code); fflush(stdout);

    code = mysystem("./wait-sig");
    printf("exit caused by signal, code: %d\n", code); fflush(stdout);
    return 0;
}
```

when running `./mysystem`, it runs `./exit-code`

```
/*
 * exit-code.c
 */
#include "csapp.h"

int main(int argc, char* argv[]) {
    exit(10);
}
```

should output

```
normally exit, code 10
```

and runs `./wait-sig`, stuck here

```
/*
 * wait-sig.c
 */
#include "csapp.h"

int main(int argc, char* argv[]) {
    while (1);
}
```

fall into dead loop. open another terminal, type

```
kill -<n> <wait-sig's pid>
```

`./mysystem` will return and output

```
exit caused by signal, code <n>
```

8.23

SIGUSR2	SIGUSR2	SIGUSR2	SIGUSR2	SIGUSR2
being handled	Pending	Canceled	Canceled	Canceled
need 1 sec				

there's only one pending signal all the time. other same type signals will be canceled.

if you modify code

remove sleep

```
/* sleep(1); */
```

or

send more signals

```
for (i = 0; i < 500000; i++)
```

you will get different counter output.

```
/* $begin waitpid1 */
#include "csapp.h"
#define N 2
#define LEN 100

int main() {
    int status, i;
    pid_t pid;

    /* Parent creates N children */
    for (i = 0; i < N; i++)
        if ((pid = Fork()) == 0) {
            /* access address 0, cause fault */
            char* cptr = NULL;
            *cptr = 'd';
        }

    /* Parent reaps N children in no particular order */
    while ((pid = waitpid(-1, &status, 0)) > 0) {
        if (WIFEXITED(status))
            printf("child %d terminated normally with exit status=%d\n",
                pid, WEXITSTATUS(status));
        else if (WIFSIGNALED(status)) {
            /* print signal that cause process exit */
            char buf[LEN];
            sprintf(buf, "child %d terminated by signal %d", pid, WTERMSIG(status));
            psignal(WTERMSIG(status), buf);
        }
        else
            printf("child %d terminated abnormally\n", pid);
    }

    /* The only normal termination is if there are no more children */
    if (errno != ECHILD)
        unix_error("waitpid error");

    exit(0);
}
/* $end waitpid1 */
```

8.25

```
/*
 * 8.25.c
 */
#include <stdio.h>
#include "csapp.h"

sigjmp_buf buf;

void handler(int sig) {
    /* jump */
    siglongjmp(buf, 1);
}

char* tfgets(char* s, int size, FILE* stream) {
    char* result;

    if (!sigsetjmp(buf, 1)) {
        alarm(5);
        if (signal(SIGALRM, handler) == SIG_ERR)
            unix_error("set alarm handler error");
        return fgets(s, size, stream);
    } else {
        /* run out of time */
        return NULL;
    }
}

#define LEN 100

int main(int argc, char* argv[]) {
    char buf[LEN];
    char* input = tfgets(buf, LEN, stdin);

    if (input == NULL) {
        printf("nothing input: NULL\n");
    } else {
        printf("%s", input);
    }
    return 0;
}
```


8.26

learn

- origin shell code in pic 8-23, pic 8-24, pic 8-25
- how to use waitpid in pic 8-18
- how to wait child process in pic 8-42

and write a job management module

compose them and the little shell is born.

add `bg` `fg` `jobs` command to origin shell program

```
#include <assert.h>
#include "../csapp.h"
#include "shell.h"
#include "job.h"

void eval(char *cmdline)
{
    char *argv[MAXARGS]; /* Argument list execve() */
    char buf[MAXLINE];    /* Holds modified command line */
    int bg;               /* Should the job run in bg or fg? */
    pid_t pid;            /* Process id */

    strcpy(buf, cmdline);
    bg = parse_line(buf, argv);
    if (argv[0] == NULL)
        return; /* Ignore empty lines */

    if (!builtin_command(argv)) {
        sigset_t mask_one, prev_one;
        Sigemptyset(&mask_one);
        Sigaddset(&mask_one, SIGCHLD);

        /* block signal child */
        Sigprocmask(SIG_BLOCK, &mask_one, &prev_one);
        if ((pid = Fork()) == 0) {
            /* unblock in child process */
            Sigprocmask(SIG_SETMASK, &prev_one, NULL);

            /* set gid same like pid */
            Setpgid(0, 0);

            if (execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }
        }

        sigset_t mask_all, prev_all;
        Sigfillset(&mask_all);
        // save job info
```

```

Sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
Jid new_jid = new_job(pid, cmdline, !bg);
Sigprocmask(SIG_SETMASK, &prev_all, NULL);

if (!bg) {
    set_fg_pid(pid);
    while(get_fg_pid())
        sigsuspend(&prev_one);
}
else
    printf("[%d] %d %s \t %s\n", new_jid, pid, "Running", cmdline);

/* unblock child signal */
Sigprocmask(SIG_SETMASK, &prev_one, NULL);
}
return;
}

/*
 * If first arg is a builtin command, run it and return true;
 * else return false.
 */
int builtin_command(char **argv)
{
    if (!strcmp(argv[0], "quit")) /* quit command */
        exit(0);
    if (!strcmp(argv[0], "&")) /* Ignore singleton & */
        return 1;
    if (!strcmp(argv[0], "jobs")) {
        print_jobs();
        return 1;
    }
    // > fg
    if (!strcmp(argv[0], "fg")) {
        int id;
        // right format: fg %ddd or fg ddd
        if ((id = parse_id(argv[1])) != -1 && argv[2] == NULL) {
            sigset_t mask_one, prev_one;
            Sigemptyset(&mask_one);
            Sigaddset(&mask_one, SIGCHLD);
            Sigprocmask(SIG_BLOCK, &mask_one, &prev_one);

            pid_t pid = id;
            // if param is jid
            if (argv[1][0] == '%') {
                JobPtr jp = find_job_by_jid(id);
                pid = jp->pid;
            }
            Kill(pid, SIGCONT);
            set_fg_pid(pid);
            while(get_fg_pid())
                sigsuspend(&prev_one);
        }
    }
}

```

```

    Sigprocmask(SIG_SETMASK, &prev_one, NULL);

} else {
    printf("format error, e.g. fg %%12 || fg 1498\n");
}

return 1;
}

// > bg
if (!strcmp(argv[0], "bg")) {
    int id;
    // right format: bg %ddd or bg ddd
    if ((id = parse_id(argv[1])) != -1 && argv[2] == NULL) {
        pid_t pid = id;
        // jid param
        if (argv[1][0] == '%') {
            JobPtr jp = find_job_by_jid(id);
            pid = jp->pid;
        }
        Kill(pid, SIGCONT);
    } else {
        printf("format error, e.g. bg %%12 or bg 1498\n");
    }
    return 1;
}

return 0; /* Not a builtin command */
}

/* parse_line - Parse the command line and build the argv array */
int parse_line(char *buf, char **argv)
{
    char *delim; /* Points to first space delimiter */
    int argc; /* Number of args */
    int bg; /* Background job? */

    buf[strlen(buf)-1] = ' '; /* Replace trailing '\n' with space */
    while (*buf && (*buf == ' ')) /* Ignore leading spaces */
        buf++;

    /* Build the argv list */
    argc = 0;
    while ((delim = strchr(buf, ' '))) {
        argv[argc++] = buf;
        *delim = '\0';
        buf = delim + 1;
        while (*buf && (*buf == ' ')) /* Ignore spaces */
            buf++;
    }
    argv[argc] = NULL;

    if (argc == 0) /* Ignore blank line */
        return 1;
}

```

```

/* Should the job run in the background? */
if ((bg = (*argv[argc-1] == '&')) != 0)
    argv[--argc] = NULL;

return bg;
}

static int is_number_str(char* s) {
    int len = strlen(s);
    for (int i = 0; i < len; i++)
        if (!isdigit(s[i]))
            return 0;

    return 1;
}

int parse_id(char* s) {
    int error = -1;
    if (s == NULL)
        return error;

    /* format: %dddddd */
    if (s[0] == '%') {
        if (!is_number_str(s+1))
            return error;

        return atoi(s+1);
    }
    /* format: ddddddd */
    if (is_number_str(s))
        return atoi(s);

    /* not right */
    return error;
}

void test_shell() {
    // parse id
    assert(-1 == parse_id("ns"));
    assert(-1 == parse_id("%%"));
    assert(0 == parse_id("%0"));
    assert(0 == parse_id("0"));
    assert(98 == parse_id("%98"));
    assert(98 == parse_id("98"));
}

```

job management module. signal handlers are the key part

```

/*
 * job.c
 */

```

```

#include <stdio.h>
#include <assert.h>
#include "job.h"
#include "../csapp.h"

static volatile sig_atomic_t fg_pid;
static Job jobs[MAXJOBS];

int is_fg_pid(pid_t pid) {
    return fg_pid == pid;
}

pid_t get_fg_pid() {
    return fg_pid;
}

void set_fg_pid(pid_t pid) {
    fg_pid = pid;
}

/* SIGCONT signal */
void sigchild_handler(int sig) {
    int old_errno = errno;
    int status;
    pid_t pid;

    sigset_t mask_all, prev_all;
    Sigfillset(&mask_all);

    /* exit or be stopped or continue */
    while ((pid = waitpid(-1, &status, WNOHANG|WUNTRACED|WCONTINUED)) > 0) {
        /* exit normally */
        if (WIFEXITED(status) || WIFSIGNALED(status)) {
            if (is_fg_pid(pid)) {
                set_fg_pid(0);
            } else {
                Sio_puts("pid "); Sio_putl(pid); Sio_puts(" terminates\n");
            }
            Sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
            del_job_by_pid(pid);
            Sigprocmask(SIG_SETMASK, &prev_all, NULL);
        }

        /* be stopped */
        if (WIFSTOPPED(status)) {
            if (is_fg_pid(pid)) {
                set_fg_pid(0);
            }
            // set pid status stopped
            Sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
            JobPtr jp = find_job_by_pid(pid);
            set_job_status(jp, Stopped);
            Sigprocmask(SIG_SETMASK, &prev_all, NULL);

            Sio_puts("pid "); Sio_putl(pid); Sio_puts(" be stopped\n");
        }
    }
}

```

```

    }

    /* continue */
    if(WIFCONTINUED(status)) {
        set_fg_pid(pid);
        // set pid status running
        Sigprocmask(SIG_BLOCK, &mask_all, &prev_all);
        JobPtr jp = find_job_by_pid(pid);
        set_job_status(jp, Running);
        Sigprocmask(SIG_SETMASK, &prev_all, NULL);

        Sio_puts("pid "); Sio_putl(pid); Sio_puts(" continue\n");
    }
}

errno = old_errno;
}

void sigint_handler(int sig) {
    /* when fg_pid == 0, stop shell itself, it'll be a dead loop */
    if (is_fg_pid(0)) {
        Signal(SIGINT, SIG_DFL);
        Kill(getpid(), SIGINT);
    } else {
        Kill(get_fg_pid(), SIGINT);
    }
}

void sigstop_handler(int sig) {
    /* same like int handler */
    if (is_fg_pid(0)) {
        Signal(SIGTSTP, SIG_DFL);
        Kill(getpid(), SIGTSTP);
    } else {
        Kill(get_fg_pid(), SIGTSTP);
    }
}

JobPtr find_job_by_jid(Jid jid) {
    return &(jobs[jid]);
}

JobPtr find_job_by_pid(pid_t pid) {
    for (int i = 0; i < MAXJOBS; i++) {
        Job j = jobs[i];
        if (j.using && j.pid == pid) {
            return &(jobs[i]);
        }
    }
    /* no such job */
    return NULL;
}

```

```

void set_job_status(JobPtr jp, enum JobStatus status) {
    if (jp)
        jp->status = status;
}

// seek a spare place for new job
static int find_spare_jid() {
    Jid jid = -1;
    for (int i = 0; i < MAXJOBS; i++) {
        if (jobs[i].using == 0) {
            jid = i;
            break;
        }
    }
    return jid;
}

int new_job(pid_t pid, char* cmdline, int fg) {
    // find a jid
    Jid jid = find_spare_jid();
    if (jid == -1)
        unix_error("no more jid to use");

    // save process info
    jobs[jid].jid = jid;
    jobs[jid].pid = pid;
    jobs[jid].status = Running;
    strcpy(jobs[jid].cmdline, cmdline);
    jobs[jid].using = 1;

    return jid;
}

void del_job_by_pid(pid_t pid) {
    // search job whose pid is pid
    for (int i = 0; i < MAXJOBS; i++) {
        if (jobs[i].using && jobs[i].pid == pid) {
            // delete job
            jobs[i].using = 0;
        }
    }
}

void print_jobs() {
    for (int i = 0; i < MAXJOBS; i++) {
        Job j = jobs[i];
        if (j.using) {
            printf("[%d] %d %s \t %s\n", j.jid, j.pid,
                j.status == Running ? "Running" : "Stopped", j.cmdline);
        }
    }
}

```

```
void init_jobs() {  
    memset(jobs, 0, sizeof(jobs));  
}  
  
void test_job() {  
  
}
```

test it

```
(cd chapter8/code/shell; make && ./shell)
```

ps: `./loop` is a dead loop program, `./sleep` sleeps 5 secs and exit.

```
> jobs  
> ./loop  
^Zpid 4948 be stopped  
> jobs  
[0] 4948 Stopped      ./loop  
  
> fg %0  
pid 4948 continue  
^Zpid 4948 be stopped  
> ./loop &  
[1] 4950 Running      ./loop &  
  
> jobs  
[0] 4948 Stopped      ./loop  
[1] 4950 Running      ./loop &  
  
> bg 4948  
> pid 4948 continue  
  
> jobs  
[0] 4948 Running      ./loop  
[1] 4950 Running      ./loop &  
  
> quit
```


Virtual Memory

I wanted to have virtual memory, at least as it's coupled with file systems.
by Ken Thompson

9.1 - 9.10 visit book

9.11 - 9.20 visit here

test

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right

solution	code file	test way
9.11	----	----
9.12	----	----
9.13	----	----
9.14	9.14.c	assert
9.15	----	----
9.16	----	----
9.17	vm/(mm.9.17, 9.17.c)	assert
9.18	vm/(mm.9.18, 9.18.c)	assert
9.19	----	----
9.20	malloc/*	measure performance with std malloc

9.11

VA: 0x027c

A.

13	12	11	10	9	8	7	6	5	4	3	2	1	0									
+	-	-	-	-	-	-	-	-	-	-	-	-	+									
	0		0		0		1		0		1		1		1		1		0		0	
+	-	-	-	-	-	-	-	-	-	-	-	-	+									

B.

param	value
VPN	0x09
TLBI	0x01
TLBT	0x02
hit?	No
page falut?	No
PPN	0x17

C.

11	10	9	8	7	6	5	4	3	2	1	0											
+	-	-	-	-	-	-	-	-	-	-	+											
	0		1		0		1		1		1		1		1		1		0		0	
+	-	-	-	-	-	-	-	-	-	-	-	+										

D.

param	value
CO	0x00
CI	0x0F
CT	0x17
hit?	No
value	----

9.12

VA: 0x03a9

A.

13	12	11	10	9	8	7	6	5	4	3	2	1	0							
+-	-	-	-	-	-	-	-	-	-	-	-	-	-							
	0		0		0		1		1		0		1		0		0		1	
+-	-	-	-	-	-	-	-	-	-	-	-	-	-							

B.

param	value
VPN	0x0E
TLBI	0x02
TLBT	0x03
hit?	No
page falut?	No
PPN	0x11

C.

11	10	9	8	7	6	5	4	3	2	1	0											
+-	-	-	-	-	-	-	-	-	-	-	+											
	0		1		0		0		1		1		0		1		0		0		1	
+-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	-	+

D.

param	value
CO	0x01
CI	0x0A
CT	0x11
hit?	No
value	----

9.13

VA: 0x0040

A.

13	12	11	10	9	8	7	6	5	4	3	2	1	0
+	-	-	-	-	-	-	-	-	-	-	-	-	+
	0		0		0		0		0		1		0
	0		0		0		0		0		0		0
+	-	-	-	-	-	-	-	-	-	-	-	-	+

B.

param	value
VPN	0x01
TLBI	0x01
TLBT	0x00
hit?	No
page falut?	Yes

```
/*
 * 9.14.c
 */
#include <stdio.h>
#include <assert.h>
#include "vm/csapp.h"

void test(char* filename, char* content) {
    int fd;
    char buf[20];
    fd = Open(filename, O_RDONLY, 0);
    Read(fd, buf, strlen(content));
    assert( !strcmp(buf, content, strlen(content)) );
}

int touch(char* filename, char* content) {
    int fd;
    umask(DEF_UMASK);
    fd = Open(filename, O_WRONLY|O_CREAT|O_TRUNC, DEF_MODE);
    Write(fd, content, strlen(content));
    Close(fd);
}

int main(int argc, char* argv[]) {
    touch("hello.txt", "Hello, world!");
    test("hello.txt", "Hello, world!");

    struct stat stat;
    int fd;
    char* bufp;
    size_t size;

    fd = Open("hello.txt", O_RDWR, 0);
    fstat(fd, &stat);
    size = stat.st_size;

    bufp = Mmap(NULL, size, PROT_WRITE, MAP_SHARED, fd, 0);
    *bufp = 'J';
    Munmap(bufp, size);

    test("hello.txt", "Jello, world!");
    return 0;
}
```

9.15

malloc	size	header
malloc(3)	8	0x9
malloc(11)	16	0x11
malloc(20)	24	0x19
malloc(21)	24	0x19

9.16

alignment	allocated block	spare block	min block size
word	Header & Footer	Header & Footer	16
word	Header	Header & Footer	16
double word	Header & Footer	Header & Footer	16
double word	Header	Header & Footer	16

9.17

complete code in `chapter9/code/vm/mm.9.17.c`

```

--- mm.c      2017-11-09 02:57:43.679935907 +0000
+++ mm.9.17.c  2017-11-09 02:57:43.679935907 +0000
@@ -41,6 +41,7 @@

/* Global variables */
static char *heap_listp = 0; /* Pointer to first block */
+static char *rover;          /* Next fit rover */

/* Function prototypes for internal helper routines */
static void *extend_heap(size_t words);
@@ -69,6 +70,7 @@
    heap_listp += (2*WSIZE);          //line:vm:mm:endinit
    /* $end mminit */

+ rover = heap_listp;
    /* $begin mminit */

    /* Extend the empty heap with a free block of CHUNKSIZE bytes */
@@ -177,6 +179,10 @@
    bp = PREV_BLKp(bp);
}
/* $end mmfree */
+ /* Make sure the rover isn't pointing into the free block */
+ /* that we just coalesced */
+ if ((rover > (char *)bp) && (rover < NEXT_BLKp(bp)))
+   rover = bp;
    /* $begin mmfree */
    return bp;
}
@@ -290,16 +296,20 @@
{
    /* $end mmfirstfit */

- /* $begin mmfirstfit */
- /* First-fit search */
- void *bp;
-
- for (bp = heap_listp; GET_SIZE(HDRP(bp)) > 0; bp = NEXT_BLKp(bp)) {
-   if (!GET_ALLOC(HDRP(bp)) && (asize <= GET_SIZE(HDRP(bp)))) {
-     return bp;
-   }
- }
- return NULL; /* No fit */
+ /* Next fit search */
+ char *oldrover = rover;
+
+ /* Search from the rover to the end of list */
+ for ( ; GET_SIZE(HDRP(rover)) > 0; rover = NEXT_BLKp(rover))
+   if (!GET_ALLOC(HDRP(rover)) && (asize <= GET_SIZE(HDRP(rover))))

```



```
+     return rover;
+
+ /* search from start of list to old rover */
+ for (rover = heap_listp; rover < oldrover; rover = NEXT_BLKPTR(rover))
+     if (!GET_ALLOC(HDRP(rover)) && (asize <= GET_SIZE(HDRP(rover))))
+         return rover;
+
+ return NULL; /* no fit found */
}
/* $end mmfirstfit */
```

9.18

complete code in `chapter9/code/vm/mm.9.18.c`

```

--- mm.c      2017-11-09 02:57:43.679935907 +0000
+++ mm.9.18.c  2018-02-01 10:08:47.563747327 +0000
@@ -20,7 +20,7 @@
#define MAX(x, y) ((x) > (y)? (x) : (y))

/* Pack a size and allocated bit into a word */
-#define PACK(size, alloc) ((size) | (alloc)) //line:vm:mm:pack
+#define PACK(size, alloc, prev_alloc) ((size) | (alloc) | (prev_alloc << 1)) //line:vm:mm:pack

/* Read and write a word at address p */
#define GET(p) (*(unsigned int *)(p)) //line:vm:mm:get
@@ -29,6 +29,7 @@
/* Read the size and allocated fields from address p */
#define GET_SIZE(p) (GET(p) & ~0x7) //line:vm:mm:getsize
#define GET_ALLOC(p) (GET(p) & 0x1) //line:vm:mm:getalloc
+#define GET_PREV_ALLOC(p) ((GET(p) >> 1) & 0x1)

/* Given block ptr bp, compute address of its header and footer */
#define HDRP(bp) ((char *)(bp) - WSIZE) //line:vm:mm:hdp
@@ -63,9 +64,9 @@
if ((heap_listp = mem_sbrk(4*WSIZE)) == (void *)-1) //line:vm:mm:begininit
    return -1;
PUT(heap_listp, 0); // Alignment padding */
- PUT(heap_listp + (1*WSIZE), PACK(DSIZE, 1)); /* Prologue header */
- PUT(heap_listp + (2*WSIZE), PACK(DSIZE, 1)); /* Prologue footer */
- PUT(heap_listp + (3*WSIZE), PACK(0, 1)); /* Epilogue header */
+ PUT(heap_listp + (1*WSIZE), PACK(DSIZE, 1, 1)); /* Prologue header */
+ PUT(heap_listp + (2*WSIZE), PACK(DSIZE, 1, 1)); /* Prologue footer */
+ PUT(heap_listp + (3*WSIZE), PACK(0, 1, 1)); /* Epilogue header */
heap_listp += (2*WSIZE); //line:vm:mm:endinit
/* $end mminit */

@@ -98,10 +99,10 @@
return NULL;

/* Adjust block size to include overhead and alignment reqs. */
- if (size <= DSIZE) //line:vm:mm:sizeadjust1
-     asize = 2*DSIZE; //line:vm:mm:sizeadjust2
+ if (size <= WSIZE)
+     asize = DSIZE;
else
-     asize = DSIZE * ((size + (DSIZE) + (DSIZE-1)) / DSIZE); //line:vm:mm:sizeadjust3
+     asize = DSIZE * ((size + (WSIZE) + (DSIZE-1)) / DSIZE); //line:vm:mm:sizeadjust3

/* Search the free list for a fit */
if ((bp = find_fit(asize)) != NULL) { //line:vm:mm:findfitcall
@@ -136,8 +137,16 @@
}

```

```

/* $begin mmfree */

- PUT(HDRP(bp), PACK(size, 0));
- PUT(FTRP(bp), PACK(size, 0));
+ PUT(HDRP(bp), PACK(size, 0, GET_PREV_ALLOC(HDRP(bp))));
+ PUT(FTRP(bp), PACK(size, 0, GET_PREV_ALLOC(HDRP(bp))));
+
+ if (GET_ALLOC(HDRP(NEXT_BLK_P(bp))))
+   PUT(HDRP(NEXT_BLK_P(bp)), PACK(GET_SIZE(HDRP(NEXT_BLK_P(bp))), 1, 0));
+ else {
+   PUT(HDRP(NEXT_BLK_P(bp)), PACK(GET_SIZE(HDRP(NEXT_BLK_P(bp))), 0, 0));
+   PUT(FTRP(NEXT_BLK_P(bp)), PACK(GET_SIZE(HDRP(NEXT_BLK_P(bp))), 0, 0));
+ }
+
  coalesce(bp);
}

@@ -148,7 +157,7 @@
/* $begin mmfree */
static void *coalesce(void *bp)
{
- size_t prev_alloc = GET_ALLOC(FTRP(PREV_BLK_P(bp)));
+ size_t prev_alloc = GET_PREV_ALLOC(HDRP(bp));
  size_t next_alloc = GET_ALLOC(HDRP(NEXT_BLK_P(bp)));
  size_t size = GET_SIZE(HDRP(bp));

@@ -158,22 +167,22 @@

  else if (prev_alloc && !next_alloc) {      /* Case 2 */
    size += GET_SIZE(HDRP(NEXT_BLK_P(bp)));
-   PUT(HDRP(bp), PACK(size, 0));
-   PUT(FTRP(bp), PACK(size, 0));
+   PUT(HDRP(bp), PACK(size, 0, 1));
+   PUT(FTRP(bp), PACK(size, 0, 1));
  }

  else if (!prev_alloc && next_alloc) {      /* Case 3 */
    size += GET_SIZE(HDRP(PREV_BLK_P(bp)));
-   PUT(FTRP(bp), PACK(size, 0));
-   PUT(HDRP(PREV_BLK_P(bp)), PACK(size, 0));
+   PUT(FTRP(bp), PACK(size, 0, 1));
+   PUT(HDRP(PREV_BLK_P(bp)), PACK(size, 0, 1));
    bp = PREV_BLK_P(bp);
  }

  else {                                     /* Case 4 */
    size += GET_SIZE(HDRP(PREV_BLK_P(bp))) +
            GET_SIZE(FTRP(NEXT_BLK_P(bp)));
-   PUT(HDRP(PREV_BLK_P(bp)), PACK(size, 0));
-   PUT(FTRP(NEXT_BLK_P(bp)), PACK(size, 0));
+   PUT(HDRP(PREV_BLK_P(bp)), PACK(size, 0, 1));
+   PUT(FTRP(NEXT_BLK_P(bp)), PACK(size, 0, 1));
    bp = PREV_BLK_P(bp);
  }

```

```

    }
    /* $end mmfree */
@@ -246,9 +255,9 @@
    return NULL;                                //line:vm:mm:endextend

    /* Initialize free block header/footer and the epilogue header */
-   PUT(HDRP(bp), PACK(size, 0));                /* Free block header */ //line:vm:mm:freeblo
ckhdr
-   PUT(FTRP(bp), PACK(size, 0));                /* Free block footer */ //line:vm:mm:freeblo
ckftr
-   PUT(HDRP(NEXT_BLKP(bp)), PACK(0, 1)); /* New epilogue header */ //line:vm:mm:newepih
dr
+   PUT(HDRP(bp), PACK(size, 0, GET_PREV_ALLOC(HDRP(bp))));          /* Free block header
*/ //line:vm:mm:freeblockhdr
+   PUT(FTRP(bp), PACK(size, 0, GET_PREV_ALLOC(HDRP(bp))));          /* Free block footer
*/ //line:vm:mm:freeblockftr
+   PUT(HDRP(NEXT_BLKP(bp)), PACK(0, 1, 0)); /* New epilogue header */ //line:vm:mm:newe
pihdr

    /* Coalesce if the previous block was free */
    return coalesce(bp);                        //line:vm:mm:returnblo
ck
@@ -267,15 +276,14 @@
    size_t csize = GET_SIZE(HDRP(bp));

    if ((csize - asize) >= (2*DSIZE)) {
-       PUT(HDRP(bp), PACK(asize, 1));
-       PUT(FTRP(bp), PACK(asize, 1));
+       PUT(HDRP(bp), PACK(asize, 1, 1));
        bp = NEXT_BLKP(bp);
-       PUT(HDRP(bp), PACK(csize-asize, 0));
-       PUT(FTRP(bp), PACK(csize-asize, 0));
+       PUT(HDRP(bp), PACK(csize-asize, 0, 1));
+       PUT(FTRP(bp), PACK(csize-asize, 0, 1));
    }
    else {
-       PUT(HDRP(bp), PACK(csize, 1));
-       PUT(FTRP(bp), PACK(csize, 1));
+       PUT(HDRP(bp), PACK(csize, 1, 1));
+       PUT(HDRP(NEXT_BLKP(bp)), PACK(csize, 1, 1));
    }
}
/* $end mmplace */

```

9.19

1)

a: Right. $2^{(k+1)}$ block size for $(2^k)+1$ allocation request

b: wrong.

c: wrong. LIFO is also fast.

d: wrong. almost every strategy has external fragmentation problem.

2)

a: wrong. first fit should be fast in this condition

b: wrong. should be order by block size

c: wrong. min spare size not max

d: Right.

3)

b

ref 9.10.3

9.20

using malloc lib code from csapp site

- implicit idle list
- allocated block with header & footer
- idle block with header & footer
- no GC
- first fit strategy

another modification to malloc lib file

mm.h

```
--- ../vm/mm.h      2017-11-09 02:57:43.679935907 +0000
+++ mm.h           2017-11-09 02:57:43.679935907 +0000
@@ -5,6 +5,9 @@
extern void mm_free (void *ptr);
/* $end mallocinterface */

+#define malloc(size) mm_malloc(size)
+#define free(ptr) mm_free(ptr)
+
extern void *mm_realloc(void *ptr, size_t size);
extern void *mm_calloc (size_t nmemb, size_t size);
extern void mm_checkheap(int verbose);
```

memlib.c

```

--- ../vm/memlib.c      2017-11-09 02:57:43.679935907 +0000
+++ memlib.c           2017-11-09 02:57:43.675935936 +0000
@@ -15,23 +15,18 @@
#include "csapp.h"
#include "memlib.h"

-#define MAX_HEAP (20*(1<20)) /* 20 MB */
-
-
/* $begin memlib */
/* Private global variables */
static char *mem_heap; /* Points to first byte of heap */
static char *mem_brk; /* Points to last byte of heap plus 1 */
-static char *mem_max_addr; /* Max legal heap addr plus 1*/

/*
 * mem_init - Initialize the memory system model
 */
void mem_init(void)
{
- mem_heap = (char *)Malloc(MAX_HEAP);
+ mem_heap = (char *)sbrk(0);
  mem_brk = (char *)mem_heap;
- mem_max_addr = (char *) (mem_heap + MAX_HEAP);
}

/*
@@ -43,7 +38,7 @@
{
  char *old_brk = mem_brk;

- if ( (incr < 0) || ((mem_brk + incr) > mem_max_addr)) {
+ if ( (incr < 0) || ((mem_brk = sbrk(incr)) == (void *)-1)) {
  errno = ENOMEM;
  fprintf(stderr, "ERROR: mem_sbrk failed. Ran out of memory...\n");
  return (void *)-1;

```

main.c file measure malloc performance

```
/*
 * main.c
 */
#include <stdio.h>
#include "csapp.h"

#ifdef CUS_MALLOC
#include "mm.h"
#include "memlib.h"
#else
#include <stdlib.h>
#endif

#define LOOP 10000

int main(int argc, char* argv[]) {
    void* m_start = sbrk(0);
    size_t malloc_size = 0;

    int i;
    for (i = 0; i < LOOP; i+=2) {
        void* ptr_f = malloc(i);
        void* ptr = malloc(i+1);
        free(ptr_f);

        malloc_size += i+1;
    }

    void* m_end = sbrk(0);
    size_t heap_size = (size_t)(m_end - m_start);

    printf("malloc size: %ld, heap_size: %ld\n", malloc_size, heap_size);

    return 0;
}
```

run `make` to generate both origin main executable file and custom version(using `-DCUS_MALLOC`)


```

CC = gcc
CFLAGS = -m64 -pthread -DCUS_MALLOC
SRCS = mm.c memlib.c csapp.c

all: origin custom diff

measure:
    time ./origin.main
    time ./custom.main

origin:
    $(CC) -m64 main.c -o origin.main

custom:
    $(CC) $(CFLAGS) $(SRCS) main.c -o custom.main

diff:
    (diff -u ../vm/mm.h mm.h > mm.h.diff; cd .)
    (diff -u ../vm/memlib.c memlib.c > memlib.c.diff; cd .)

test:

.PHONY: clean
clean:
    find . -type f -executable -print0 | xargs -0 rm -f --

```

measurement

```

(cd chapter9/code/malloc; make measure)

time ./origin.main
malloc size: 25000000, heap_size: 28311552
0.00user 0.01system 0:00.01elapsed 100%CPU (0avgtext+0avgdata 19256maxresident)k
0inputs+0outputs (0major+4547minor)pagefaults 0swaps

time ./custom.main
malloc size: 25000000, heap_size: 31327104
0.58user 0.00system 0:00.58elapsed 99%CPU (0avgtext+0avgdata 30592maxresident)k
0inputs+0outputs (0major+7339minor)pagefaults 0swaps

```

System-Level I/O

I think the major good idea in Unix was its clean and simple interface: open, close, read, and write.
by Ken Thompson

10.1 - 10.5 visit book

10.6 - 10.10 visit here

test

code directory: `./code`

test way:

- assert means assert function from `<assert.h>`
- output means to watch code output to judge if it works right

solution	code file	test way
10.6	10.6.c	output
10.7	10.7.c	output
10.8	10.8.c	output
10.9	----	----
10.10	10.10.c	output

$f_d = 4$

10.7

```
/*
 * 10.7.c
 */
#include <stdio.h>
#include "csapp.h"

int main(int argc, char* argv[]) {
    int n;
    char buf[MAXBUF];

    while ((n = Rio_readn(STDIN_FILENO, buf, MAXBUF)) != 0)
        Rio_writen(STDOUT_FILENO, buf, n);

    return 0;
}
```

10.8

```
/*
 * 10.8.c
 */
#include <stdio.h>
#include "csapp.h"

int main(int argc, char* argv[]) {
    struct stat stat;
    char *type, *readok;

    int fd;
    if (argc <= 1)
        fd = 0; // stdin
    else
        fd = atoi(argv[1]);

    Fstat(fd, &stat);

    if (S_ISREG(stat.st_mode))
        type = "regular";
    else if (S_ISDIR(stat.st_mode))
        type = "dir";
    else
        type = "other";
    if ((stat.st_mode & S_IRUSR))
        readok = "yes";
    else
        readok = "no";

    printf("type: %s, read: %s\n", type, readok);

    return 0;
}
```

10.9

```
if (Fork() == 0) {  
    Dup2(0, 3);  
    Execve("fstatcheck", argv, envp);  
}
```

10.10

```
/*
 * 10.10.c
 */
#include <stdio.h>
#include "csapp.h"

int main(int argc, char* argv[]) {
    int n;
    rio_t rio;
    char buf[MAXLINE];

    if (argc == 2) {
        int fd = Open(argv[1], O_RDONLY, 0);
        while ((n = Rio_readn(fd, buf, MAXBUF)) != 0)
            Rio_writen(STDOUT_FILENO, buf, n);
        exit(0);
    }

    Rio_readinitb(&rio, STDIN_FILENO);
    while ((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0)
        Rio_writen(STDOUT_FILENO, buf, n);

    return 0;
}
```

Network Programming

640k is enough for anyone, and by the way, what's a network?

by William Gates III

11.1 - 11.5 visit book

11.6 - 11.13 visit here

test

code directory: `./code`

test way:

- browser means start server and use browser visit server and watch result

solution	code file	test way
11.6	tiny.6.c	browser
11.7	tiny.7.c	browser
11.8	tiny.8.c	browser
11.9	tiny.9.c	browser
11.10	cgi-bin/form-adder.c, tiny.origin.c	browser
11.11	cgi-bin/head-adder.c, tiny.11.c	browser
11.12	cgi-bin/post-adder.c, tiny.12.c	browser
11.13	tiny.13.c	browser

11.6

A.

```

--- tiny.origin.c      2017-11-09 02:57:43.651936112 +0000
+++ tiny.6.c          2017-11-09 02:57:43.651936112 +0000
@@ -13,6 +13,8 @@
void clienterror(int fd, char *cause, char *errnum,
                 char *shortmsg, char *longmsg);

+void echo(int connfd);
+
int main(int argc, char **argv)
{
    int listenfd, connfd;
@@ -33,11 +35,24 @@
    Getnameinfo((SA *) &clientaddr, clientlen, hostname, MAXLINE,
                port, MAXLINE, 0);
    printf("Accepted connection from (%s, %s)\n", hostname, port);
-    doit(connfd);                                //line:netp:tiny:doit
+    echo(connfd);
    Close(connfd);                                //line:netp:tiny:close
}
}

+void echo(int connfd) {
+    size_t n;
+    char buf[MAXLINE];
+    rio_t rio;
+
+    Rio_readinitb(&rio, connfd);
+    while ((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
+        if (strcmp(buf, "\r\n") == 0)
+            break;
+        Rio_writen(connfd, buf, n);
+    }
+}
+
/*
 * doit - handle one HTTP request/response transaction
 */

```

B.

```
./tiny.6 5000
```

run server and firefox visit localhost:5000

```

GET / HTTP/1.1
Host: localhost:5000
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:53.0) Gecko/20100101 Firefox/53.0
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Language: en-US,en;q=0.5
Accept-Encoding: gzip, deflate
Connection: keep-alive

```

C.

HTTP 1.1

D.

visit [rfc](#) section 14 Header Field Definitions

Accept: 14.1

The Accept request-header field can be used to specify certain media types which are acceptable for the response. Accept headers can be used to indicate that the request is specifically limited to a small set of desired types, as in the case of a request for an in-line image.

Accept-Encoding: 14.3

The Accept-Encoding request-header field is similar to Accept, but restricts the content-codings that are acceptable in the response.

Accept-Language: 14.4

The Accept-Language request-header field is similar to Accept, but restricts the set of natural languages that are preferred as a response to the request. Language tags are defined in section 3.10.

Connection: 14.10

The Connection general-header field allows the sender to specify options that are desired for that particular connection and MUST NOT be communicated by proxies over further connections.

Host: 14.23

The Host request-header field specifies the Internet host and port number of the resource being requested, as obtained from the original URI given by the user or referring resource. The Host field value MUST represent the naming authority of the origin server or gateway given by the original URL. This allows the origin server or gateway to differentiate between internally-ambiguous URLs, such as the root “/” URL of a server for multiple host names on a single IP address.

User-Agent: 14.43

The User-Agent request-header field contains information about the user agent originating the request. This is for statistical purposes, the tracing of protocol violations, and automated recognition of user agents for the sake of tailoring responses to avoid particular user agent limitations. User agents SHOULD include this field with requests. The field can contain multiple product tokens (section 3.8) and comments identifying the agent and any subproducts which form a significant part of the user agent. By convention, the product tokens are listed in order of their significance for identifying the application.

11.7

visit [mdn](#) to check all mime type

```
--- tiny.origin.c    2017-11-09 02:57:43.651936112 +0000
+++ tiny.7.c        2017-11-09 02:57:43.651936112 +0000
@@ -173,6 +173,8 @@
     strcpy filetype, "image/png");
     else if (strstr(filename, ".jpg"))
         strcpy filetype, "image/jpeg");
+   else if (strstr(filename, ".mpeg"))
+       strcpy filetype, "video/mpeg");
     else
         strcpy filetype, "text/plain");
 }
```

run server

```
(cd chapter11/code; make && ./tiny.7 5000)
```

browser visit

```
http://localhost:5000/ghost-in-shell.mpeg
```

```

--- tiny.origin.c    2017-11-09 02:57:43.651936112 +0000
+++ tiny.8.c        2017-11-09 02:57:43.651936112 +0000
@@ -12,6 +12,7 @@
void serve_dynamic(int fd, char *filename, char *cgiargs);
void clienterror(int fd, char *cause, char *errnum,
    char *shortmsg, char *longmsg);
+void sigchild_handler(int sig);

int main(int argc, char **argv)
{
@@ -26,6 +27,9 @@
    exit(1);
}

+ if (Signal(SIGCHLD, sigchild_handler) == SIG_ERR)
+     unix_error("signal child handler error");
+
    listenfd = Open_listenfd(argv[1]);
    while (1) {
        clientlen = sizeof(clientaddr);
@@ -38,6 +42,15 @@
    }
}

+void sigchild_handler(int sig) {
+    int old_errno = errno;
+    int status;
+    pid_t pid;
+    while ((pid = waitpid(-1, &status, WNOHANG)) > 0) {
+    }
+    errno = old_errno;
+}
+
/*
 * doit - handle one HTTP request/response transaction
 */
@@ -196,7 +209,6 @@
    Dup2(fd, STDOUT_FILENO);          /* Redirect stdout to client */ //line:netp:serve
dynamic:dup2
    Execve(filename, emptylist, environ); /* Run CGI program */ //line:netp:servedynam
ic:execve
}
- Wait(NULL); /* Parent waits for and reaps child */ //line:netp:servedynamic:wait
}

/*

```


11.9

```
--- tiny.origin.c    2017-11-09 02:57:43.651936112 +0000
+++ tiny.9.c        2017-11-09 02:57:43.651936112 +0000
@@ -152,12 +152,12 @@
     printf("Response headers:\n");
     printf("%s", buf);

- /* Send response body to client */
+ srcfd = Open(filename, O_RDONLY, 0);    //line:netp:servestatic:open
- srcp = Mmap(0, filesize, PROT_READ, MAP_PRIVATE, srcfd, 0); //line:netp:servestatic:m
map
+ srcp = (char*)Malloc(filesize);
+ Rio_readn(srcfd, srcp, filesize);
    Close(srcfd);                        //line:netp:servestatic:close
    Rio_writen(fd, srcp, filesize);      //line:netp:servestatic:write
- Munmap(srcp, filesize);               //line:netp:servestatic:munmap
+ free(srcp);
}

/*
```

11.10

A.

```
<!DOCTYPE html>
<html>
<head>
<meta charset="utf-8" />
<title>Tiny Server</title>
</head>
<body>
<form action="/cgi-bin/form-adder" method="GET">
  <p>first number: <input type="text" name="first"/></p>
  <p>second number: <input type="text" name="second"/></p>
  <input type="submit" value="Submit"/>
</form>
</body>
</html>
```

browser should visit `localhost:5000/cgi-bin/form-adder?first=222&second=333`

B.

`form-adder` handle query string like `first=ddd&second=dddd`

```
--- adder.c      2017-11-09 02:57:43.647936141 +0000
+++ form-adder.c 2017-11-09 02:57:43.647936141 +0000
@@ -1,5 +1,5 @@
/*
- * adder.c - a minimal CGI program that adds two numbers together
+ * form-adder.c - a minimal CGI program that adds two numbers together
 */
#include "../csapp.h"

@@ -12,10 +12,8 @@
if ((buf = getenv("QUERY_STRING")) != NULL) {
    p = strchr(buf, '&');
    *p = '\0';
-   strcpy(arg1, buf);
-   strcpy(arg2, p+1);
-   n1 = atoi(arg1);
-   n2 = atoi(arg2);
+   sscanf(buf, "first=%d", &n1);
+   sscanf(p+1, "second=%d", &n2);
}

/* Make the response body */
```


11.11

from rfc2626 section 9.4 HEAD

The HEAD method is identical to GET except that the server MUST NOT return a message-body in the response.

tiny.c changes

```

--- tiny.origin.c      2017-11-09 02:57:43.651936112 +0000
+++ tiny.11.c         2017-11-09 02:57:43.651936112 +0000
@@ -7,9 +7,9 @@
 void doit(int fd);
 void read_requesthdrs(rio_t *rp);
 int parse_uri(char *uri, char *filename, char *cgiargs);
-void serve_static(int fd, char *filename, int filesize);
+void serve_static(int fd, char *filename, int filesize, char *method);
 void get_filetype(char *filename, char *filetype);
-void serve_dynamic(int fd, char *filename, char *cgiargs);
+void serve_dynamic(int fd, char *filename, char *cgiargs, char *method);
 void clienterror(int fd, char *cause, char *errnum,
                  char *shortmsg, char *longmsg);

@@ -55,7 +55,7 @@
     return;
     printf("%s", buf);
     sscanf(buf, "%s %s %s", method, uri, version);           //line:netp:doit:parserequest
-    if (strcasecmp(method, "GET")) {                          //line:netp:doit:beginrequesterr
+
+    if (!(strcasecmp(method, "GET") == 0 || strcasecmp(method, "HEAD") == 0)) {
         clienterror(fd, method, "501", "Not Implemented",
                     "Tiny does not implement this method");
         return;
@@ -76,7 +76,7 @@
     "Tiny couldn't read the file");
     return;
 }
-    serve_static(fd, filename, sbuf.st_size);                 //line:netp:doit:servestatic
+    serve_static(fd, filename, sbuf.st_size, method);         //line:netp:doit:servestatic
+
+    }
+
+    else { /* Serve dynamic content */
+        if (!(S_ISREG(sbuf.st_mode)) || !(S_IXUSR & sbuf.st_mode)) { //line:netp:doit:exec
+            clienterror(fd, method, "500", "Internal Error",
+                        "Tiny couldn't run the CGI program");
+            return;
+        }
+        serve_dynamic(fd, filename, cgiargs);                 //line:netp:doit:servedynamic
+        serve_dynamic(fd, filename, cgiargs, method);         //line:netp:doit:servedynamic
+    }
+}

```

```

@@ -136,7 +136,7 @@
/*
 * serve_static - copy a file back to the client
 */
-void serve_static(int fd, char *filename, int filesize)
+void serve_static(int fd, char *filename, int filesize, char *method)
{
    int srcfd;
    char *srcp, filetype[MAXLINE], buf[MAXBUF];
@@ -152,6 +152,9 @@
    printf("Response headers:\n");
    printf("%s", buf);

+   if (strcasecmp(method, "HEAD") == 0)
+       return;
+
    /* Send response body to client */
    srcfd = Open(filename, O_RDONLY, 0); //line:netp:serve_static:open
    srcp = Mmap(0, filesize, PROT_READ, MAP_PRIVATE, srcfd, 0); //line:netp:serve_static:m
map
@@ -180,7 +183,7 @@
/*
 * serve_dynamic - run a CGI program on behalf of the client
 */
-void serve_dynamic(int fd, char *filename, char *cgiargs)
+void serve_dynamic(int fd, char *filename, char *cgiargs, char *method)
{
    char buf[MAXLINE], *emptylist[] = { NULL };

@@ -193,6 +196,7 @@
    if (Fork() == 0) { /* Child */ //line:netp:servedynamic:fork
        /* Real server would set all CGI vars here */
        setenv("QUERY_STRING", cgiargs, 1); //line:netp:servedynamic:setenv
+       setenv("REQUEST_METHOD", method, 1);
        Dup2(fd, STDOUT_FILENO); // Redirect stdout to client */ //line:netp:serve
dynamic:dup2
        Execve(filename, emptylist, environ); /* Run CGI program */ //line:netp:servedynam
ic:execve
    }

```

adder.c changes

```
--- adder.c      2017-11-09 02:57:43.647936141 +0000
+++ head-adder.c 2017-11-09 02:57:43.647936141 +0000
@@ -1,10 +1,10 @@
/*
- * adder.c - a minimal CGI program that adds two numbers together
+ * head-adder.c - a minimal CGI program that adds two numbers together
 */
#include "../csapp.h"

int main(void) {
- char *buf, *p;
+ char *buf, *p, *method;
    char arg1[MAXLINE], arg2[MAXLINE], content[MAXLINE];
    int n1=0, n2=0;

@@ -18,6 +18,8 @@
    n2 = atoi(arg2);
}

+ method = getenv("REQUEST_METHOD");
+
    /* Make the response body */
    sprintf(content, "Welcome to add.com: ");
    sprintf(content, "%sTHE Internet addition portal.\r\n<p>", content);
@@ -29,7 +31,10 @@
    printf("Connection: close\r\n");
    printf("Content-length: %d\r\n", (int)strlen(content));
    printf("Content-type: text/html\r\n\r\n");
- printf("%s", content);
+
+ if (strcasecmp(method, "HEAD") != 0)
+     printf("%s", content);
+
    fflush(stdout);

    exit(0);
```

11.12

run server

```
(cd chapter11/code; make && ./tiny.12 5000)
```

visit `http://localhost:5000/post-home.html` and submit

POST method pass param by message-body behind Headers part. When we input 9 and 10 and submit, socket pass content

```
POST /cgi-bin/post-adder HTTP/1.1
Host: localhost:5000
User-Agent: Mozilla/5.0 (X11; Ubuntu; Linux x86_64; rv:53.0) Gecko/20100101 Firefox/53.0
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Language: en-US,en;q=0.5
Accept-Encoding: gzip, deflate
Content-Type: application/x-www-form-urlencoded
Content-Length: 16
Referer: http://localhost:5000/post-home.html
Cookie: _ga=GA1.1.1286836072.1494744693
Connection: keep-alive
Upgrade-Insecure-Requests: 1
(CRLF here, message-body below)
first=9&second=10
```

Content-Length marks the length of message-body. use `readnb` to fetch post param.

```
--- tiny.origin.c    2017-11-09 02:57:43.651936112 +0000
+++ tiny.12.c       2017-11-09 02:57:43.651936112 +0000
@@ -5,7 +5,7 @@
#include "csapp.h"

void doit(int fd);
-void read_requesthdrs(rio_t *rp);
+int read_requesthdrs(rio_t *rp, char *method);
int parse_uri(char *uri, char *filename, char *cgiargs);
void serve_static(int fd, char *filename, int filesize);
void get_filetype(char *filename, char *filetype);
@@ -55,12 +55,14 @@
return;
printf("%s", buf);
sscanf(buf, "%s %s %s", method, uri, version);           //line:netp:doit:parserequest
- if (strcasecmp(method, "GET")) {                          //line:netp:doit:beginrequesterr
+ if (!(strcasecmp(method, "GET") == 0 || strcasecmp(method, "POST") == 0)) {
clienterror(fd, method, "501", "Not Implemented",
"Tiny does not implement this method");
return;
}                                                         //line:netp:doit:endrequesterr
```

```

- read_requesthdrs(&rio); //line:netp:doit:readrequesthdrs

+ int param_len = read_requesthdrs(&rio, method);
+
+ Rio_readnb(&rio, buf, param_len);

/* Parse URI from GET request */
is_static = parse_uri(uri, filename, cgiargs); //line:netp:doit:staticcheck
@@ -84,24 +86,29 @@
    "Tiny couldn't run the CGI program");
    return;
}
- serve_dynamic(fd, filename, cgiargs); //line:netp:doit:servedynamic
+ if (strcasecmp(method, "GET") == 0)
+     serve_dynamic(fd, filename, cgiargs);
+ else
+     serve_dynamic(fd, filename, buf);
}
}

/*
 * read_requesthdrs - read HTTP request headers
 */
-void read_requesthdrs(rio_t *rp)
+int read_requesthdrs(rio_t *rp, char *method)
{
    char buf[MAXLINE];
+ int len = 0;

- Rio_readlineb(rp, buf, MAXLINE);
- printf("%s", buf);
- while(strcmp(buf, "\r\n")) { //line:netp:readhdrs:checkterm
+ do {
    Rio_readlineb(rp, buf, MAXLINE);
    printf("%s", buf);
- }
- return;
+ if (strcasecmp(method, "POST") == 0 && strncasecmp(buf, "Content-Length:", 15) ==
0)
+     sscanf(buf, "Content-Length: %d", &len);
+ } while(strcmp(buf, "\r\n"));
+
+ return len;
}

/*
 * parse_uri - parse URI into filename and CGI args

```

post-adder code

```
/*
 * post-adder.c - a minimal CGI program that adds two numbers together
 */
#include "../csapp.h"

int main(void) {
    char *buf, *p;
    char arg1[MAXLINE], arg2[MAXLINE], content[MAXLINE];
    int n1=0, n2=0;

    /* Extract the two arguments */
    if ((buf = getenv("QUERY_STRING")) != NULL) {
        p = strchr(buf, '&');
        *p = '\0';
        sscanf(buf, "first=%d", &n1);
        sscanf(p+1, "second=%d", &n2);
    }

    /* Make the response body */
    sprintf(content, "Welcome to add.com: ");
    sprintf(content, "%sTHE Internet addition portal.\r\n<p>", content);
    sprintf(content, "%sThe answer is: %d + %d = %d\r\n<p>",
        content, n1, n2, n1 + n2);
    sprintf(content, "%sThanks for visiting!\r\n", content);

    /* Generate the HTTP response */
    printf("Connection: close\r\n");
    printf("Content-length: %d\r\n", (int)strlen(content));
    printf("Content-type: text/html\r\n\r\n");
    printf("%s", content);
    fflush(stdout);

    exit(0);
}
```

11.13

```

--- tiny.origin.c    2017-11-09 02:57:43.651936112 +0000
+++ tiny.13.c       2017-11-09 02:57:43.651936112 +0000
@@ -13,6 +13,17 @@
void clienterror(int fd, char *cause, char *errnum,
                 char *shortmsg, char *longmsg);

+// improved rio written
+void Im_rio_writen(int fd, void *usrbuf, size_t n) {
+  if (rio_writen(fd, usrbuf, n) != n) {
+    if (errno == EPIPE)
+      fprintf(stderr, "EPIPE error");
+
+    fprintf(stderr, "%s ", strerror(errno));
+    unix_error("client side has ended connection");
+  }
+}
+
int main(int argc, char **argv)
{
  int listenfd, connfd;
@@ -26,6 +37,9 @@
  exit(1);
}

+ if (Signal(SIGPIPE, SIG_IGN) == SIG_ERR)
+   unix_error("mask signal pipe error");
+
  listenfd = Open_listenfd(argv[1]);
  while (1) {
    clientlen = sizeof(clientaddr);
@@ -148,7 +162,7 @@
    sprintf(buf, "%sConnection: close\r\n", buf);
    sprintf(buf, "%sContent-length: %d\r\n", buf, filesize);
    sprintf(buf, "%sContent-type: %s\r\n\r\n", buf, filetype);
-   Rio_writen(fd, buf, strlen(buf));           //line:netp:servestatic:endserve
+   Im_rio_writen(fd, buf, strlen(buf));        //line:netp:servestatic:endserve
    printf("Response headers:\n");
    printf("%s", buf);

@@ -156,7 +170,7 @@
    srcfd = Open(filename, O_RDONLY, 0);        //line:netp:servestatic:open
    srcp = Mmap(0, filesize, PROT_READ, MAP_PRIVATE, srcfd, 0); //line:netp:servestatic:m
map
    Close(srcfd);                               //line:netp:servestatic:close
-   Rio_writen(fd, srcp, filesize);             //line:netp:servestatic:write
+   Im_rio_writen(fd, srcp, filesize);          //line:netp:servestatic:write
    Munmap(srcp, filesize);                     //line:netp:servestatic:munmap
  }

@@ -186,11 +200,13 @@

```

```

/* Return first part of HTTP response */
sprintf(buf, "HTTP/1.0 200 OK\r\n");
- Rio_writen(fd, buf, strlen(buf));
+ Im_rio_writen(fd, buf, strlen(buf));
sprintf(buf, "Server: Tiny Web Server\r\n");
- Rio_writen(fd, buf, strlen(buf));
+ Im_rio_writen(fd, buf, strlen(buf));

if (Fork() == 0) { /* Child */ //line:netp:servedynamic:fork
+   if (Signal(SIGPIPE, SIG_DFL) == SIG_ERR)
+       unix_error("unmask signal pipe error");
    /* Real server would set all CGI vars here */
    setenv("QUERY_STRING", cgiargs, 1); //line:netp:servedynamic:setenv
    Dup2(fd, STDOUT_FILENO);          /* Redirect stdout to client */ //line:netp:serve
dynamic:dup2
@@ -216,10 +232,10 @@

    /* Print the HTTP response */
    sprintf(buf, "HTTP/1.0 %s %s\r\n", errnum, shortmsg);
-   Rio_writen(fd, buf, strlen(buf));
+   Im_rio_writen(fd, buf, strlen(buf));
    sprintf(buf, "Content-type: text/html\r\n");
-   Rio_writen(fd, buf, strlen(buf));
+   Im_rio_writen(fd, buf, strlen(buf));
    sprintf(buf, "Content-length: %d\r\n\r\n", (int)strlen(body));
-   Rio_writen(fd, buf, strlen(buf));
-   Rio_writen(fd, body, strlen(body));
+   Im_rio_writen(fd, buf, strlen(buf));
+   Im_rio_writen(fd, body, strlen(body));
}

```

ignore SIGPIPE and show more friendly error message.

unmask SIGPIPE, leave child process to handle it.

Concurrent Programming

Multi-tasking arises out of distraction itself.

by Marilyn vos Savant

12.1 - 12.15 visit book

12.16 - 12.39 visit here

test

prerequisite

- need [wrk](#) to benchmark server in 35, 36, 37, 38
- need package `apache2-utils` (required by command line `ab`) to benchmark proxy server

code directory: `./code`

test way:

- output means to watch code output to judge if it works right
- benchmark means using `wrk/ab` to make lots of connections at same time to test server

solution	code file	test way
12.16	12.16.c	output
12.17	12.17.c	output
12.18	----	----
12.19	12.19.c	output
12.20	12.20.c	output
12.21	12.21.c	output
12.22	12.22.c	using telnet, more visit 12.22.md
12.23	12.23.c	visit 12.23.md
12.24	----	----
12.25	----	----
12.26	12.26.c	output
12.27	----	----
12.28	----	----
12.29	----	----
12.30	----	----
12.31	12.31.c	run and input or BOOM after 5 seconds
12.32	12.32.c	run and input or BOOM after 5 seconds
12.33	12.33.c	run and input or BOOM after 5 seconds
12.34	12.34.c	visit 12.34.md, how to measure performance
12.35	12.35.c	benchmark
12.36	12.36/*	benchmark
12.37	12.37.c	benchmark
12.38	12.38/*	benchmark
12.39	12.39/*	benchmark, more see 12.39.md

12.16

```
/*
 * 12.16.c
 */
#include <stdio.h>
#include "csapp.h"

void *thread(void *vargp);

#define DEFAULT 4

int main(int argc, char* argv[]) {
    int N;
    if (argc > 2)
        unix_error("too many param");
    else if (argc == 2)
        N = atoi(argv[1]);
    else
        N = DEFAULT;

    int i;
    pthread_t tid;
    for (i = 0; i < N; i++) {
        Pthread_create(&tid, NULL, thread, NULL);
    }
    Pthread_exit(NULL);
}

void *thread(void *vargp) {
    printf("Hello, world\n");
    return NULL;
}
```

12.17

A.

main thread didn't wait other thread.

B.

pthread_exit

```
/*
 * 12.17.c
 */
#include "csapp.h"
void *thread(void *vargp);

int main()
{
    pthread_t tid;

    Pthread_create(&tid, NULL, thread, NULL);
    // exit(0);
    Pthread_exit(NULL);
}

/* Thread routine */
void *thread(void *vargp)
{
    Sleep(1);
    printf("Hello, world!\n");
    return NULL;
}
```

12.18

A unsafe

B safe

C unsafe

12.19

pay attention to `static int reader_first`

```
/*
 * 12.19.c
 */
#include <stdio.h>
#include "csapp.h"

#define WRITE_LIMIT 100000
#define PEOPLE 4

static int readtimes;
static int writetimes;
static int readcnt;
// if a reader is waiting when writing, reader first next round
static int reader_first;
sem_t mutex, w;

void *reader(void *vargp) {
    while (1) {
        P(&mutex);
        readcnt++;
        if (readcnt == 1)
            P(&w);
        V(&mutex);

        /* Critical section */
        readtimes++;
        reader_first = 0;
        /* Critical section */

        P(&mutex);
        readcnt--;
        if (readcnt == 0)
            V(&w);
        V(&mutex);
    }
}

void *writer(void *vargp) {
    while (1) {
        if (reader_first == 1)
            continue;

        P(&w);

        /* Critical section */
        writetimes++;
        if (writetimes == WRITE_LIMIT) {
            printf("read/write: %d/%d\n", readtimes, writetimes);
            exit(0);
        }
    }
}
```

```
    }
    /* Critical section */

    // if a reader is waiting, reader first next round
    if (readcnt == 1)
        reader_first = 1;
    V(&w);
}
}

void init(void) {
    readcnt = 0;
    readtimes = 0;
    writetimes = 0;
    reader_first = 0;
    Sem_init(&w, 0, 1);
    Sem_init(&mutex, 0, 1);
}

int main(int argc, char* argv[]) {
    int i;
    pthread_t tid;

    init();

    for (i = 0; i < PEOPLE; i++)
        if (i%2 == 0)
            Pthread_create(&tid, NULL, reader, NULL);
        else
            Pthread_create(&tid, NULL, writer, NULL);

    Pthread_exit(NULL);
    exit(0);
}
```

12.20

`readercnt` limits the readers number at the same time.

```
/*
 * 12.20.c
 */
#include <stdio.h>
#include "csapp.h"

#define WRITE_LIMIT 100000
#define PEOPLE 20 // 10 reader and 10 writer
#define N 5

static int readtimes;
static int writetimes;
sem_t mutex;
sem_t readercnt;

void *reader(void *vargp) {
    while (1) {
        P(&readercnt);
        P(&mutex);

        readtimes++;

        V(&mutex);
        V(&readercnt);
    }
}

void *writer(void *vargp) {
    while (1) {
        P(&mutex);

        writetimes++;
        if (writetimes == WRITE_LIMIT) {
            printf("read/write: %d/%d\n", readtimes, writetimes);
            exit(0);
        }

        V(&mutex);
    }
}

void init(void) {
    readtimes = 0;
    writetimes = 0;
    Sem_init(&mutex, 0, 1);
    Sem_init(&readercnt, 0, N);
}
```



```
int main(int argc, char* argv[]) {
    int i;
    pthread_t tid;

    init();

    for (i = 0; i < PEOPLE; i++) {
        if (i%2 == 0)
            Pthread_create(&tid, NULL, reader, NULL);
        else
            Pthread_create(&tid, NULL, writer, NULL);
    }

    Pthread_exit(NULL);
    exit(0);
}
```

12.21

`writectnt` record how many writers are waiting.

```
/*
 * 12.21.c
 */
#include <stdio.h>
#include "csapp.h"

#define WRITE_LIMIT 100000
#define PEOPLE 4

static int readtimes;
static int writetimes;
static int writectnt;
sem_t mutex, w;

static int number;

void *reader(void *vargp) {
    while (1) {
        // writer first
        if (writectnt > 0)
            continue;

        P(&w);

        /* Critical section */
        readtimes++;
        /* Critical section */

        V(&w);
    }
}

void *writer(void *vargp) {
    while (1) {
        P(&mutex);
        // one more writer wait to write
        writectnt++;
        V(&mutex);

        P(&w);

        /* Critical section */
        writetimes++;
        if (writetimes == WRITE_LIMIT) {
            printf("read/write: %d/%d\n", readtimes, writetimes);
            exit(0);
        }
        /* Critical section */
    }
}
```

```
V(&w);

P(&mutex);
// writer has written
writecnt--;
V(&mutex);
}
}

void init(void) {
    writecnt = 0;
    readtimes = 0;
    writetimes = 0;
    Sem_init(&w, 0, 1);
    Sem_init(&mutex, 0, 1);
}

int main(int argc, char* argv[]) {
    int i;
    pthread_t tid;

    init();

    for (i = 0; i < PEOPLE; i++) {
        if (i%2 == 0)
            Pthread_create(&tid, NULL, reader, NULL);
        else
            Pthread_create(&tid, NULL, writer, NULL);
    }

    Pthread_exit(NULL);
    exit(0);
}
```

12.22

```

#include "csapp.h"

/* read line from connfd and echo line to connfd */
int echo_line(int connfd);

void command(void);

int main(int argc, char **argv)
{
    int listenfd, connfd;
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;
    fd_set read_set, ready_set;

    if (argc != 2) {
        fprintf(stderr, "usage: %s <port>\nuse port 5000 here\n", argv[0]);
        // default port 5000
        listenfd = Open_listenfd("5000");
    } else {
        listenfd = Open_listenfd(argv[1]); //line:conc:select:openlistenfd
    }

    FD_ZERO(&read_set); // Clear read set */ //line:conc:select:clearreadset
    FD_SET(STDIN_FILENO, &read_set); // Add stdin to read set */ //line:conc:select:addstd
    //din
    FD_SET(listenfd, &read_set); // Add listenfd to read set */ //line:conc:select:ad
    //dlistenfd

    // max n for select
    int n = listenfd+1;

    while (1) {
        ready_set = read_set;
        Select(n, &ready_set, NULL, NULL, NULL); //line:conc:select:select

        if (FD_ISSET(STDIN_FILENO, &ready_set)) //line:conc:select:stdinready
            command(); // Read command line from stdin */

        if (FD_ISSET(listenfd, &ready_set)) { //line:conc:select:listenfdready
            clientlen = sizeof(struct sockaddr_storage);
            connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);

            // listen to accepted io ports
            if (connfd+1 > FD_SETSIZE) {
                fprintf(stderr, "too many clients\n");
                Close(connfd);
            }
            n = n > connfd+1 ? n : connfd+1;
            FD_SET(connfd, &read_set);
        }
    }
}

```

```
// echo one line every time
int fd;
for (fd = listenfd+1; fd < n; fd++)
    if (FD_ISSET(fd, &ready_set))
        if (echo_line(fd) == -1) {
            Close(fd);
            FD_CLR(fd, &read_set);
        }
}

void command(void) {
    char buf[MAXLINE];
    if (!Fgets(buf, MAXLINE, stdin))
        exit(0); /* EOF */
    printf("%s", buf); /* Process the input command */
}

int echo_line(int connfd) {
    ssize_t n;
    char buf[1];

    while ((n = Rio_readn(connfd, buf, 1)) > 0) {
        Rio_writen(connfd, buf, n);
        if (buf[0] == '\n')
            return 0;
    }
    return -1;
}
```

run server

```
(cd chapter12/code; make && ./12.22 5000)
```

input command in terminal or telnet to test server

```
telnet 127.0.0.1 5000
```

12.23

code in pic 12-8 save as file `12.23.bug.c`

run server

```
(cd chapter12/code; make && ./12.23.bug)
```

let's figure out when server will fail, see code `12.23.client.c`

```
/*
 * 12.23.client.c - An echo client
 */
#include "csapp.h"

int main(int argc, char **argv)
{
    int clientfd;
    char *host, *port;
    char *buf = "something to send\n";
    rio_t rio;

    host = "127.0.0.1";
    port = "5000";

    clientfd = Open_clientfd(host, port);

    Rio_readinitb(&rio, clientfd);
    Rio_writen(clientfd, buf, strlen(buf));
    /*Close(clientfd);*/
    exit(0);
}
```

line 20 `Close(clientfd)` are commented. when server is running, open another terminal and run `./12.23.client`, server will exit with error

```
Rio_readlineb error: Connection reset by peer
```

server can't read EOF because client doesn't call `Close`

how to fix:

just output error information and doesn't exit

```

--- 12.23.bug.c      2017-11-09 02:57:43.655936083 +0000
+++ 12.23.c         2017-11-09 02:57:43.655936083 +0000
@@ -1,7 +1,5 @@
/*
- * 12.23.bug.c - A concurrent echo server based on select
- *
- * bug in this file
+ * 12.23.c - A concurrent echo server based on select
 */
#include "csapp.h"

@@ -105,15 +103,21 @@
/* If the descriptor is ready, echo a text line from it */
if ((connfd > 0) && (FD_ISSET(connfd, &p->ready_set))) {
    p->nready--;
-    if ((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
+    if ((n = rio_readlineb(&rio, buf, MAXLINE)) > 0) {
        byte_cnt += n; //line:conc:echoservers:beginecho
        printf("Server received %d (%d total) bytes on fd %d\n",
            n, byte_cnt, connfd);
        Rio_writen(connfd, buf, n); //line:conc:echoservers:endecho
    }
-
    /* EOF detected, remove descriptor from pool */
+    else if (n == 0) {
+        Close(connfd); //line:conc:echoservers:closeconnfd
+        FD_CLR(connfd, &p->read_set); //line:conc:echoservers:beginremove
+        p->clientfd[i] = -1; //line:conc:echoservers:endremove
+    }
+    /* n == -1, it's an error */
    else {
+        fprintf(stderr, "error in fd %d, close fd %d connection\n", connfd, connfd);
        Close(connfd); //line:conc:echoservers:closeconnfd
        FD_CLR(connfd, &p->read_set); //line:conc:echoservers:beginremove
        p->clientfd[i] = -1; //line:conc:echoservers:endremove
    }
}

```

12.24

if don't pass pointer param which points to same data block, functions

```
rio_readn  
rio_writen  
rio_readinitb  
rio_readlineb  
rio_readnb
```

are all implicitly reentrant functions

12.25

```

/*
 * A thread-safe version of echo that counts the total number
 * of bytes received from clients.
 */
/* $begin echo_cnt */
#include "csapp.h"

static int byte_cnt; /* Byte counter */
static sem_t mutex; /* and the mutex that protects it */

static void init_echo_cnt(void)
{
    Sem_init(&mutex, 0, 1);
    byte_cnt = 0;
}

void echo_cnt(int connfd)
{
    int n;
    char buf[MAXLINE];
    rio_t rio;
    static pthread_once_t once = PTHREAD_ONCE_INIT;

    Pthread_once(&once, init_echo_cnt); //line:conc:pre:pthreadonce
    Rio_readinitb(&rio, connfd); //line:conc:pre:rioinitb
    while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
        P(&mutex);
        byte_cnt += n; //line:conc:pre:cntaccess1
        printf("server received %d (%d total) bytes on fd %d\n",
            n, byte_cnt, connfd); //line:conc:pre:cntaccess2
        V(&mutex);
        Rio_writen(connfd, buf, n);
    }
}
/* $end echo_cnt */

```

thread safe?

Yes, mutex make it safe

reentrant?

No, share the same mutex

12.26

```

struct hostent *gethostbyname(const char *name)

struct hostent {
    char *h_name;
    char **h_aliases;
    int  h_addrtype;
    int  h_length;
    char **h_addr_list;
}

```

copy int, copy char, *copy char** in struct hostent are 3 different ways.

```

/*
 * 12.26.c
 */
#include <stdio.h>
#include "csapp.h"

/*
 * struct hostent *gethostbyname(const char *name)
 *
 * struct hostent {
 *     char *h_name;
 *     char **h_aliases;
 *     int  h_addrtype;
 *     int  h_length;
 *     char **h_addr_list;
 * }
 */
static sem_t mutex;

static void init_mutex(void) {
    Sem_init(&mutex, 0, 1);
}

struct hostent *gethostbyname_ts(const char *name, struct hostent *host) {
    struct hostent *sharehost;

    P(&mutex);
    sharehost = gethostbyname(name);
    // copy int
    host->h_addrtype = sharehost->h_addrtype;
    host->h_length = sharehost->h_length;
    // copy char *
    host->h_name = (char*)Malloc(strlen(sharehost->h_name));
    strcpy(host->h_name, sharehost->h_name);
    // copy char **
    int i;
    for (i = 0; sharehost->h_aliases[i] != NULL; i++) {}
    host->h_aliases = (char**)Malloc(sizeof(char*) * (i+1));
}

```

```
for (i = 0; sharehost->h_aliases[i] != NULL; i++) {
    // copy every char *
    host->h_aliases[i] = (char*)Malloc(strlen(sharehost->h_aliases[i]));
    strcpy(host->h_aliases[i], sharehost->h_aliases[i]);
}
host->h_aliases[i] = NULL;

for (i = 0; sharehost->h_addr_list[i] != NULL; i++) {}
host->h_addr_list = (char**)Malloc(sizeof(char*) * (i+1));
for (i = 0; sharehost->h_addr_list[i] != NULL; i++) {
    // copy every char *
    host->h_addr_list[i] = (char*)Malloc(strlen(sharehost->h_addr_list[i]));
    strcpy(host->h_addr_list[i], sharehost->h_addr_list[i]);
}
host->h_addr_list[i] = NULL;
V(&mutex);

return host;
}

int main(int argc, char* argv[]) {
    init_mutex();

    struct hostent host;
    gethostbyname_ts("127.0.0.1", &host);
    // result in &host

    return 0;
}
```

12.27

```
FILE *fpin, *fpout;
fpin = fdopen(sockfd, "r");
fpout = fdopen(sockfd, "w");

// read and write

fclose(fpin);
fclose(fpout);
```

ref: 10.11

`fdopen` open 2 stream on same sockfd, `fclose` will close sockfd under stream. if you call `fclose` 2 stream on the same sockfd, the second `fclose` will fail.

image one thread execute code and open 2 stream on fd N. after execution of line `fclose(fpin);`, program create another thread and execute the same code.

but

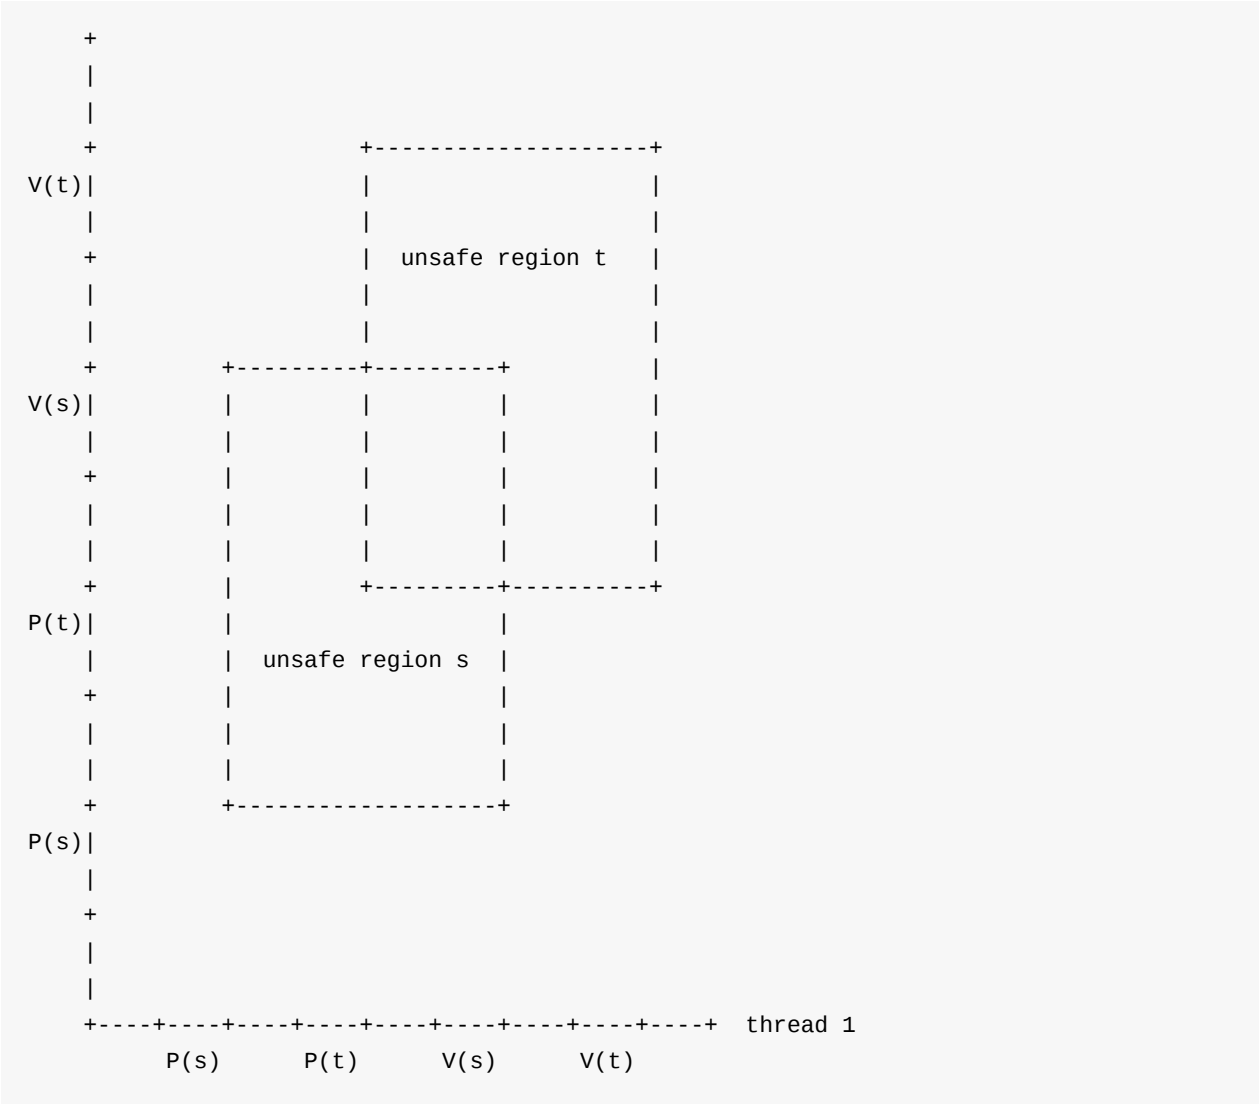
after `fclose(fpin);` in thread 1, fd N is reusable again. assume thread 2 use the fd N again: thread 1 execute line `fclose(fpout);` close the fd that thread 2 is using. it'll cause something unpredicted.

12.28

No effect on deadlock

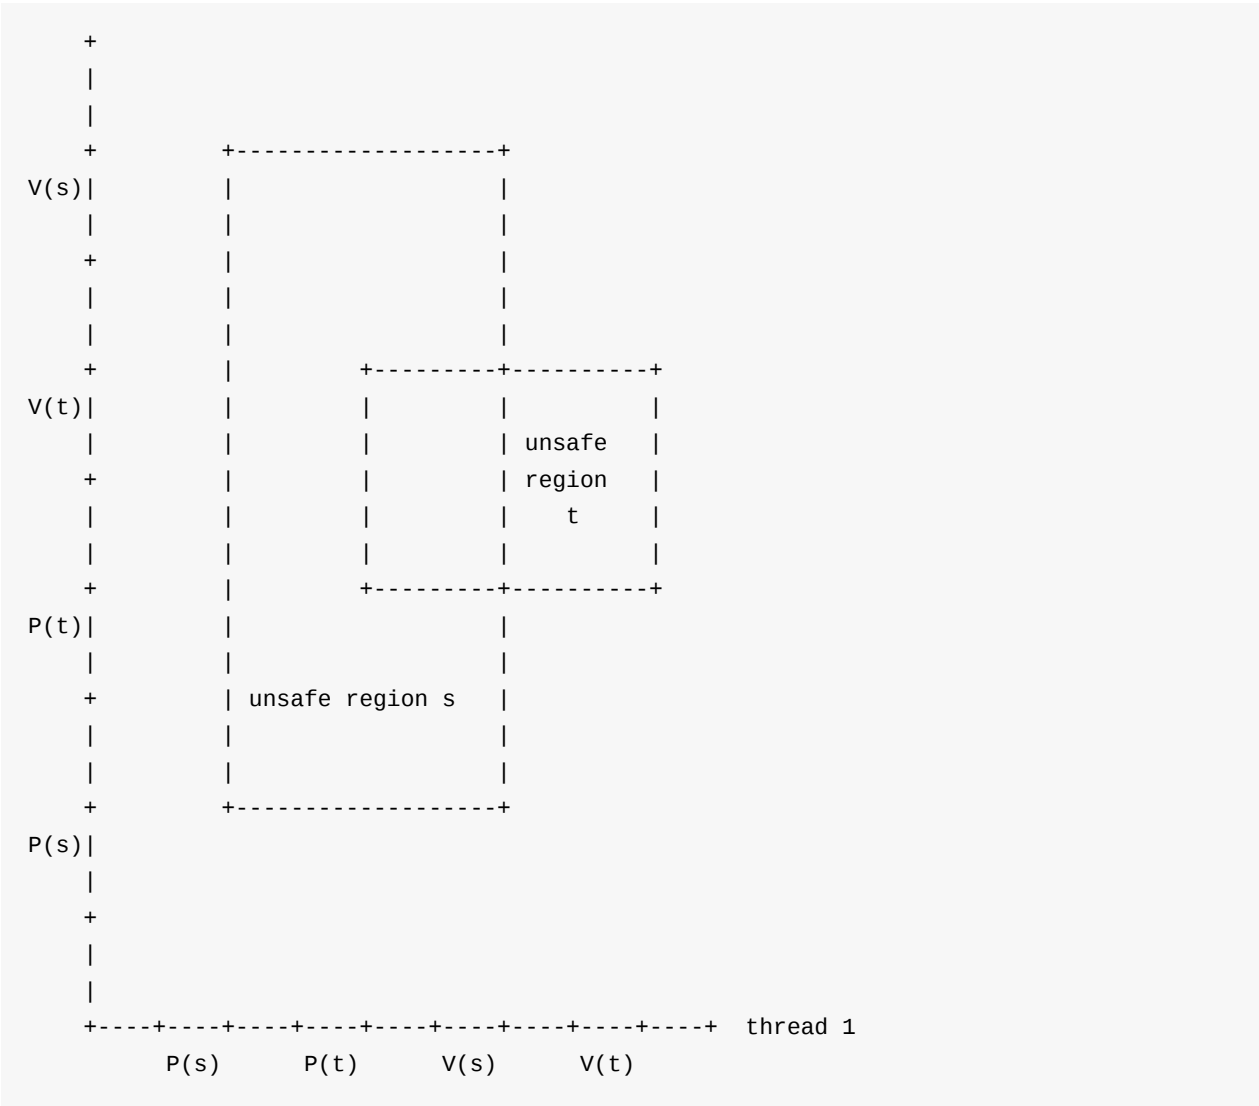
1

thread 1	thread 2
P(s)	P(s)
P(t)	P(t)
V(s)	V(s)
V(t)	V(t)



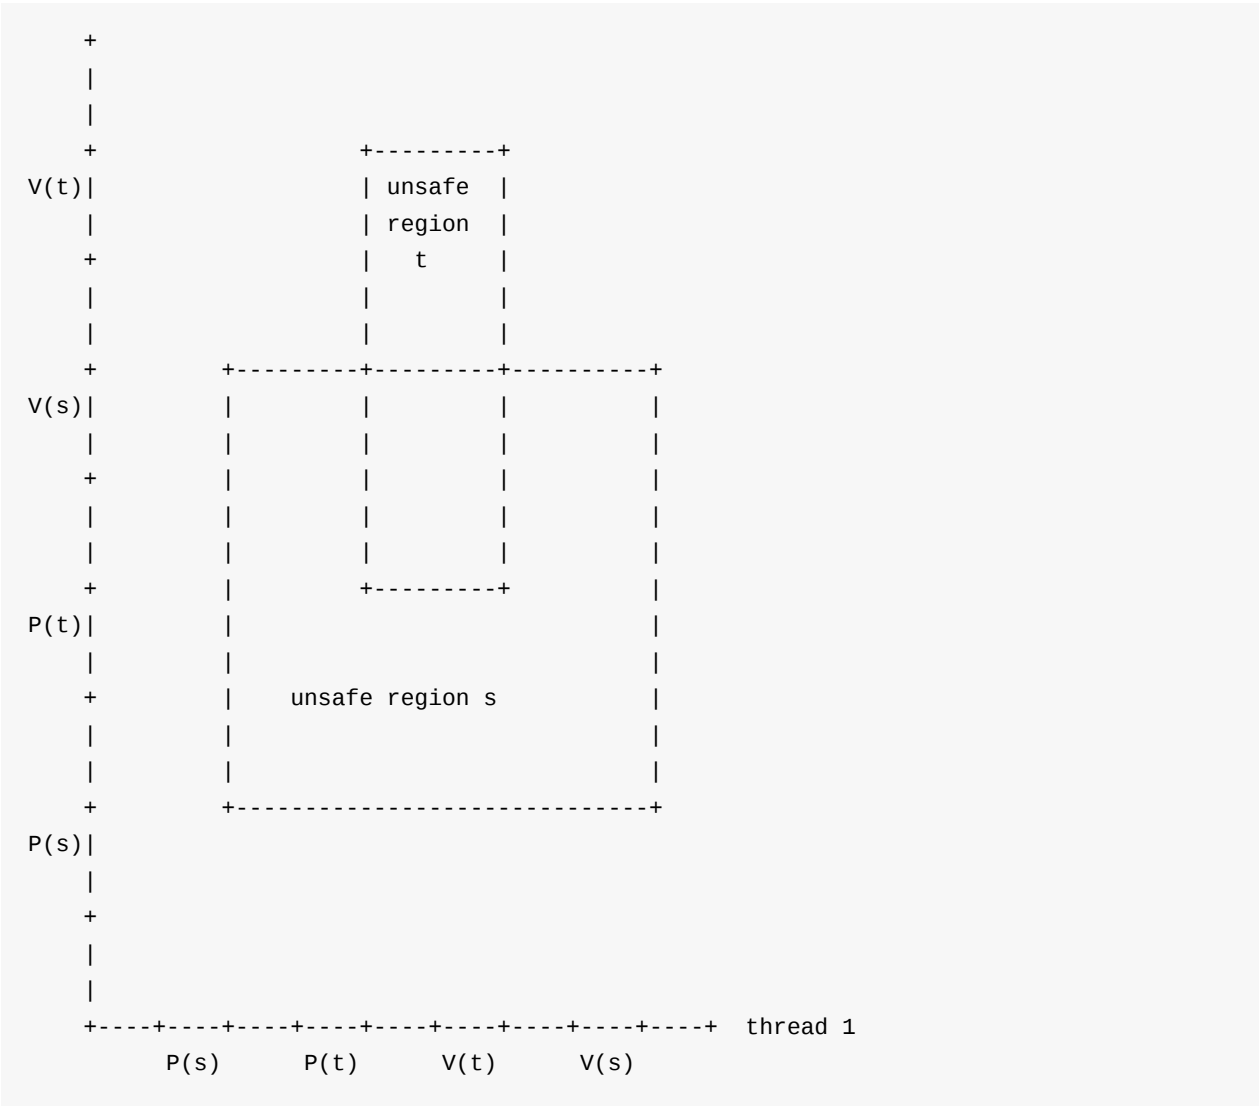
2

thread 1	thread 2
P(s)	P(s)
P(t)	P(t)
V(s)	V(t)
V(t)	V(s)



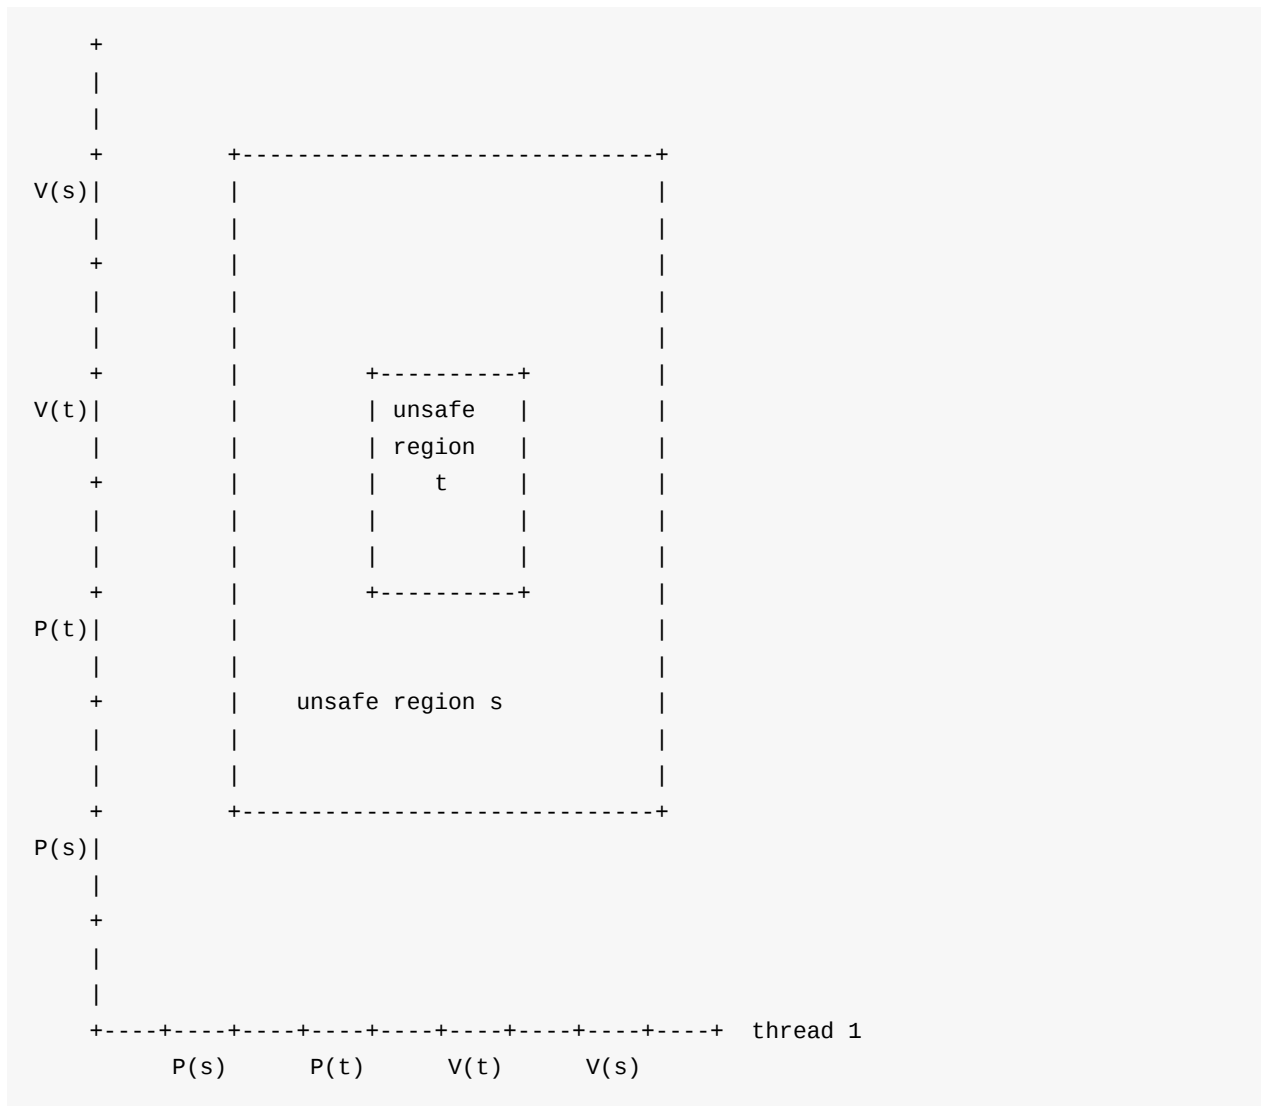
3

thread 1	thread 2
P(s)	P(s)
P(t)	P(t)
V(t)	V(s)
V(s)	V(t)



4

thread 1	thread 2
$P(s)$	$P(s)$
$P(t)$	$P(t)$
$V(t)$	$V(t)$
$V(s)$	$V(s)$



12.29

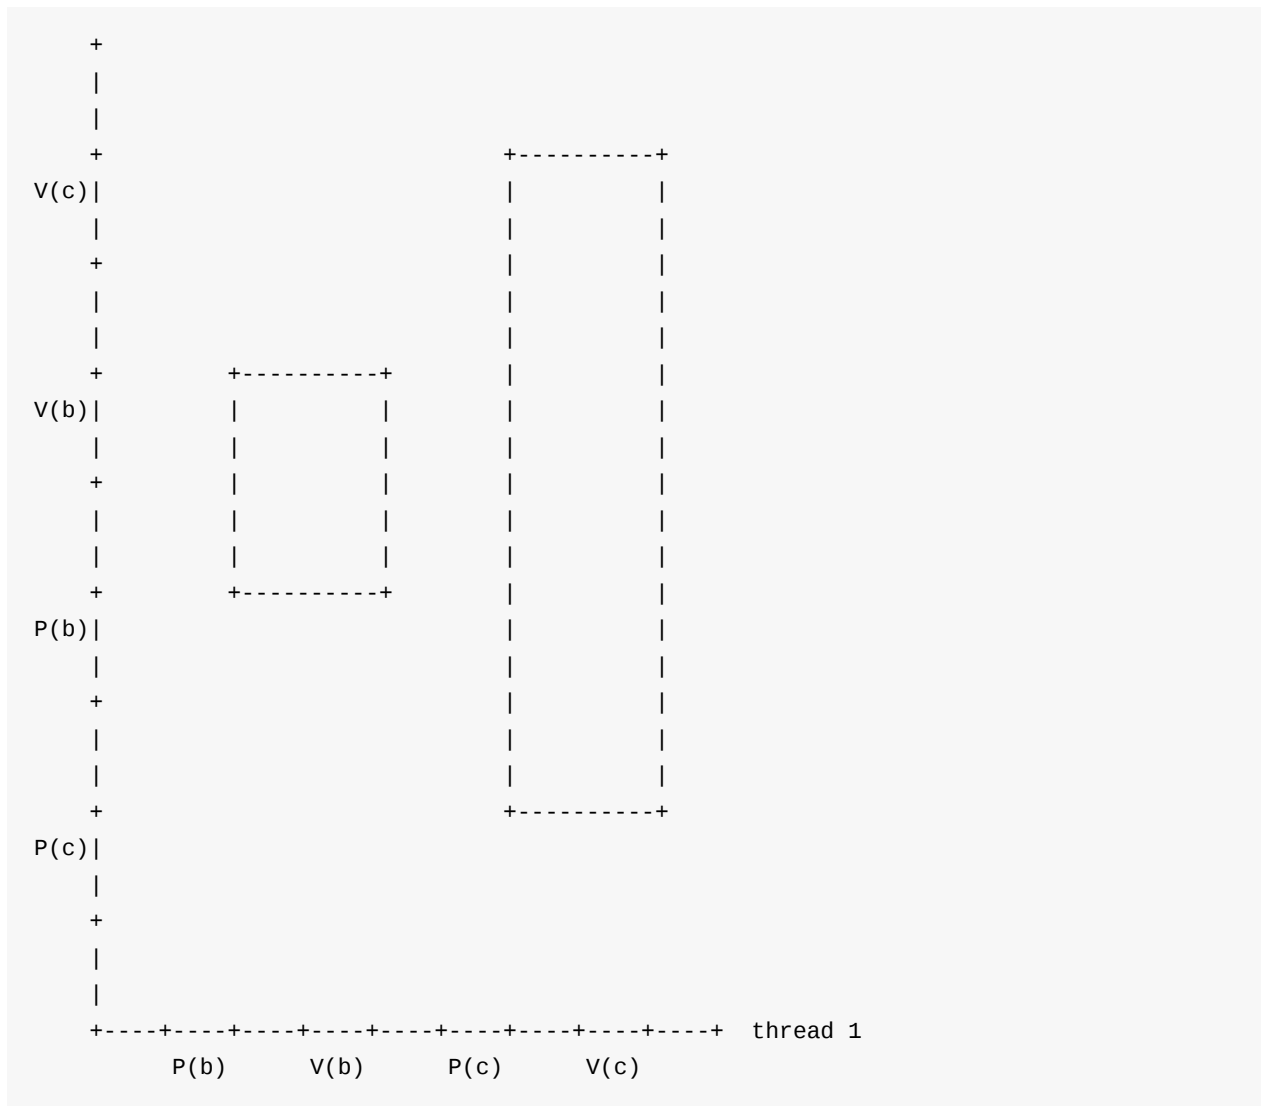
no deadlock

initial: a = 1, b = 1, c = 1

thread 1	thread 2
P(a)	P(c)
P(b)	P(b)
V(b)	V(b)
P(c)	V(c)
V(c)	----
V(a)	----

thread 2 doesn't manipulate mutex a and initial a is 1, so P(a), V(a) don't affect deadlock status.

thread 1	thread 2
P(b)	P(c)
V(b)	P(b)
P(c)	V(b)
V(c)	V(c)



12.30

initial: a = 1, b = 1, c = 1

thread 1	thread 2	thread 3
P(a)	P(c)	P(c)
P(b)	P(b)	V(c)
V(b)	V(b)	P(b)
P(c)	V(c)	P(a)
V(c)	P(a)	V(a)
V(a)	V(a)	V(b)

A.

thread 1: a&b, a&c

thread 2: b&c

thread 3: a&b

B.

thread 2 and thread 3

C.

keep same order P(a), P(b), P(c) in every thread

12.31

```
/*
 * 12.31.c
 */
#include <stdio.h>
#include "csapp.h"

sigjmp_buf buf;

void sigchild_handler(int sig) {
    siglongjmp(buf, 1);
}

char *tfgets(char *s, int size, FILE *stream) {
    if (Fork() == 0) {
        Sleep(5);
        exit(0);
    }

    switch (sigsetjmp(buf, 1)) {
        case 0:
            Signal(SIGCHLD, sigchild_handler);
            return fgets(s, size, stream);
        case 1:
            return NULL;
    }
}

int main(int argc, char* argv[]) {
    char buf[MAXLINE];

    if (tfgets(buf, MAXLINE, stdin) == NULL)
        printf("BOOM!\n");
    else
        printf("%s", buf);

    return 0;
}
```

12.32

key is last param of `select`

```
/*
 * 12.32.c
 */
#include <stdio.h>
#include "csapp.h"

char *tfgets(char *s, int size, FILE *stream) {
    fd_set read_set;
    FD_ZERO(&read_set);
    FD_SET(STDIN_FILENO, &read_set);

    struct timeval timeout;
    timeout.tv_sec = 5;
    timeout.tv_usec = 0;

    Select(1, &read_set, NULL, NULL, &timeout);
    if (FD_ISSET(STDIN_FILENO, &read_set))
        return fgets(s, size, stream);
    else
        return NULL;
}

int main(int argc, char* argv[]) {
    char buf[MAXLINE];

    if (tfgets(buf, MAXLINE, stdin) == NULL)
        printf("BOOM!\n");
    else
        printf("%s", buf);

    return 0;
}
```

12.33

```
/*
 * 12.33.c
 */
#include <stdio.h>
#include "csapp.h"

struct pack {
    char *s;
    int size;
    FILE *stream;
};

char *ptr = NULL;
int timeout = -1;

void *thread_read(void *vargp) {
    struct pack p = *(struct pack *)vargp;
    ptr = fgets(p.s, p.size, p.stream);
    timeout = 0;
}

void *thread_sleep(void *vargp) {
    Sleep(5);
    timeout = 1;
}

char *tfgets(char *s, int size, FILE *stream) {
    pthread_t tid_read;
    pthread_t tid_sleep;
    struct pack p;

    p.s = s;
    p.size = size;
    p.stream = stream;
    Pthread_create(&tid_read, NULL, thread_read, (void*)&p);

    Pthread_create(&tid_sleep, NULL, thread_sleep, NULL);

    // wait 2 thread race result
    while(timeout == -1) {}

    if (timeout == 1) {
        Pthread_cancel(tid_read);
        return NULL;
    } else {
        Pthread_cancel(tid_sleep);
        return ptr;
    }
}

int main(int argc, char* argv[]) {
```

```
char buf[MAXLINE];

if (tfgets(buf, MAXLINE, stdin) == NULL)
    printf("BOOM!\n");
else
    printf("%s", buf);

return 0;
}
```

12.34

matrix size and thread number

```
/*
 * 12.34.h
 */

#define N 640
#define M 640

#define THREAD (1<<4)
#define ROWS_PER_THREAD (N / THREAD)
```

non concurrent version

```
/*
 * 12.34.non.concurrent.c
 */
#include <stdio.h>
#include "csapp.h"
#include "12.34.h"

int M1[N][M];
int M2[N][M];

int MUL12[N][M];

void non_concurrent_mul(void) {
    int i, j, k;
    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++) {
            int sum = 0;
            for (k = 0; k < M; k++) {
                sum += M1[i][k] * M2[k][j];
            }
            MUL12[i][j] = sum;
        }
}

int main(int argc, char* argv[]) {
    non_concurrent_mul();
    return 0;
}
```

concurrent version


```

/*
 * 12.34.concurrent.c
 */
#include <stdio.h>
#include "csapp.h"
#include "12.34.h"

int M1[N][M];
int M2[N][M];

int MUL12[N][M];

void *thread_mul(void *vargp) {
    int idx = *(int*)vargp;
    int start = ROWS_PER_THREAD * idx;
    int i, j, k;
    for (i = start; i < start+ROWS_PER_THREAD; i++)
        for (j = 0; j < N; j++) {
            int sum = 0;
            for (k = 0; k < M; k++) {
                sum += M1[i][k] * M2[k][j];
            }
            MUL12[i][j] = sum;
        }
}

void concurrent_mul(void) {
    pthread_t tid[THREAD];
    int param[THREAD];
    int i;

    for (i = 0; i < THREAD; i++) {
        param[i] = i;
        Pthread_create(&tid[i], NULL, thread_mul, &param[i]);
    }
    for (i = 0; i < THREAD; i++) {
        Pthread_join(tid[i], NULL);
    }
}

int main(int argc, char* argv[]) {
    concurrent_mul();
    return 0;
}

```

measure performance

```
(cd chapter12/code; make clean && make && make measure)
```

output

```
(time ./12.34.non.concurrent)
0.90user 0.00system 0:00.90elapsed 99%CPU (0avgtext+0avgdata 3704maxresident)k
0inputs+0outputs (0major+756minor)pagefaults 0swaps

// single cpu run time 0.9s

(time ./12.34.concurrent)
2.20user 0.00system 0:00.64elapsed 341%CPU (0avgtext+0avgdata 3896maxresident)k
0inputs+0outputs (0major+1462minor)pagefaults 0swaps

// 4 cpu, total run time 2.2s(every cpu run 0.55s, concurrent!!)
```

more detailed

thread(t)	1	2	4	8	16
core(p)	1	2	4	4	4
time(Tp)	0.86	0.466	0.626	0.627	0.628
speedup(Sp)	1	1.84	1.37	1.37	1.37
efficiency(Ep)	100%	92.2%	34.3%	34.3%	34.3%

12.35

see origin tiny server in section 11.6 on book

the key is `close(connfd)` in both parent and child process to ensure close fd and reuse it or it'll failed: `Accept error: too many open files`

```

--- 12.tiny.c      2017-11-09 02:57:43.655936083 +0000
+++ 12.35.c       2017-11-09 02:57:43.655936083 +0000
@@ -1,6 +1,8 @@
/*
- * tiny.c - A simple, iterative HTTP/1.0 Web server that uses the
+ * 12.35.c - A simple, iterative HTTP/1.0 Web server that uses the
+ *      GET method to serve static and dynamic content.
+ *
+ *      concurrent server in multi process way.
+ */
#include "csapp.h"

@@ -35,8 +37,14 @@
    Getnameinfo((SA *) &clientaddr, clientlen, hostname, MAXLINE,
                port, MAXLINE, 0);
    printf("Accepted connection from (%s, %s)\n", hostname, port);
-    doit(connfd);                                //line:netp:tiny:doit
-    Close(connfd);                                //line:netp:tiny:close
+
+    if (Fork() == 0) {
+        Close(listenfd);
+        doit(connfd);                                //line:netp:tiny:doit
+        Close(connfd);                                //line:netp:tiny:close
+        exit(0);
+    }
+    Close(connfd);
}
}

```

run server

```
(cd chapter12/code; make && ./12.35)
```

open another terminal and benchmark it

```
wrk -d4 http://localhost:5000
```

output

```
Running 4s test @ http://localhost:5000
2 threads and 10 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    62.20ms   87.89ms 348.01ms   82.24%
  Req/Sec   397.69    184.95  767.00    58.33%
2585 requests in 4.00s, 578.09KB read
Requests/sec:   645.64
Transfer/sec:   144.39KB
```

12.36

check section 11.6 and 12.2.1, combine code pieces together

code directory: chapter12/code/12.36

origin tiny server: tiny.h tiny.c

select echoserver: echoserver.h echoserver.c

key points:

- move main loop to main.c file
- `check_clients` in echoserver call `doit` (from tiny server) and close fd

benchmark it

```
Running 4s test @ http://localhost:5000
2 threads and 10 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    62.72ms   91.68ms  368.55ms   82.94%
  Req/Sec   615.27    398.96    1.46k    65.91%
3649 requests in 4.00s, 816.04KB read
Requests/sec:    911.45
Transfer/sec:    203.83KB
```

12.37

see origin tiny server in section 11.6 on book

the key is how to pass connfd into thread

```

--- 12.tiny.c      2017-11-09 02:57:43.655936083 +0000
+++ 12.37.c       2017-11-09 02:57:43.655936083 +0000
@@ -13,9 +13,13 @@
void clienterror(int fd, char *cause, char *errnum,
                 char *shortmsg, char *longmsg);

+void *thread(void *vargp);
+
int main(int argc, char **argv)
{
    int listenfd, connfd;
+   int *connfdp;
+   pthread_t tid;
    char hostname[MAXLINE], port[MAXLINE];
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;
@@ -35,11 +39,23 @@
    Getnameinfo((SA *) &clientaddr, clientlen, hostname, MAXLINE,
                port, MAXLINE, 0);
    printf("Accepted connection from (%s, %s)\n", hostname, port);
-   doit(connfd);                                //line:netp:tiny:doit
-   Close(connfd);                               //line:netp:tiny:close
+
+   connfdp = (int*)Malloc(sizeof(int));
+   *connfdp = connfd;
+   Pthread_create(&tid, NULL, thread, connfdp);
}
}

+void *thread(void *vargp) {
+   int connfd = *(int*)vargp;
+   Pthread_detach(Pthread_self());
+   Free(vargp);
+
+   doit(connfd);
+   Close(connfd);
+   return NULL;
+}
+
/*
 * doit - handle one HTTP request/response transaction
 */

```

run server

```
(cd chapter12/code; make && ./12.37)
```

open another terminal and benchmark it

```
wrk -d4 http://localhost:5000
```

output

```
Running 4s test @ http://localhost:5000
 2 threads and 10 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    60.24ms   87.59ms  354.24ms   82.32%
  Req/Sec   481.88    276.48    1.09k    66.67%
 3151 requests in 4.10s, 704.67KB read
Requests/sec:    768.63
Transfer/sec:    171.89KB
```

12.38

code in directory `chapter12/code/12.38`

files:

- `sbuf.h`, `sbuf.c`: prethreading package from section 12.5.4 in book. differences are that add new functions `sbuf_full` , `sbuf_empty` .
- `tiny.h`, `tiny.c`: origin tiny server from section 11.6 in book. differences are we separate function declarations into a `.h` file
- `main.c`: listen new connections, create threads and adjust threads number dynamicly

see new functions in `sbuf`, `sbuf_full` and `sbuf_empty`

```
/* Empty buf? */
int sbuf_empty(sbuf_t *sp) {
    int e;
    P(&sp->mutex);                /* Lock the buffer */
    e = sp->front == sp->rear;
    V(&sp->mutex);                /* Lock the buffer */
    return e;
}

/* Full buf? */
int sbuf_full(sbuf_t *sp) {
    int f;
    P(&sp->mutex);                /* Lock the buffer */
    f = (sp->rear - sp->front) == sp->n;
    V(&sp->mutex);                /* Lock the buffer */
    return f;
}
```

`main.c`

```
/*
 * main.c
 */
#include <stdio.h>
#include "../csapp.h"
#include "tiny.h"
#include "sbuf.h"

#define SBUFSIZE 4
#define INIT_THREAD_N 1
#define THREAD_LIMIT 4096

static int nthreads;
static sbuf_t sbuf; /* Shared buffer of connected descriptors */

// thread info
typedef struct {
    pthread_t tid;
```



```

    sem_t mutex;
} ithread;

static ithread threads[THREAD_LIMIT];

// init work
void init(void);
// function for create server thread
void *serve_thread(void *vargp);
/*
 * creating thread that adjust total thread count according to sbuf situation
 *
 * if sbuf is empty, double threads
 * if sbuf is full, half threads
 */
void *adjust_threads(void *);
// from start to end, create (end - start) new server threads
void create_threads(int start, int end);

int main(int argc, char **argv) {
    int i, listenfd, connfd;
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;
    pthread_t tid;

    if (argc != 2) {
        fprintf(stderr, "usage: %s <port>\n", argv[0]);
        fprintf(stderr, "use default port 5000\n");
        listenfd = Open_listenfd("5000");
    } else {
        listenfd = Open_listenfd(argv[1]);
    }

    init();

    Pthread_create(&tid, NULL, adjust_threads, NULL);

    while (1) {
        clientlen = sizeof(struct sockaddr_storage);
        connfd = Accept(listenfd, (SA *) &clientaddr, &clientlen);
        sbuf_insert(&sbuf, connfd); /* Insert connfd in buffer */
    }
}

void init(void) {
    nthreads = INIT_THREAD_N;
    sbuf_init(&sbuf, SBUFSIZE);

    // create initail server threads
    create_threads(0, nthreads);
}

void *serve_thread(void *vargp) {

```

```

int idx = *(int*)vargp;
Free(vargp);

while (1) {
    // get lock first
    // thread can't be kill now
    P(&(threads[idx].mutex));

    int connfd = sbuf_remove(&sbuf);
    doit(connfd);
    Close(connfd);

    // service ends and release lock
    // so thread can be kill at this time
    V(&(threads[idx].mutex));
}
}

void create_threads(int start, int end) {
    int i;
    for (i = start; i < end; i++) {
        // init mutex for every new thread
        Sem_init(&(threads[i].mutex), 0, 1);
        // create thread
        int *arg = (int*)Malloc(sizeof(int));
        *arg = i;
        // pass thread index in array into thread inside
        Pthread_create(&(threads[i].tid), NULL, serve_thread, arg);
    }
}

void *adjust_threads(void *vargp) {
    sbuf_t *sp = &sbuf;

    while (1) {
        // if sbuf is full, double threads
        if (sbuf_full(sp)) {
            if (nthreads == THREAD_LIMIT) {
                fprintf(stderr, "too many threads, can't double\n");
                continue;
            }

            // double n
            int dn = 2 * nthreads;
            create_threads(nthreads, dn);
            nthreads = dn;
            continue;
        }

        // half threads
        if (sbuf_empty(sp)) {
            if (nthreads == 1)
                continue;
        }
    }
}

```

```
// half n
int hn = nthreads / 2;
/*
 * all server thread are divide to 2 parts
 *
 * keep [0, hn] running
 * kill [hn, nthreads] threads
 *
 * if you want to kill a thread, you must get the lock before it so you
 * won't kill a thread which is offering service.
 */
int i;
for (i = hn; i < nthreads; i++) {
    P(&(threads[i].mutex));
    Pthread_cancel(threads[i].tid);
    V(&(threads[i].mutex));
}
nthreads = hn;
continue;
}
}
```

run server

```
(cd chapter12/code/12.38; make && ./main)
```

open another terminal and benchmark it

```
wrk -d4 http://localhost:5000
```

output

```
Running 4s test @ http://localhost:5000
2 threads and 10 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
Latency    130.44us  183.40us   9.03ms   98.06%
Req/Sec    11.08k     5.09k   17.78k   53.66%
90380 requests in 4.10s, 19.74MB read
Requests/sec: 22039.72
Transfer/sec:    4.81MB
```

12.39

```

+-----+      +-----+      +-----+
|         |----->|         |----->|         |
|client|      |proxy |      |server|
|         |<-----|         |<-----|         |
+-----+      +-----+      +-----+
              as server      as client

```

proxy play with client as a server, with server as a client, interesting.

the key point is when normal situation

```

+-----+      +-----+
|         |----->|         |
|client|      |server|
|         |<-----|         |
+-----+      +-----+

```

HTTP request from client is something like

```

GET / HTTP/1.1
Host: address:port

```

but with proxy, it'll be

```

GET http://address:port/ HTTP/1.1
Host: address:port

```

so proxy have to separate host from uri

A.

```

/*
 * proxy.c
 *
 * visited url log to file log.list
 * block url base on entry from file block.list
 */
#include <stdio.h>
#include "../csapp.h"

// block.list limit entry num
#define MAXENTRY 100

int separate_uri(char *uri, char *host, char *port, char *path);
void parse_block_file(char *filename, char list[MAXENTRY][MAXLINE], int limit);
int blocked_uri(char *uri, char list[MAXENTRY][MAXLINE]);

int main(int argc, char **argv) {

```

```

int i, listenfd, connfd;
int clientfd;

socklen_t clientlen;
struct sockaddr_storage clientaddr;

rio_t client_rio, server_rio;
char c_buf[MAXLINE], s_buf[MAXLINE];
ssize_t sn, cn;

char method[MAXLINE], uri[MAXLINE], version[MAXLINE];
char host[MAXLINE], port[MAXLINE], path[MAXLINE];

char block_list[MAXENTRY][MAXLINE];
int logfd;
char log_buf[MAXLINE];

if (argc != 2) {
    fprintf(stderr, "usage: %s <port>\n", argv[0]);
    fprintf(stderr, "use default port 5000\n");
    listenfd = Open_listenfd("5000");
} else {
    listenfd = Open_listenfd(argv[1]);
}

logfd = Open("log.list", O_WRONLY | O_APPEND, 0);

memset(block_list, '\0', MAXLINE * MAXENTRY);
parse_block_file("block.list", block_list, MAXENTRY);

while (1) {
    // wait for connection as a server
    clientlen = sizeof(struct sockaddr_storage);
    connfd = Accept(listenfd, (SA *) &clientaddr, &clientlen);
    Rio_readinitb(&server_rio, connfd);

    /*
     * if uri is full path url like http://localhost:8000/server.c
     * remove host part http://localhost:8000
     * only pass /server.c to server
     */
    // parse HTTP request first line
    if (!Rio_readlineb(&server_rio, s_buf, MAXLINE)) {
        Close(connfd);
        continue;
    }
    sscanf(s_buf, "%s %s %s", method, uri, version);
    // if uri is blocked?
    if (blocked_uri(uri, block_list)) {
        printf("%s is blocked\n", uri);
        Close(connfd);
        continue;
    }
}

```

```

// log visit
sprintf(log_buf, "visit url: %s\n", uri);
Write(logfd, log_buf, strlen(log_buf));

memset(host, '\0', MAXLINE);
memset(port, '\0', MAXLINE);
memset(path, '\0', MAXLINE);
int res;
if ((res = separate_uri(uri, host, port, path)) == -1) {
    fprintf(stderr, "not http protocol\n");
    Close(connfd);
    continue;
} else if (res == 0) {
    fprintf(stderr, "not a absolute request path\n");
    Close(connfd);
    continue;
}

// connect server as a client
clientfd = Open_clientfd(host, port);
Rio_readinitb(&client_rio, clientfd);

/*
 * browser --> proxy --> server
 *
 * send requests
 */
// write first request line
sprintf(s_buf, "%s %s %s\n", method, path, version);
Rio_writen(clientfd, s_buf, strlen(s_buf));
printf("%s", s_buf);
do {
    // pass next http requests
    sn = Rio_readlineb(&server_rio, s_buf, MAXLINE);
    printf("%s", s_buf);
    Rio_writen(clientfd, s_buf, sn);
} while(strcmp(s_buf, "\r\n"));

/*
 * server --> proxy --> browser
 *
 * server send response back
 */
while ((cn = Rio_readlineb(&client_rio, c_buf, MAXLINE)) != 0)
    Rio_writen(connfd, c_buf, cn);

Close(connfd);
Close(clientfd);
}

Close(logfd);
}

```

```
/*
 * if uri is absolute path url like
 *   http://localhost:8888/something
 *   or
 *   http://localhost/something (port default is 80)
 * separate into three part and return 1
 *
 * if uri is relative path like /something
 * do nothing and return 0
 *
 * if uri is absolute path and not http protocol like https/ftp/etc
 * do nothing, return -1, it's error
 */
int separate_uri(char *uri, char *host, char *port, char *path) {
    // relative path
    if (uri[0] == '/')
        return 0;

    // absolute path
    char *prefix = "http://";
    int prelen = strlen(prefix);
    // if not http protocol, error
    if (strncmp(uri, prefix, prelen) != 0)
        return -1;

    char *start, *end;
    start = uri + prelen;
    end = start;

    // copy host
    while (*end != ':' && *end != '/') {
        end++;
    }
    strncpy(host, start, end-start);

    // port is provided
    if (*end == ':') {
        // skip ':'
        ++end;
        start = end;
        // copy port
        while (*end != '/')
            end++;
        strncpy(port, start, end-start);
    } else {
        // port is not provided, default 80
        strncpy(port, "80", 2);
    }

    // copy path
    strcpy(path, end);
}
```

```

/*
 * read block file, parse all the entries and save into list
 * entries count no more than limit
 */
void parse_block_file(char *filename, char list[MAXENTRY][MAXLINE], int limit) {
    int blockfd;
    char block_buf[MAXLINE];
    rio_t block_rio;
    ssize_t block_n;

    blockfd = Open(filename, O_RDONLY, 0);
    Rio_readinitb(&block_rio, blockfd);

    memset(block_buf, '\0', MAXLINE);
    int num = 0;
    while ((block_n = Rio_readlineb(&block_rio, block_buf, MAXLINE)) != 0) {
        // exceed limit
        if (num == limit)
            break;

        // right entry
        if (strncmp(block_buf, "http://", 7) == 0) {
            strcpy(list[num], block_buf);
            num++;
        }

        // comment or not right format entry
        // do nothing
    }

    Close(blockfd);
}

/*
 * if uri is in list, return true
 * if not, return false
 */
int blocked_uri(char *uri, char list[MAXENTRY][MAXLINE]) {
    int i;
    for (i = 0; list[i][0] != '\0'; i++)
        if (strncmp(uri, list[i], strlen(uri)) == 0)
            return 1;

    return 0;
}

```

B.

move process into thread function

```

/*
 * proxy-thread.c multi thread deal with concurrent

```



```

*
* visited url log to file log.list
* block url base on entry from file block.list
*/
#include <stdio.h>
#include "../csapp.h"

// block.list limit entry num
#define MAXENTRY 100

int separate_uri(char *uri, char *host, char *port, char *path);
void parse_block_file(char *filename, char list[MAXENTRY][MAXLINE], int limit);
int blocked_uri(char *uri, char list[MAXENTRY][MAXLINE]);
void *proxy_thread(void *vargp);

// url black list
static char block_list[MAXENTRY][MAXLINE];
// log file fd
static int logfd;

int main(int argc, char **argv) {
    int listenfd;
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;
    int *connfdp;
    pthread_t tid;

    if (argc != 2) {
        fprintf(stderr, "usage: %s <port>\n", argv[0]);
        fprintf(stderr, "use default port 5000\n");
        listenfd = Open_listenfd("5000");
    } else {
        listenfd = Open_listenfd(argv[1]);
    }

    logfd = Open("log.list", O_WRONLY | O_APPEND, 0);
    memset(block_list, '\0', MAXLINE * MAXENTRY);
    parse_block_file("block.list", block_list, MAXENTRY);

    while (1) {
        // wait for connection as a server
        clientlen = sizeof(struct sockaddr_storage);
        connfdp = Malloc(sizeof(int));
        *connfdp = Accept(listenfd, (SA *) &clientaddr, &clientlen);
        // new thread
        Pthread_create(&tid, NULL, proxy_thread, connfdp);
    }

    Close(logfd);
}

void *proxy_thread(void *vargp) {
    pthread_t tid = Pthread_self();

```

```

Pthread_detach(tid);

int connfd = *(int*)vargp;
Free(vargp);

rio_t client_rio, server_rio;
char c_buf[MAXLINE], s_buf[MAXLINE];
ssize_t sn, cn;

char method[MAXLINE], uri[MAXLINE], version[MAXLINE];
char host[MAXLINE], port[MAXLINE], path[MAXLINE];

char log_buf[MAXLINE];

int clientfd;
Rio_readinitb(&server_rio, connfd);
/*
 * if uri is full path url like http://localhost:8000/server.c
 * remove host part http://localhost:8000
 * only pass /server.c to server
 */
// parse HTTP request first line
if (!Rio_readlineb(&server_rio, s_buf, MAXLINE)) {
    Close(connfd);
    return NULL;
}
sscanf(s_buf, "%s %s %s", method, uri, version);
// if uri is blocked?
if (blocked_uri(uri, block_list)) {
    printf("thread %ld: %s is blocked\n", tid, uri);
    Close(connfd);
    return NULL;
}
// log visit
sprintf(log_buf, "thread %ld: visit url: %s\n", tid, uri);
Write(logfd, log_buf, strlen(log_buf));

memset(host, '\0', MAXLINE);
memset(port, '\0', MAXLINE);
memset(path, '\0', MAXLINE);
int res;
if ((res = separate_uri(uri, host, port, path)) == -1) {
    fprintf(stderr, "tid %ld: not http protocol\n", tid);
    Close(connfd);
    return NULL;
} else if (res == 0) {
    fprintf(stderr, "tid %ld: not a absolute request path\n", tid);
    Close(connfd);
    return NULL;
}

// connect server as a client
clientfd = Open_clientfd(host, port);

```

```

Rio_readinitb(&client_rio, clientfd);

/*
 * browser --> proxy --> server
 *
 * send requests
 */
// write first request line
sprintf(s_buf, "%s %s %s\n", method, path, version);
Rio_writen(clientfd, s_buf, strlen(s_buf));
printf("tid %ld: %s", tid, s_buf);
do {
    // pass next http requests
    sn = Rio_readlineb(&server_rio, s_buf, MAXLINE);
    printf("tid %ld: %s", tid, s_buf);
    Rio_writen(clientfd, s_buf, sn);
} while(strcmp(s_buf, "\r\n"));

/*
 * server --> proxy --> browser
 *
 * server send response back
 */
while ((cn = Rio_readlineb(&client_rio, c_buf, MAXLINE)) != 0)
    Rio_writen(connfd, c_buf, cn);

Close(connfd);
Close(clientfd);
}

/*
 * if uri is absolute path url like
 * http://localhost:8888/something
 * or
 * http://localhost/something (port default is 80)
 * separate into three part and return 1
 *
 * if uri is relative path like /something
 * do nothing and return 0
 *
 * if uri is absolute path and not http protocol like https/ftp/etc
 * do nothing, return -1, it's error
 */
int separate_uri(char *uri, char *host, char *port, char *path) {
    // relative path
    if (uri[0] == '/')
        return 0;

    // absolute path
    char *prefix = "http://";
    int prelen = strlen(prefix);
    // if not http protocol, error
    if (strncmp(uri, prefix, prelen) != 0)

```

```

    return -1;

    char *start, *end;
    start = uri + prelen;
    end = start;

    // copy host
    while (*end != ':' && *end != '/') {
        end++;
    }
    strncpy(host, start, end-start);

    // port is provided
    if (*end == ':') {
        // skip ':'
        ++end;
        start = end;
        // copy port
        while (*end != '/')
            end++;
        strncpy(port, start, end-start);
    } else {
        // port is not provided, default 80
        strncpy(port, "80", 2);
    }

    // copy path
    strcpy(path, end);
}

void parse_block_file(char *filename, char list[MAXENTRY][MAXLINE], int limit) {
    int blockfd;
    char block_buf[MAXLINE];
    rio_t block_rio;
    ssize_t block_n;

    blockfd = Open(filename, O_RDONLY, 0);
    Rio_readinitb(&block_rio, blockfd);

    memset(block_buf, '\0', MAXLINE);
    int num = 0;
    while ((block_n = Rio_readlineb(&block_rio, block_buf, MAXLINE)) != 0) {
        // exceed limit
        if (num == limit)
            break;

        // right entry
        if (strncmp(block_buf, "http://", 7) == 0) {
            strcpy(list[num], block_buf);
            num++;
        }

        // comment or not right format entry
    }
}

```

```
// do nothing
}

Close(blockfd);
}

/*
 * if uri is in list, return true
 * if not, return false
 */
int blocked_uri(char *uri, char list[MAXENTRY][MAXLINE]) {
    int i;
    for (i = 0; list[i][0] != '\0'; i++)
        if (strncmp(uri, list[i], strlen(uri)) == 0)
            return 1;

    return 0;
}
```

how to benchmark:

run proxy

```
(cd chapter12/code/12.39; make && ./proxy)
# or
(cd chapter12/code/12.39; make && ./proxy-thread)
```

open another terminal, start a local http server

```
(python3 -m http.server 8000)
```

open one another terminal, benchmark

```
# -X means proxy, -n means how many requests, -c means concurrent
(ab -X localhost:5000 -n 100 -c 10 http://localhost:8000/)
```